

Rahnuma

रहनुमा رہنما

A working guide to the ILM, PANINI and Zistgah stack, from bytes and scripts to compilers, Pāṇini and quantum substrates

प्रकृति

	written	Devanagari hub	Romenagri	and back
Devanagari	प्रकृति	→ प्रकृति	→ prak_ri_ti	→ प्रकृति
Telugu	ప్రకృతి	→ प्रकृति	→ prak_ri_ti	→ ప్రకృతి
Bengali	প্রকৃতি	→ प्रकृति	→ prak_ri_ti	→ the same identity
Urdu, as usually written	ہندوی	→ हनदवी	→ _hana_dawee	→ हनदवी
Urdu, with zer and sukun	ہندوی	→ हिन्दवी	→ _hin_dawee	→ हिन्दवी

One identity, many scripts. Telugu and Bengali meet Devanagari in one hub, become the same Romenagri word, and come back exactly. Urdu meets the hub at what is written: without its short-vowel marks the word reads hanadawee; with zer and sukun it is hindawi. Every cell here was produced by the Hindawi tools while this guide was built.

Abhishek Choudhary, AyeAI. Version 2.0.0, 3 October 2026.

Contents

Part I: Orientation 3

1. How to use this guide	4
2. The work, not the worker	5
3. Start here: find your way in	6

Part II: One architecture 8

4. Where the work comes from	9
5. The record, with dates	11
6. Architecture: core, projections, layers	13
7. The dependency graph	18
8. The ecosystem, term by term	20
9. The three Acts	24
10. Three families, set apart: the AyeAI Triad, GramSheel and Project VIKRAM, Humanesque and its projections	26
11. Projections: names, and raising a component into Humanesque	30
12. The repositories	34

Part III: Setting up 40

13. Choosing a platform	41
14. The baseline environment and a first run	43
15. Running every component	46
16. Labs: containers now, launch scripts next	55

Part IV: Foundations 56

17. Numbers inside the machine	57
18. Characters inside the machine	59
19. Writing systems, sound, and reversible transliteration	62
20. Pāṇini's grammar as a formal system	68
21. Compilers, from lexer to microcode	72
22. Below the software: logic, chips, ternary and qubits	81
23. Quantum computing and PANINIQ, with the mathematics	83

Part V: Working with the estate 86

24. AI, from MYCIN to transformers, and how to use any AI here	87
25. Open source, provenance and contributing	89
26. GATE: sit it, whatever your age	91
27. Syllabus links, and our syllabus	93
28. Questions that come up, answered technically	97
29. The roadmap, for every collaborator and every domain	101
30. What the work would cost: a model-based estimate	106

Part VI: Status 107

31. This release and the next	108
-------------------------------------	-----

Appendices 110

A. The estate's terms, and the nearest common terms	111
B. Reading the architecture without flattening it	113
C. For researchers in the humanities and social sciences	116
D. For researchers in science, technology, engineering and mathematics	118
E. For practitioners: teachers, policy makers, communities, makers and engineers	120
F. Formal definitions	121
G. Review register	124

Part I: Orientation

1. How to use this guide

In plain words. This guide explains a large body of work on computers, languages and intelligence. You can read it, run the programs on your own computer, or ask an AI to help you, and every example in it was really run.

This guide is for anyone who wants to use, check or extend the work gathered in the [zistgah](#), [project-ilm](#) and [hindawiai](#) organisations on GitHub: students preparing for GATE, engineers who write compilers, firmware or chip layouts, linguists who work with scripts and sound, and researchers in quantum computing and AI. It asks for curiosity and a laptop, and nothing else.

There are four ways in.

- **Understand.** Part II gives the whole architecture first: where the work comes from, how the parts compose, each term with its nearest analogue and the difference, and the three Acts.
- **Read.** Parts IV and V explain the ideas the work rests on, with the mathematics written out: how numbers and letters sit in memory, what makes a transliteration reversible, how Pāṇini's grammar computes, how a compiler turns text into machine instructions, and how an oscillator network and a quantum circuit attack the same problem.
- **Run.** Part III gives the exact commands to set up Linux, Windows with WSL, or macOS, and to run every component of the estate on your own machine.
- **Ask.** Chapter 24 shows how to put any AI assistant to work on these repositories, and how to check what it tells you.

Conventions used throughout:

- Every example that could run while this guide was being built did run, and the text shows its real output, marked *Executed while this guide was built*. The package's own gate, `bash ops/verify.sh`, runs those examples again and fails if any output differs from what is printed here.
- Status words are used exactly. **Tested:** a test in the repository asserts it. **Executed:** it ran and its output is shown. **Written:** the code exists but has not been run here. **Mocked:** a stand-in replaces hardware that is not attached. **Planned:** described, not built.
- Everything said about the repositories was retrieved from the repositories themselves on 29 September 2026: their listings on GitHub, their `CITATION.cff` files and their READMEs. Where a repository does not record something, this guide says so instead of filling the gap.
- A dagger, †, marks a statement that follows the author's working notes and write-ups as supplied on 29 September and 3 October 2026, rather than the repositories.
- Sanskrit is transliterated in IAST. The spelling is British.

The work lives at github.com/zistgah, github.com/project-ilm and github.com/hindawiai, on the site zistgah.org, and in the Zenodo records listed in Chapter 12.

2. The work, not the worker

In plain words. What matters is whether the work holds up, whether the code runs and the reasoning is sound, not who made it. The aim is that people build technology in their own languages.

We know Pāṇini almost entirely through his grammar. Tradition gives a birthplace, Śalātura in Gandhāra, and very little else that can be checked. Close to four thousand sūtras survive; the person does not. The grammar still works, and Chapter 20 runs part of it. That is the standard this guide keeps. Work has to stand without its author, because titles, affiliations and biographies do not compile, and code, derivations and measurements do.

Two consequences follow for how the estate is built.

The first is enablement. Linguistic and cultural equity is not achieved by making AI available everywhere. Systems have to work within the languages, scripts, cultural contexts and knowledge traditions of the people they serve, so that those people become producers of technology in their own languages rather than remaining users of someone else's.

The second is openness. Every repository in Chapter 12 can be read, run, forked and criticised. Nothing here needs anyone's permission to use; each repository states its licence.

The programme stands on two pillars. The linguistic pillar standardises at the level of the dialect, with the idiolect as a later overlay, and grows through linguists working across the world. The technological pillar, the older of the two, is now the enabler: the transliteration engines, compilers, parsers and substrates that make a language usable for building. This guide deals mostly with the second pillar, because that is where a newcomer can start contributing today.

3. Start here: find your way in

In plain words. Four steps to find your part of the guide: who you are, your board or field, what you want to do, and your path with a PDF made for you.

This guide serves a school student and a senior researcher; a village health worker and an investor. Nobody needs all of it. Four steps find your part, on the site or with this page:

1. **Who you are.** A school student, a college student, a postgraduate or doctoral researcher; a teacher; an academic researcher; a working professional, someone in public service, policy or law, someone in community, NGO or village work, an artist, writer or maker; an industry partner or investor; or an AI agent working for a person.
2. **Your board or your field.** A school student chooses a board, ICSE, ISC, CBSE, IGCSE, IB, a state board or another, and a class. Everyone else chooses a field from the international standard classification of fields of study, ISCED-F 2013, the same one FAKIR's lattice uses. It places each subject where its own scholars would: political science is under social and behavioural sciences, not the humanities; law is under business, administration and law; linguistics is under languages.
3. **What you want to do.** Learn the basics, run the software, study or research, teach, contribute, or adopt, fund or invest.
4. **Your path and your edition.** The site lists your chapters in order, with reading times and the components to try first, and offers a PDF edition made for you. Every edition below is a printed selection from this one guide, so all of them stay in step.

Edition	PDF
School student	role-school.pdf
College student	role-college.pdf
Postgraduate or doctoral researcher	role-pg.pdf
Teacher or lecturer	role-teacher.pdf
Academic researcher	role-researcher.pdf
Working professional or engineer	role-professional.pdf
Public service, policy or law	role-public.pdf
Community, NGO or village work	role-community.pdf
Artist, writer or maker	role-creator.pdf
Industry partner, investor or funder	role-industry.pdf
Generic programmes and qualifications	field-00.pdf
Education	field-01.pdf
Arts and humanities	field-02.pdf
Social sciences, journalism and information	field-03.pdf
Business, administration and law	field-04.pdf
Natural sciences, mathematics and statistics	field-05.pdf
Information and Communication Technologies (ICTs)	field-06.pdf
Engineering, manufacturing and construction	field-07.pdf
Agriculture, forestry, fisheries and veterinary	field-08.pdf
Health and welfare	field-09.pdf

Edition	PDF
Services	field-10.pdf

The first public version of this guide, written for advanced researchers and senior professionals, stays as it was at [version 1](#).

Part II: One architecture

4. Where the work comes from

In plain words. This work began long before today's AI. It grew along five connected lines: languages and computers, learning machines, robots and medicine, organised knowledge, and the future of people and their homes.

This work did not begin with today's language models, and it is not a collection of separate projects. It is one architecture built along several lines of work that started at different times, and it reads correctly only in that order. The dated record, with the place where each entry can be checked, is the next chapter; this one gives the lines themselves.

Five lines that intersect

Linguistic and computational identity. Romenagri (2003), the reversible ASCII-7 kernel, came first; the Hindawi Programming System (15 August 2004) built a complete systems-programming stack on it; ILM, Integrative Linguistic Multiscript, succeeds Hindawi and generalises it to every script and language; PANINI's construct model now sits above ILM as the common core of programming, prompting and realization.




Figure 1. PEDLER's branches: cognition, embodiment, physics and silicon.

Cognition and event-driven computation. PEDLER, the Point Event-Driven Learner (November 2001, Indian patent application 3033/CHE/2011), is not one genealogy but two branches: a cognitive branch through LVF and MLCNE to the cognitive enablement modules, CEMb and CEMs, and an embodiment branch through NI2A2 (2003) to the cognitive robot ANGEL (2003). QEDLER carries PEDLER into physics. PRATIK and Zamin take its balanced-ternary, event-driven logic towards silicon.

Embodiment, medicine and field systems. HMSEI, a medical wearable (2002); ANGEL; TARA; RDK, among Nokia's global top ten in 2010; Dr Rho, the TARA and RDK medical telepresence platform, Highly Commended at the IET Innovation Awards in 2014; GramSheel's Village Knowledge Center, a Stockholm Challenge finalist in 2010. The AyeAI Triad closes this line formally: AyeAM for embodiment, AyeAI for cognition, AyeCNSe for coordination and communication.

Knowledge and reasoning. UKOP, the reference architecture for human knowledge across ISIC, ISCO and ISCED; FAKIR, its kernel; Dhancha, the spine every domain is built on; the domains themselves; and the working layer that runs on them: GENIE, the cyclers, AAB and Mez.

Continuity and civilisation. Humanesque, its constitutional derivation from the Proclamation of Individual Equity, TransEg, the transferred alter ego, and the habitats on Earth, the Moon and Mars.

The lines cross constantly. PEDLER supplies the Act, the executable intent that the constitutional corpus is built from; Romenagri and ILM carry every language into FAKIR's lattice; PANINI runs the cyclers and also drives the realization compilers; TransEg receives what yadein records. That is why the chapters that follow teach the pieces as nodes of one graph rather than as a list.

The invariant

† What the architecture preserves, whatever changes around it, is meaning, intent, identity, agency, capability, provenance, equity and continuity, while the representation, the language, the script, the formalism, the machine, the substrate, the embodiment, the social setting and eventually the habitat all change. Romenagri is the smallest instance of the principle: a name keeps its identity through every tool of a toolchain. TransEg is the largest: a person's constituted identity keeps its continuity beyond the substrate that first carried it.

Sovereignty and equity

The Proclamation of Individual Equity, PoIE (zistgah/poie, doi:10.5281/zenodo.21397274), makes the individual the primary locus of equity, identity and agency. Sovereignty is recursive from there: individual, family, community, institution, province, nation and beyond, across political, economic, social and technological dimensions. Technological sovereignty, in this sense, is the capacity to determine, govern, develop, reproduce, modify and deploy one's own domain; the precise, testable form of it for software is given in Chapter 21. The aim that follows is enablement: that people become producers, builders, researchers and owners of computing in their own languages, not permanent consumers of systems built elsewhere.

The orientation map

His poster "The Expanse of Our Epistemic Realms" runs one axis from the inner to the outer: being, meaning, knowledge, language, cognition, code, systems, society, civilisation, planet. At its apex stands Kaivalyik AGI; beneath it, the Equitable Mind, one humanity with one identity and infinite expressions; and the rule that every sentient individual, bioform or synthetiform, is a sovereign. Read the rest of this Part along that axis: language and identity first, then computation, cognition, knowledge and agency, realization, and the Acts that carry them to civilisational scale.

5. The record, with dates

In plain words. A list of what was made and when, from 1993 onwards, with where each item can be checked.

Dates settle questions of priority without argument, so the record of the work is given here with the place where each entry can be checked.

When	What	Where to check
1993	The Nature Welfare Society, which became the GramSheel Foundation	The author's account
September 1998	Labyrinths of Xanadu, a game, on the PCQuest cover CD	The author's provenance record
November 2001	PEDLER, the Point Event-Driven Learner; Indian patent application 3033/CHE/2011	doi:10.5281/zenodo.17497559
2002	HMSEI, a medical wearable	The author's provenance record
2003	Romenagri, the reversible transliteration kernel, released under the GPL; ANGEL, a cognitive robot; NI2A2, the thesis built on PEDLER	doi:10.5281/zenodo.17254607 for NI2A2
15 August 2004	Hindawi Programming System, first public release	sourceforge.net/projects/hindawi/ ; Wikipedia
2005	CSI National Young IT Professional Award; Infocomm showcase	Certificate held by the author
December 2005	GNU Savannah project "Hindawi Vernacular Programming System"; the savannah-register list archive timestamps Romenagri	savannah.nongnu.org/projects/hindawi
19 December 2005	Announcement to the Sarai PRC list of Hindi and Bangla equivalents of C, C++, lex, yacc, assembly and Java, with Hindi and Bangla DOS, BASIC and LOGO; 7-bit ASCII left unaltered; GCC as the host compiler; filters between ISCII, Romenagri, Unicode, APCISR and HP-PCL	The ILUGC list archive
2005 to 2006	Sarai/CSDS FLOSS Fellowship	The author's provenance record
2006	LinuxAsia showcase	The author's provenance record
2007	Manthan Award nomination	Cited in the Wikipedia article
April 2008	"Hindawi: From Assembler to Lisp in Indic", <i>Linux For You</i>	The magazine's April 2008 issue
2008	FOSS India Award; Stockholm Challenge nomination for Hindawi; profile in <i>DataQuest</i>	Certificate held by the author; the magazine's archive
15 September 2010	Nokia GEVC global top ten: RDK	The PRNewswire release of that date
2010	Stockholm Challenge finalist: GramSheel Foundation, Village Knowledge Center (Universal Device)	Certificate held by the author
2011	Intel India Embedded Challenge	Certificate held by the author
2014	IET Innovation Awards, finalist and Highly Commended: Dr Rho, the TARA and RDK medical telepresence platform	Certificate held by the author
3 May 2020	Project VIKRAM launched as the Project Vikram Journal by the GramSheel Foundation	pvjournal.github.io
2026	The repositories and records of this estate	Chapter 12

Hindawi was also reviewed in *Linux Magazine*, issue 76, and has been cited in academic work, including at the University of the Punjab and the University of Ibadan. Its current source trees are in the [hindawiai](#) organisation, and its description, with the SourceForge home, is on the [AyeAI site](#). Every item the author publishes carries his ORCID,

0009-0002-0684-8320.

6. Architecture: core, projections, layers

In plain words. One shared core, Humanesque, is seen through three cultures: Indian, Persian and Western. Around it sit the parts that turn an idea, in any language, into something real.

One core, three projections

Humanesque is the core: one shared, multidimensional, executable typed hypergraph, with a federated ontology of 26 positions, O_0 to O_{25} , a common execution and evidence spine, and components that can be implemented independently. Its GitHub organisation is hmnsq, where every core stem repository is to live.




Figure 2. Humanesque and its projections: parallel cultural traversals, not stages.

Kaivalyik, Zistgah and Cosmopolis are three projections over it, running in parallel: the Indic, including the Advaita framing; the Persian and Islamicate; and the Greek, Western and classical. A projection is a domain-specific view and traversal of the whole architecture, with its own culturally sensitive naming and branding. The projections are not stages, not levels of a hierarchy and not versions of one another; they intersect, diverge and reconverge, and none completes another. Zistgah is not the parent of the other two. Their organisations are kaivalyik, zistgah and c-polis; today zistgah carries the working repositories.

The PANINI stack

Layer	What it holds	Where to look
Front end	The languages a program or a prompt can arrive in, Sanskrit first: ILM and the language parsers	project-ilm/ilm.codes, zistgah/tajziya

Layer	What it holds	Where to look
Middleware	PANINI's own front ends, from <code>ada.pni</code> to <code>zig.pni</code> and <code>sysml.pni</code> , and the PANINI language itself	<code>zistgah/panini</code> , <code>zistgah/humanesque</code>
Core PANINI	The construct model: common semantics plus decorators	<code>zistgah/humanesque</code> , <code>ilm/constructs.csv</code> and <code>ilm/decorators.csv</code>
Realization backends	Where a program meets a substrate: host toolchains for code; PANINIq for oscillator and quantum substrates; PANINib for biology; PANINIphy for physical systems	<code>zistgah/paniniq</code> , <code>zistgah/paninib</code> , <code>zistgah/paniniphy</code>




Figure 3. The PANINI stack, with Romenagri as the script layer beneath every front end.

Beneath every front end lies the script layer, Romenagri. Every keyword and every name of a program becomes a word over A to Z, a to z and the underscore, reversibly, so that the whole existing toolchain, from lexer to linker to debugger, carries it unchanged, and the inverse renders it back in the script wherever a person reads it. Chapters 18, 19 and 21 explain it and run it.

A construct such as a counted loop is defined once in the core and translated once per human language. A decorator carries what a particular host language drags in with it, such as `stdio.h` for printing in C. Core PANINI, in his words, “does not violate anything”; the front ends and the realization backends are all projections of it.

How the parts compose

Humanesque	the core; Kaivalyik, Zistgah and Cosmopolis project it in parallel; everything below lives in it, nothing below leads to it
front end	ILM (script language standard) and tajziya, Sanskrit first the script axis is Romenagri, which persists as the symbol bridge through compiler, linker, ELF, DWARF and debugger
middleware	PANINI's own front ends and the PANINI language
core	core PANINI: construct model = common semantics + decorators
realization	host toolchains for code; PANINIq (oscillator, quantum);

Around the stack: the cyclers are written in PANINI and run by one engine; GENIE composes and runs research cycles; AAB paints tasks and runs the process cyclers; FAKIR, UKOP and Dhancha supply the domain lattice and its invariants; Mez, the Cognitive Workbench, composes any of them for an exercise without absorbing them; the provenance systems seal every artifact; VGC gates the work between people and AI. Two readings must be kept apart. As genealogy, the order is Romenagri, Hindawi, ILM, PANINI and the realization arms. As architecture, Romenagri sits inside ILM's script axis, the cyclers are written in PANINI rather than downstream of it, and Humanesque is the core that contains the whole, not its end point.

The realization spine

The realization repositories state the spine they share. PANINI lowers declarative intent through language, semantic IR, a realization requirement, domain resolution and a domain IR, and then a common layer of geometry, trajectory, animation, observation, verification and an ArtifactGraph. PANINIB lowers biological intent through semantic, resolution, CISC, RISC and sequence IRs and emits two coupled artifacts, a physical realization plan and a digital execution run; PANINIphy lowers physical intent through a physical IR and physical-domain engines. "IR contracts are the product. Laboratories and machines are external adapters." And the same repositories keep the states of a result apart: desired, predicted, observed and validated phenotype are four different things, as are compilation success, execution success, target satisfaction and biological validation.




Figure 4. The realization spine the realization arms share.

The habitat elements

Zamin is the ground: the balanced-ternary hardware of PRATIK. AAB is the water: the paint program for systems, from painted intent to verified, sealed artifact, where the verification-gated quests run. Fiza is the air: the environmental replica. Chakra is the turning sky: the observatory and its calendars. All of them mount in one shared virtual dome, which the habitats on the Moon and Mars replicate.

The cyclers

A cycler is a sequence of prompts and the outputs they produce, written down so that it can be edited, shared and run with any AI. Six are classified by what they produce: matba (print), khwab (visual), awaz (audio), tilasm (immersive), pench (embodied) and yadein (the record); genie runs research cycles. Each has its own purpose, contract, context, state model, invariants, failure modes, evidence requirements, workflow and artifact model; only the engine is shared. The protocol behind every cycle, in the author's words: intent, then context, then a meaningful prompt; any AI answers and the human inspects; the artifacts and responses are kept under configuration management; the next prompt follows, until a final artifact that the human authored with intention.




Figure 5. The cyclers: one engine, configured in PANINI, and the protocol every cycle follows.

Provenance, in a fixed order

Four systems, never collapsed into one. **Tok DOI** proves that a thing existed by a given time. **spiguard** is the disclosure gate, which fails closed. **Candor** records signed intent as in-toto statements. **Misty DOI** is the one that reaches the world, through Zenodo. The order is always seal, clear, attest, mint. Chapter 25 explains the machinery.





Figure 6. Provenance in its fixed order: seal, clear, attest, mint.

7. The dependency graph

In plain words. A map of how the parts connect, like a city map with roads between places. Any route you follow is one journey; the map is the whole picture.

The estate is a typed graph, and any account of it in words is one traversal through that graph. Read a traversal, then return to the graph: adjacency is not dependency, a projection is not a hierarchy, composition is not absorption, and an Act is not a stage of maturity.

Every node and every edge below is taken from the retrieved record: the repositories, their READMEs and the author's statements. Clusters group the nodes by family, not by layer; and an edge label names the relation it records. The graph's source is `data/graph.json`; `python3 tools/graph.py` renders it, and the build keeps the rendered `dependency-graph.svg` and its Graphviz source `dependency-graph.dot` beside this guide.



Three traversals

A name, from a Devanagari source to a debugger. Romenagri is the script layer of HPS and the script axis of ILM; through the symbol bridge it carries every name into the host toolchain, and the inverse brings it back (Chapters 19 and 21).

A task, from a domain and a language to a sealed artifact. FAKIR seeds the domain and ILM the language; AAB builds; the cyclers, written in the PANINI language and composed in Mez, produce the work; matba publishes it through Kitab; Tok DOI, spiguard, Candor and Misty DOI seal, clear, attest and mint it, in that order; VGC gates the work between people and AI throughout.

A mind, from an event to a substrate. PEDLER's cognitive branch runs through LVF and MLCNE to the cognitive enablement modules, which Interglial Healthcare operates on the biological side; its embodiment branch runs through NI2A2 to ANGEL; its balanced-ternary logic runs into PRATIK and down to Zamin, the ground of the dome; and its engine drives PANINIq.

8. The ecosystem, term by term

In plain words. A dictionary of the special words used here, each with the nearest everyday idea and what makes it different.

Each entry gives what the thing is in this architecture, the nearest analogue a newcomer is likely to reach for, and the exact difference. The analogue is a door into the idea, never a definition of it. The state in brackets says what exists: established (historical record), released (on the forge, with its DOI where minted), executed (run for this guide), under construction, specification, or planned. Entries marked † follow the author's working notes and write-ups as supplied on 29 September and 3 October 2026, rather than the repositories.

Language and identity

ILM, Integrative Linguistic Multiscript (released; ilm.codes live). *In the architecture:* the linguistic pillar and the identity layer for language, with script, language and standard kept as separate axes, and layers for phonology, transliteration, orthography, lexicon, syntax and semantics, and interfaces, supported by data, validation, development tools, language specifications and a language server; its registry lists 7,867 languages and 226 scripts. *Nearest analogue:* localisation, an NLP toolkit, a transliteration library. *Difference:* it keeps the three axes apart, so that any language can travel in any script, and it gives every script a reversible path through the whole computing stack, down to the symbol table.

Romenagri (established 2003; released; executed). *In the architecture:* the reversible kernel between scripts and ASCII-7. Its paper states six constraints met together: ASCII-7 closure, case independence, diacritic independence, legality as a C identifier, reversibility and linear time, over the Brahmi, Perso-Arabic and Northwest Semitic families, with a canonical-form layer and an ASCII-reduction layer. *Nearest analogue:* romanisation, such as IAST, ISO 15919 or ITRANS. *Difference:* its output is a valid identifier for every tool from lexer to debugger, and it inverts, so a name survives compilation and comes back.

HPS, the Hindawi Programming System (established 15 August 2004; released; executed). *In the architecture:* nine shailis over nine host languages, from assembly, lex and yacc to C++, Java, Python, BASIC and LOGO; a script layer, a language layer and the host standard, composed as transducers over unmodified host toolchains; now HPS's Indic projection under ILM, with an Urdu edition beside it. *Nearest analogue:* a localised programming language. *Difference:* a complete systems stack, not a keyword translation, with names carried through Romenagri.

tajziya (released, doi:10.5281/zenodo.22982238). *In the architecture:* the parser federation for the classical languages, by family, with the author's Sanskrit parser; the front end of the PANINI stack beside ILM. *Nearest analogue:* NLP parsers. *Difference:* organised to feed PANINI's construct model and AAB's paintings.

Computation

PANINI (several members released; some under construction). *In the architecture:* the construct model at the core, with ILM and tajziya in front, PANINI's own front ends and language as middleware, and realization backends behind; it runs through all three Acts. Its family includes the prompt-cycle language in which the cyclers are written (zistgah/panini), the general-purpose PANINI of the merged release (zistgah/humanesque), and the realization arms. *Nearest analogue:* a compiler; a programming language, an intermediate representation. *Difference:* programming languages, prompt cycles and physical, biological and quantum realizations are all projections of one construct model.

PANINIq (released, doi:10.5281/zenodo.23020648; executed). *In the architecture:* a realization backend: an oscillator substrate, the PEDLER engine and quantum validation. *Nearest analogue:* a quantum simulator. *Difference:* the substrate is an oscillator Ising machine with a learning automaton on it; the quantum circuit validates it rather than replacing it.

PANINIB and PANINiphy (released research platforms, version 1.0). *In the architecture:* PANINIB, the biological realization arm, lowers biological intent through semantic, resolution, CISC, RISC and sequence IRs to a realization plan and an execution run; PANINiphy, the physical realization arm, lowers physical intent through a physical IR and

physical-domain engines, down to modular blocks and the selection of parts by force, stroke and voltage. *Nearest analogue*: computer-aided design, genome design tools. *Difference*: both share PANINI's common spine, language, verification, provenance and ArtifactGraph, adding only their own domain IR and engines; laboratories and machines are external adapters, and neither claims a result its evidence does not show.

Cognition

PEDLER and QEDLER (established November 2001; record doi:10.5281/zenodo.17497559; QEDLER a research framework). *In the architecture*: the point-event-driven learner; a six-tuple extension of the Turing machine over balanced ternary; the source of the primitives Act (executable intent), Inclination (a signed directional gradient on transition edges) and Natural Justice (a gate on the passage from intent to Act); QEDLER extends it to physics. *Nearest analogue*: an event-driven or neuromorphic learning model. *Difference*: it predates the present generation of models and carries the executable-intent semantics the constitutional corpus is built from.

PRATIK and Zamin (released: pratik_core_mvp doi:10.5281/zenodo.21288232, zamin doi:10.5281/zenodo.21297556). *In the architecture*: PRATIK, Participatory Recursive Adaptive Trans-Intelligence Kernels, on the divider between Acts I and II, with a kernel in C++ and CUDA; Zamin, its physical balanced-ternary substrate with a poised zero. *Nearest analogue*: a neuromorphic kernel on novel hardware. *Difference*: balanced ternary and event-driven from the logic up.

CEM, with CEMb, CEMs and Eco-CEM (specification; Act II). *In the architecture*: cognitive enablement modules over one kernel, $\mathcal{E} = (E, C, \Pi, W)$, realised on biological substrates (BCI, prosthetics, closed-loop cognition, operated through Interglial Healthcare), on synthetic ones (processors, distributed systems, robotics) and on ecological ones. *Nearest analogue*: robotics, BCI, prosthetics. *Difference*: one invariant kernel across substrates; the superscript is part of the name.

The AyeAI Triad (specification with formal closure). *In the architecture*: AyeAM = $\langle S, R, C \rangle$ for embodiment, AyeAI = $\langle M, I, G \rangle$ for cognition, AyeCNSe = $\langle T, Ch, \Sigma \rangle$ for coordination and communication across media, with AyeAI at the apex. *Nearest analogue*: a perception, cognition and action loop. *Difference*: a triad, not a loop and not a pipeline.

Knowledge and agency

UKOP, FAKIR and Dhancha (FAKIR released, doi:10.5281/zenodo.21436550; Dhancha released, doi:10.5281/zenodo.22821645; UKOP a specification). *In the architecture*: UKOP, the reference architecture for human knowledge; FAKIR, its kernel over ISIC, ISCO and ISCED crossed with AGI layers and language; Dhancha, the spine every domain is built on, whose rule is that the engine is common and the workflow is not. *Nearest analogue*: a knowledge graph or reasoning engine. *Difference*: domain invariants are enforced by tests, and resolution and verification are primitives.

Cyclers (six released and minted; Act I). *In the architecture*: matba, khwab, awaz, tilasm, pench and yadein, classified by what they produce, plus genie, each with its own contract, state model and evidence rules, written in PANINI and run by one engine. *Nearest analogue*: an agentic harness or orchestration loop. *Difference*: AI-agnostic by construction, inspected by a person at every step, and the recorded method, never the content, is itself the reproducible work. † An older form of the corpus names SAFAR in the sixth position, where the current corpus has yadein.

Kitab (released). *In the architecture*: the config-driven book template that matba publishes through, and † in the author's notes the persistent artifact and publication layer that keeps an artifact's state, its creative lineage, its execution trace and evaluations, its forks and descendants. *Nearest analogue*: a publishing template. *Difference*: a cycler publishes through Kitab; Kitab is not a seventh cycler.

Research Kundali (released). *In the architecture*: a research map of a person, a lab or a paper; built from an ORCID, a lab name or a DOI (project-ilm/research-kundali); † in the author's notes, the systematic record of claims, artifacts, chronology, provenance, objections, evidence and responses, so that every claim can be checked against its evidence. *Nearest analogue*: a researcher profile. *Difference*: it records evidence and discrepancies, not reputation.

GENIE (released). *In the architecture*: the Generalized Emotive-Narrative Interaction Engine, storyteller, poet and painter, and the research cycle through its Prompt Operating System, with six verbs (create, verify, execute, measure, falsify, integrate), nine epistemic tags and a gate that a simulation cannot pass in place of an experiment. † In the author's own

formulation GENIE is itself a cyler; the creative cyler; which can decorate the modality cyclers without erasing their identities. *Nearest analogue*: a creative or research AI agent. *Difference*: governed by a constitution of primitives; orchestration and composition, not an implementation.

AAB (released; Act I). *In the architecture*: the paint program for systems, from painted intent to verified, sealed artifact; the gamified studio of verification-gated quests; home of the process cyclers; the process half of the estate's software factory, with FAKIR as its component registry. In the author's words, AAB and FAKIR together constitute the working definition of AGI used here. *Nearest analogue*: a low-code studio or a software factory. *Difference*: every task is a painting with a manifest, quest stages and oracle-gated verification, and every component keeps its provenance.

Mez, the Cognitive Workbench (released). *In the architecture*: the local-first desk that brings independently existing systems together for an exercise: composition, exercise, observation, synthesis, a new artifact. *Nearest analogue*: an IDE or a notebook. *Difference*: composition without absorption; every composed system keeps its own identity and life.

Method and evidence

VGC, Verification-Gated Human-AI Co-Development (released, doi:10.5281/zenodo.21264248). *In the architecture*: the method by which people and AI build together; with every step gated by verification. *Nearest analogue*: test-driven development or human review of AI output. *Difference*: the gate is a verified artifact, not a reviewer's impression, and the AI's report that something passed is never the evidence that it passed.

COPA, the Cost of Perceived Authority (released, doi:10.5281/zenodo.21782217). *In the architecture*: a protocol for measuring how much an AI system asserts before it verifies, through the Authority Projection Index, with its hypothesis stated beside a fair null. *Nearest analogue*: an AI evaluation benchmark. *Difference*: it measures projected authority rather than accuracy, and its target is calibrated trust.

† One rule binds all of these: a record that has not been retrieved is not thereby false. Unretrieved is a state of the reader's context, not of the world, and a claim moves from unknown to false only on evidence.

Continuity and civilisation

Humanesque (specification; the merged release on the forge). *In the architecture*: the core over which everything above is a typed hypergraph, and the constitutional realm derived from PoIE through recursive sovereignty. *Nearest analogue*: an AGI architecture. *Difference*: a constitutional and epistemic realm as well as a technical substrate, not a product.

Kaivalyik, Zistgah and Cosmopolis (projections). As in Chapter 6: parallel cultural projections of the whole, not stages. Kaivalyik carries Synthematic Pragmatic Realism, SPR, the framework that heads Act III ("Kaivalyik: Towards a Synthematic Pragmatic Realism for the AGI Singularity"); Zistgah carries the habitats and the elements.

TransEg (released, doi:10.5281/zenodo.21321558). *In the architecture*: the transferred alter ego, the mechanism by which a constituted identity continues; a local digital-twin reference implementation; yadein's staggered upload feeds it. *Nearest analogue*: a digital avatar. *Difference*: continuity of an identity under the constitutional corpus, not a persona.

PAT.AL, VIDYA, GramSheel, Project VIKRAM and TWISHA (established). PAT.AL, the Participatory Alliance for Technology, Access and Livelihoods, is participatory infrastructure; VIDYA bridges to AyeAM; Project VIKRAM, Virtualized Infrastructure for Knowledge-driven Rural Ascension Management, runs under the GramSheel Foundation beside TWISHA. They sit on the dividers between the Acts. The question the author puts to them is whether technology can distribute capability without concentrating power:

The reference lab (planned; its launch scripts are the next release). An AGI-capable laboratory for 30 to 40 thousand, able to carry 80 to 90 per cent of the research behind high-end AI, robotics and automation papers, repeated as a pattern for every domain and aligned with GATE-level courses in computer science and in robotics and automation.

Also in the working notes

Modular self-reconfiguring robots †: very simple modular blocks with face-to-face connections and electro-permanent-magnet contacts, assembling into larger structures, ruggedised for decentralised manufacture with low-cost embedded control; PANINIphy's modular-block examples are the compiler side of this. **Further realization families** †: chemical, photonic, and a computational-biological intermediate layer of the AlphaFold class, built as an independent spine rather than a dependency. **The build engine** †: the direction of the whole is a self-reflecting, self-evolving and self-sustaining build engine in which PANINI captures and expresses the construct and the intent, and agency, human or machine, biological or physical, does the building. **INDUS and KHĀK** † are named in the notes among the lineages and the robots, alongside ANGEL, TARA and NI2A2.

9. The three Acts

In plain words. Three stages of one long story: machines that think with us; machines that help us build with living, synthetic and natural materials; and people and their worlds carrying on into the future.

The author's note "3 Act ASI ^ Panini" and his dictation of 27 September 2026 place the components in time. An Act says when a piece becomes most consequential, not when work on it began: much of Act III is being worked on now. The Acts are a temporal coordinate; the projections of Chapter 6 are cultural ones; neither is a ladder of maturity.




Figure 7. The three Acts, with their dividers; PANINI runs through all three.

Act	What it holds
Act I, AGI complete	Sanskrit and Pāṇini, with cyclers and paradigms; the front end, ILM and tajziya; the cyclers named by output with genie, which in market terms are agents and harnesses; the process cyclers, designed and run in AAB
Divider between I and II	PAT.AL, PRATIK, VIDYA
Act II, post-AGI to ASI	The cognitive enablement modules, CEMb and CEMs, plus Eco-CEM, with FAKIR; the accelerators
Divider between II and III	GramSheel and Project VIKRAM
Act III, post-ASI	All the BCI work; Synthemantic Pragmatic Realism, then TransEg, the transferred alter ego; the habitats: Humanesque, with Zistgah on Earth (Asli Zistgah), the Moon (Zistgah-e-Mahtab) and Mars (Zistgah-e-Bahram), Kaivalyik, Cosmopolis, and more left open
Across all three	PANINI
Beyond the Acts	The reference lab and the aligned courses

Act I: intelligence that works through language

The machinery of Act I is the front end (ILM and tajziya), the cyclers with genie, and AAB, with PANINI as the notation of the cycles. † In the author's notes, Act I is complete when intelligence can take an intent through language, in any language and script, into executable specifications, run it through cycles a person inspects, research and create with it, and keep the provenance of what it makes. On the first divider stand PAT.AL, the participatory infrastructure, PRATIK, the kernels, and VIDYA, the bridge to AyeAM.

Act II: from agency to realization

† In the author's notes, Act II is where the architecture moves from cognitive and digital agency to cognitive enablement and realization on substrates. The cognitive enablement modules carry one kernel to biological, synthetic and ecological substrates; FAKIR supplies the knowledge they act on; the accelerators carry the load. The notes tie this Act to realization on substrates, which PANINI's realization arms carry out: PANINIb for biology and PANINIphy for physical systems, each lowering intent through the common spine of Chapter 6 and each keeping desired, predicted, observed and validated results apart. † In the author's notes the same spine runs intent, specification, knowledge

resolution, construct, realization selection, planning, geometry and trajectory, execution, observation, evidence, verification and validation, provenance and a knowledge update, and the states of a result are kept apart all the way: desired, designed, predicted, planned, realized, executed, observed, verified, validated. On the second divider stand GramSheel and Project VIKRAM: the foundation with its principles, and the rural digitalisation platform inspired by them.

Act III: continuity

Act III gathers the brain-computer interface work (Quantum Neuromorphic BCI, brain-robot interfaces and NeuroMusical Therapeutics as active research; brain-to-text, cognitive prosthetics and intelligent rehabilitation as research directions needing clinical partners), Synthematic Pragmatic Realism, and TransEg, through which a constituted identity continues beyond the substrate that first carried it. Its habitats are Humanesque itself and its projections, with Zistgah's domes on Earth, the Moon and Mars; the Moon carries a lander embodiment and Mars a quadrotor one, and both replicate the Earth dome.

PANINI runs through all three Acts.

10. Three families, set apart: the AyeAI Triad, GramSheel and Project VIKRAM, Humanesque and its projections

In plain words. Three groups explained one at a time: the AyeAI Triad of thinking, body and communication; GramSheel and Project VIKRAM for villages; and Humanesque with the names it takes in each culture.

Three families of work are easy to blur into each other or into the rest of the estate. This chapter sets each apart, with what its own record says.

The AyeAI Triad

The Triad has three members, AyeAI®, AyeAM® and AyeCNSe®, with AyeAI at the apex and AyeCNSe and AyeAM at the base. It is a triad, not a pipeline and not a loop: a loop running from the environment through AyeCNSe, AyeAI and AyeAM back to the environment was proposed elsewhere and rejected. Its formal closure gives each member a tuple:




Figure 8. The AyeAI Triad: a triad with AyeAI at the apex, not a loop and not a pipeline.

$$\text{AyeAM} = \langle S, R, C \rangle, \quad \text{AyeAI} = \langle M, I, G \rangle, \quad \text{AyeCNSe} = \langle T, Ch, \Sigma \rangle.$$

AyeAI is cognition, the apex: “the inclusive AI”, in the words of its site, and the organisation [ayelai](#) holds its platforms, among them an open public stack for synthetic intelligence (opssi), virtual deployment infrastructure for scientific and cognitive computing (ayevdi), Staggered Upload™ (upload), the AyeAI Singularity Public License (spl) and the NI2A2 work.

AyeAM is embodiment: the body through which intelligence acts, across ground, air, water, space and habitat. VIDYA bridges to AyeAM without being a member of the Triad.

AyeCNSe is coordination and communication across media: the form studio at [forms.ayecnse.site](#) and the Human Context Model paper at [hcm.ayecnse.site](#) are its published faces.

Two repositories mark themselves as PEDLER and AyeAI Triad work: PRATIK’s kernel and Zamin’s ternary silicon.

GramSheel and Project VIKRAM

The GramSheel Foundation is the social architecture of the estate and the older of its lines: † it began as the Nature Welfare Society in 1993. Its principles, HEJLP, are health, education, justice, livelihood and peace. Its Village Knowledge Center (Universal Device) was a Stockholm Challenge finalist in 2010, and TWISHA runs under the same foundation.




Figure 9. GramSheel and Project VIKRAM, with VIKRAM's sections as its own pages name them.

Project VIKRAM, Virtualized Infrastructure for Knowledge-driven Rural Ascension Management, is † the technical implementation of GramSheel's principles. It was launched on 3 May 2020 as the Project Vikram Journal, “an inclusive rural digitalization platform inspired by the principles of GramSheel”, under the motto गाँव बढ़ेंगे तो सब बढ़ेंगे, when the villages grow, everyone grows. Its own pages (pvjournal.github.io) name its components:

Component	What its page covers
Education	learning and schooling
Health	healthcare, open to submissions as issues, with tools for an accessible virtual hospital
Livelihood	the ISIC domains of economic activity
Justice	social justice
Environment and Habitat	the living environment
Food Processing and Agriculture	agriculture and food
Peace	harmony through social dialogue
Industries, Technology and Space Exploration	the technical frontiers it reaches towards
Socialising, Projects and Resources	the community and its shared material

With Project Hindawi, its page says, VIKRAM works towards language neutrality across technical fields. In the three Acts, GramSheel and Project VIKRAM stand on the divider between Acts II and III.

Gram Swaraj and the questions behind GramSheel and VIKRAM

Gandhi's *Hind Swaraj* (1909) argued that self-rule is more than a transfer of authority, and his *Constructive Programme: Its Meaning and Place* (1941) placed the rebuilding of daily life inside the work of independence: its eighteen items include village industries, village sanitation, basic and adult education, the provincial languages and a national language, economic equality, the kisans, labour, the adivasis and women. The estate does not present itself as a restatement of these ideas. The point of contact is a question that becomes unavoidable once computation is infrastructure: can a community govern itself if the knowledge and the machines its decisions depend on are controlled elsewhere?

On Gandhi Jayanti, 2 October 2026, the author put the questions behind GramSheel, Project VIKRAM and Humanesque in his own words. Among them:

- Can rural communities become producers of knowledge and technology, not merely consumers of systems designed elsewhere?
- Can technology distribute capability without concentrating power?
- Can artificial intelligence strengthen individual and community sovereignty rather than replace it?
- Can knowledge remain plural, multilingual and culturally grounded while becoming computationally executable?
- Can verification, provenance and human agency become constitutional properties of intelligent systems?
- Can GramSheel build the local knowledge infrastructure, can Project VIKRAM turn that infrastructure into practical capability, and can Humanesque provide a larger architecture in which intelligence, identity, agency, provenance and equity remain connected?

His answer to the old division of the world is a direction rather than a side: to march together, not as a Global South or a Global North, but towards a Global Centre, where capability is distributed and humanity remains sovereign.

Humanesque and its projections, with their names

Humanesque is the core, with its site at humanesque.site and its organisation at hmnsq: the shared executable typed hypergraph, the ontology federation, the common execution and evidence spine, and the components every projection reuses. Its constitutional ground is PoIE, from which sovereignty recurses to every level. Its core stem repositories all belong to Humanesque; each projection gives them its own culturally sensitive names and sites.

Kaivalyik, the Indic projection, including the Advaita framing (kaivalyik.org, organisation [kaivalyik](http://kaivalyik.org)): consciousness, intelligence and transcendence; it carries Synthemetic Pragmatic Realism, Kaivalyik AGI and Jyotish. **Zistgah**, [زستگاه](http://zistgah.org), the Persian and Islamicate projection (zistgah.org, organisation [zistgah](http://zistgah.org)): life, planet and ecology, a habitat to live and work in; it carries the elements, the dome and the habitats on Earth (Asli Zistgah), the Moon (Zistgah-e-Mahtab) and Mars (Zistgah-e-Bahram, † conceived earlier as Mangal Kamna). **Cosmopolis**, the Greek, Western and classical projection (cosmopolis.org, organisation [c-polis](http://cosmopolis.org)): societies, civilisations, space and deep time.

The stem names below are the projection-neutral names of the components. Today the working repositories live in the [zistgah](http://zistgah.org) organisation, so most of them carry their Zistgah names; two stems already show a second projection, HPS in Hindawi and CHAKRA in Jyotish.

Stem, in Humanesque	Zistgah	Kaivalyik
The programming system	Urdu edition (urdu-ilm)	Hindawi
The observatory and its calendars	Chakra, the turning sky	Jyotish
The ground: balanced-ternary silicon	Zamin, زمین	
The water: the factory process	AAB, آب	
The air: the environmental replica	Fiza, فضا	
The shared virtual dome	the dome	
The Cognitive Workbench	Mez, میز	
The press	Matba, مطبع	
The cutting room, the visual cyler	Khwab, خواب	
The listening room, the audio cyler	Awaz, آواز	
The immersive cyler	Tilasm, طلسم	
The embodied cyler	Pench	
The record	Yadein, یادیں	
The book template	Kitab, کتاب	
The domain spine	Dhancha, ڈھانچہ	
Communication engineering	Ertabat, ارتباط	
Device control	Zasab, اعصاب	
The door for any AI	Darwaza, دروازہ	
Scan and store	Tasvir, تصویر	
The plates	Lawh, لوح	
The parser federation	Tajziya, تجزیہ	
The guide	Rahnuma, رہنما	
The Geometry of Becoming	Tabdili, تبدیلی	
The Geometry of Living	Zindagi, زندگی	
The Geometry of Persisting	Qaem, قائم	

Stem, in Humanesque	Zistgah	Kaivalyik
Topological Bildung	Tarbiyat, تربیت	

Some stems keep one name in every projection: Humanesque, ILM, Romenagri, PANINI with PANINIq, PANINIb and PANINIphy, PEDLER, QEDLER, PRATIK, CEM, UKOP, FAKIR, GENIE, TransEg, PoIE, VGC, COPA, and the provenance systems Tok DOI, spiguard, Candor and Misty DOI. Kaivalyik's and Cosmopolis's own names for the rest are given with their organisations, as each projection takes its repositories.

11. Projections: names, and raising a component into Humanesque

In plain words. The names each culture gives the shared parts, with new names proposed for review, and the steps for raising a part into the common core.

Humanesque holds the stems; each projection gives them names, sites and branding fit for its culture, and changes nothing in what they do. Three projections are established: Kaivalyik, Zistgah and Cosmopolis. Zistgah's names are in use. The names below for Kaivalyik and Cosmopolis, and for two further projections, one for the Chinese, Japanese, Korean and Vietnamese world and one for Polynesia, are proposals for the author's ruling. For the two new projections they also need consultation with speakers and communities of those cultures; sacred names are deliberately left out.

Stem, in Humanesque	Zistgah	Kaivalyik, proposed	Cosmopolis, proposed	Datong, proposed	Vaka, proposed
The ground: balanced-ternary silicon	Zamin زمین	Bhūmi भूमि	Gaia Γαῖα	地 dì, chí, jì, jī, địa	Whenua, Honua
The water: the factory process	AAB آب	Jala जल	Hydōr ὕδωρ	水 shuǐ, suǐ, su, 수, thủy	Wai
The air: the environmental replica	Fiza فضا	Vāyu वायु	Aēr αἴρ	風 fēng, fū, pung 풍, phong	Hau, Makani
The sky: the observatory	Chakra	Jyotish ज्योतिष (established)	Ouranos Οὐρανός	天 tiān, ten, cheon 천, thiên	Rangi, Lani
The calendars of the sky	Taqvīm تقویم	Pañcāṅga पञ्चाङ्ग	Hēmerologion ἡμερολόγιον	曆 lì, reki, ryeok 력, lịch	Maramataka
The shared virtual dome	Gumbad گنبد	Maṇḍapa मण्डप	Tholos θόλος	穹 qióng, kyū, gung 궁, khung	Whare, Hale
The desk, the Cognitive Workbench	Mez ميز	Piṭha पीठ	Trapeza τράπεζα	卓 zhuō, taku, tak 탁, trác	Papa
The press	Matba مطبع	Mudraṇa मुद्रण	Typographeion τυπογραφεῖον	印 yìn, in, in 인, ấn	Tā
The cutting room, the visual cycler	Khwab خواب	Svapna स्वप्न	Oneiros ὄνειρος	夢 mèng, mu, mong 몽, mông	Moemoeā
The listening room, the audio cycler	Awaz آواز	Svara स्वर	Phōnē φωνή	聲 shēng, sei, seong 성, thanh	Reo
The immersive cycler	Tilasm طلسم	Indrajāla इन्द्रजाल	Phantasia φαντασία	幻 huàn, gen, hwan 환, huyền	Whakaata
The embodied cycler	Pench	Yantra यन्त्र	Mēchanē μηχανή	機 jī, ki, gi 기, cơ	Mihini
The record	Yadein يادعين	Smṛti स्मृति	Mnēmē μνήμη	記 jì, kǐ, gǐ 기, ký	Mahara
Small work, cycled with any AI	Alam	Loka लोक	Kosmos κόσμος	界 jiè, kai, gye 계, giới	Ao
The cycles	Daur دور	Āvarta आवर्त	Kyklos κύκλος	環 huán, kan, hwan 환, hoàn	Porohita
The book template	Kitab کتاب	Grantha ग्रन्थ	Biblos βιβλος	書 shū, sho, seo 서, thư	Pukapuka
The guide	Rahnuma راهنما	Mārgadarśaka मार्गदर्शक	Periēgētēs περιηγητής	導 dǎo, dō, do 도, đạo	Kaiārahi
The domain spine	Dhancha دھانچہ	Merudaṇḍa मेरुदण्ड	Rhachis ράχis	骨 gǔ, kotsu, gol 골, cốt	Tuarā

Stem, in Humanesque	Zistgah	Kaivalyik, proposed	Cosmopolis, proposed	Datong, proposed	Vaka, proposed
Communication engineering	Ertabat ارتباط	Samvāda संवाद	Angelia ἀγγελία	信 xìn, shin, sin 신, tín	Kōrero
Device control	Zasab اعصاب	Nāḍī नाडी	Neuron νεῦρον	脈 mài, myaku, maek 맥, mäch	Uaua
The door for any AI	Darwaza دروازہ	Dvāra द्वार	Thyra θύρα	門 mén, mon, mun 문, mōn	Tatau
Scan and store	Tasvir تصویر	Citra चित्र	Eikōn εἰκών	像 xiàng, zō, sang 상, tuǝng	Whakaahua
The plates	Lawh لوح	Phalaka फलक	Pinax πίναξ	板 bǎn, ban, pan 판, bān	Pou
The parser federation	Tajziya تجزیہ	Vyākaraṇa व्याकरण	Analysis ἀνάλυσις	析 xī, seki, seok 석, tich	Wetewete
The Geometry of Becoming	Tabdili تبدیلی	Parivartana परिवर्तन	Genesis γένεσις	變 biàn, hen, byeon 변, biēn	Panoni
The Geometry of Living	Zindagi زندگی	Jīvana जीवन	Bios βίος	生 shēng, sei, saeng 생, sinh	Ora
The Geometry of Persisting	Qaem قائم	Sthiti स्थिति	Monē μονή	恒 héng, kō, hang 항, hāng	Mau
Topological Bildung	Tarbiyat تربیت	Śikṣā शिक्षा	Paideia παιδεία	育 yù, iku, yuk 육, dục	Ako
Kaivalyik AGI, the posters	Amal عمل	Karma कर्म	Ergon ἔργον	業 yè, gyō, eop 업, nghiệp	Mahi
The 28 nm tape-out catalogue	Tadbir28 تدبیر	Yukti युक्ति	Mētis μῆτις	巧 qiǎo, kō, gyo 교, xáo	Mōhio
The estate's register	Fehrist فهرست	Sūcī सूची	Katalogos κατάλογος	錄 lù, roku, rok 록, lục	Rārangi
The contracts	Qānūn قانون	Niyama नियम	Nomos νόμος	律 lǜ, ritsu, yul 율, luật	Ture
A research map	Zāicha زائچہ	Kuṇḍalī कुण्डली	Horoskopos ὠροσκόπος	譜 pǔ, fu, bo 보, phō	Mahere
The Paradox of Innovation	Tanāquz تناقض	Virodhābhāsa विरोधाभास	Paradoxon παράδοξον	矛盾 máodùn, mujun, mosun 모순, mâu thuẫn	to be named with speakers
Inclinations	Mayl ميل	Pravṛtti प्रवृत्ति	Rhopē ροπή	傾 qīng, kei, gyeong 경, khuynh	to be named with speakers
Health	Marham مرهم	Auśadha औषध	Pharmakon φάρμακον	藥 yào, yaku, yak 약, dược	Rongoā
Exercise	Varzish ورزش	Vyāyāma व्यायाम	Gymnasia γυμνασία	練 liàn, ren, yeon 련, luyēn	Korikori
Waking	Bedārī بیداری	Jāgrti जागृति	Egersis ἔγερσις	醒 xǐng, sei, seong 성, tỉnh	Oho
ILM's lexicon	ILM Lughat لغت	Śabdakośa शब्दकोश	Lexikon λεξικόν	辭 cí, ji, sa 사, từ	Papakupu
ILM's phonology	ILM Sautiyāt صوتيات	Dhvani ध्वनि	Phthongos φθόγγος	音 yīn, on, eum 음, âm	Oro
ILM's orthography	ILM Imlā املہ	Lipi लिपि	Graphē γραφή	字 zì, jì, ja 자, tự	Tuhi
The programming system	Urdu edition (urdu-ilm)	Hindawi (established)	HPS, Latin and Greek	HPS, Han characters	HPS, Te Reo Māori

Stem, in Humanesque	Zistgah	Kaivalyik, proposed	Cosmopolis, proposed	Datong, proposed	Vaka, proposed
humanesque	same	same	same	same	same
ilm	same	same	same	same	same
romenagri	same	same	same	same	same
panini	same	same	same	same	same
paniniq	same	same	same	same	same
paninib	same	same	same	same	same
paniniphy	same	same	same	same	same
pedler	same	same	same	same	same
qedler	same	same	same	same	same
pratik	same	same	same	same	same
fakir	same	same	same	same	same
ukop	same	same	same	same	same
genie	same	same	same	same	same
transeg	same	same	same	same	same
transeg-idgov	same	same	same	same	same
copa	same	same	same	same	same
vgc	same	same	same	same	same
candor	same	same	same	same	same
misty-doi	same	same	same	same	same
tok-doi	same	same	same	same	same
spiguard	same	same	same	same	same
dukedom	same	same	same	same	same
duke2	same	same	same	same	same
janapad	same	same	same	same	same
poie	same	same	same	same	same

The two proposed projections:

- **Datong**, 大同 (**Dàtóng**, **Daidō**, **Daedong** 대동, **Đai đồng**), for the Chinese, Japanese, Korean and Vietnamese world: the Great Unity, the ideal of a shared world in the Book of Rites; one written form read in all four languages. Alternatives: 和 (harmony: hé, wa, hwa 화, hòa). Care: the term also carries political uses in modern China; consultation should weigh that.
- **Vaka** (**vaka**, **waka**, **va'a**, **wa'a**), for the Polynesian world: the voyaging canoe that carried people across the Pacific, a habitat in motion. Alternatives: Moana (ocean), Fenua or Whenua (land). Care: Polynesian names should be confirmed with Māori, Hawaiian, Samoan, Tongan and Tahitian speakers; sacred names, such as Hawaiki, are deliberately not proposed.

Raising a component into Humanesque

A component usually begins life in one projection, as most have in Zistgah. Raising it into Humanesque, as proposed here for the author's ruling, takes six steps:

1. **Name the stem.** State what the component does in projection-neutral words, with its contract, its gate and its interfaces.
2. **Move the stem.** The repository becomes a Humanesque repository under its stem name, keeping its history and its DOI lineage.
3. **Project it.** Each projection carries the stem as its own branch or fork, with its own name, site and branding, as the author proposed: work and scripts stay atomic in the stem, and a projection never forks behaviour.
4. **Keep one contract.** Every projection runs the stem's gate unchanged; a projection may change words, names and design, never what the component computes.
5. **Keep one lineage.** A release of the stem is minted once, beneath its concept DOI; each projection's site links to that record rather than minting its own.
6. **Rule and review.** The author rules on every name; names in a new culture are confirmed with its speakers first; each name carries its review mark.

12. The repositories

In plain words. A catalogue of every public project folder, what each one does, and its permanent reference number where it has one.

This chapter was generated from what the forge itself reported on 29 September 2026: the GitHub REST API's listing of each organisation, and each repository's own `CITATION.cff` and README. The DOI column shows only what a repository records about itself. A repository with no DOI is public and pushed but not minted: a DOI is a dated public disclosure, and minting is decided repository by repository.

Organisation	Public repositories
zistgah	56
project-ilm	32
hindawiai	22
pvjournal	4
ayei	39
obonac	41
Total	194

The tables below place 113 of these repositories by what they do; 23 of them record a DOI of their own. The organisations hmnsq and kaivalyik hold no public repositories yet. Forks of third-party projects kept for toolchains and study, by organisation: ayeai: adversarial-robustness-toolbox, AIX360, athena, ayeam, bash2python, ChainOfSurvival.github.io, dockerfiles, dockersh, FHIR, go-digest, image-tools, insilico, KodExplorer, kubernetes-cobol, kui, qiskit, qiskit-aer, qiskit-terra, runc, runtime-tools, shellinabox, sway, util-linux, Virtual-Studio, x11docker, xmp; hindawiai: beautiful-racket, em-dosbox, JPC, js-dos, js-py, pyaara, qb.js, qbjc, racketscript, v86; obonac: apothecurry.github.io, autogit, avatarify, ayevdi, BLAS-Tester, BrainModel, Covid-FLutter, digihealthindia.github.io, docker-ubuntu-vnc-desktop, fft, foundry, gpconnect, HIT-Tools.github.io, ilm.codes, ishsi25, legacy-ilm, minimal-mistakes, misty-doi, NeuroMusical.github.io, ops, privacy-policy-generator, runtime-tools, services-1, spi-scan, tb, test-push, tui.editor, Twitter-API-v2-sample-code, universal_isic, urdu; project-ilm: praat, scripts.

Language, scripts and compilation

Repository	What it is	DOI
zistgah/panini, page	A language for prompt cycles	22103243
zistgah/panini_by_claude	PANINI: parser and interpreter	none recorded
zistgah/panini_by_grok, page	PANINI	none recorded
zistgah/panini_by_grok_toolchain	No description recorded on the forge.	none recorded
zistgah/panini_by_grok_toolchain_cgpt_fix	No description recorded on the forge.	none recorded
zistgah/humanesque, page	Humanesque merged release: PANINI, ILM, Hindawi, cyclers, Mez	none recorded
zistgah/paninib	PANINI Genetic Compiler	none recorded
zistgah/paniniphy	PANINI Genetic Compiler	none recorded
zistgah/tajziya, page	A parser federation for the classical languages of the human corpus, organised by family	22982238
zistgah/urdu-ilm, page	پروجیکٹ الیم / Project ILM: Urdu	none recorded
project-ilm/ilm.codes, page	ILM / Project ILM: Integrative Linguistic Multiscript (BETA)	none recorded

Repository	What it is	DOI
project-ilm/romenagri, page	Reversible transliteration library (Perso-Arabic seed). HPS/ILM. GPL. © 1993-2026 Abhishek Choudhary.	none recorded
project-ilm/ilm-phonology	Layer 0: Phonological base - IPA core, phonotactics, prosody, dialect normalization, audio alignment. Reversible sound mapping foundation.	none recorded
project-ilm/ilm-transliteration	Layer 1: Transliteration pivot - reversible, compiler-friendly Romanized layer with diacritics; bridges phonology and scripts.	none recorded
project-ilm/ilm-orthography	Layer 2: Orthographic engine - script-specific shaping, Unicode normalization, ZWJ/ZWNJ handling, grapheme clustering.	none recorded
project-ilm/ilm-lexicon	Layer 3: Morpho-lexical layer - multilingual lexicons, WordNets, morphology analyzers, clitics, affixation, reduplication.	none recorded
project-ilm/ilm-syntax-semantics	Layer 4: Syntax-semantics - UD parsing, SRL, AMR/UCCA, discourse segmentation, code-switching awareness.	none recorded
project-ilm/ilm-interface	Layer 5: Interface & application - multiscrypt IMEs, NLP/NLU APIs (TTS, ASR, QA), mother-tongue IDEs, educational tools.	none recorded
project-ilm/ilm-data	Data: canonical corpora, dialect maps, phoneme-audio pairs, aligned scripts, acoustic training data, lexical resources.	none recorded
project-ilm/ilm-devtools	Devtools: reversible compiler infra, script-aware IDEs, debug harnesses. Consumers of core ILM layers, not core layers themselves.	none recorded
project-ilm/ilm-validation	Validation: native speaker QA, annotation tools, field instruments, crowdsourced schema, cultural-linguistic QA loops.	none recorded
project-ilm/ilm-meta	Meta repo for ILM / علم: vision docs, architecture diagrams, manifestos, research summaries, roadmap, cultural-philosophical grounding.	none recorded
project-ilm/ilm-site, page	A layered, open-source effort for phonetically accurate, reversible, culturally respectful language tooling across world scripts.	none recorded
project-ilm/ilm-lsp	Language server for ILM/Hindawi	none recorded
project-ilm/language-specs	ILM/Hindawi language specifications	none recorded
project-ilm/vscode-ilm	Highlighting for localized keywords (.uhin)	none recorded
project-ilm/linguistics-labs	Dockerised labs for collaborators	none recorded
project-ilm/legacy	Initial fork of Project ILM from Project Hindawi (a fork)	none recorded
hindawiai/hindawi2020	Hindawi Programming System (Hindawi@Linux) 2008 version (a fork)	none recorded
hindawiai/hindawi2021	Bootstrapping Hindawi Programming System	none recorded
hindawiai/hindawi-legacy	Legacy files of Hindawi Programming System	none recorded
hindawiai/hindawi-tx	No description recorded on the forge.	none recorded
hindawiai/chintamani	HindawiAI in Telugu	none recorded
hindawiai/2023.12.21	Hindawi AI - Bada Din release (Coding in all ISO supported languages by 25th Dec 2023)	none recorded
hindawiai/hi, page	No description recorded on the forge.	none recorded
hindawiai/hindawiai.github.io, page	Hindawi Programming System - Programming in every person's own language (a fork)	none recorded
hindawiai/hinlin	Linux kernel source tree (a fork)	none recorded

Quantum, cognition and hardware

Repository	What it is	DOI
zistgah/paniniq, page	Panini Q: an oscillator substrate, the PEDLER engine and quantum validation, as a realization backend of PANINI	23020648
project-ilm/qedler, page	A Deterministic Event-Hypergraph Research Programme for Quantum Mechanics, Gravity, Gauge Theory and Cosmology	none recorded
zistgah/pratik_core_mvp, page	PRATIK Kernel Core: balanced-ternary, event-driven developmental kernel (PEDLER/AyeAI Triad)	21288232
zistgah/zamin, page	PRATIK ternary silicon: physicalizing the poised-zero substrate (PEDLER/AyeAI Triad)	21297556
zistgah/jugaad28, page	Integrated circuits taped out on 28 nm: evidence-gated descriptors and a low-cost open laboratory	22818408
zistgah/inclinations, page	No description recorded on the forge.	none recorded
project-ilm/cognitive-fabric	Cognitive Fabric	none recorded
ayeai/ni2a2	No description recorded on the forge.	none recorded
ayeai/opssi	Open public stack for synthetic intelligence	none recorded

Knowledge kernel, domains and instruments

Repository	What it is	DOI
zistgah/fakir, page	Foundational Architecture for Knowledge, Intelligence & Reasoning (UKOP kernel)	21436550 , 21436552
zistgah/dhancha, page	The domain-invariant skeleton every FAKIR domain is built on: ten invariants, a machine-readable domain descriptor, a validator, a cross-domain vocabulary leak check, a correlation graph, and PANINI cyclus configuration so the work can be done by any model.	22821645
zistgah/ertabat, page	A FAKIR domain for communication engineering: one capability surface for a link whose geometry is changing, nine conformance conditions with units (two honestly unestablished), a closed four-point integration boundary, and twenty-eight work packages with executable acceptance criteria.	22821651
zistgah/transeg	Project TransEg: local digital twin reference implementation (identity continuity research platform)	21321558
zistgah/transeg-idgov	TransEg identity governance: typed layers, envelope, staged policies, audit	none recorded
zistgah/transeg-research	TransEg research: 9-benchmark framework + 7 reference experiments	none recorded
zistgah/dome, page	Zistgah/dome: the Zistgah virtual dome as a reusable app/lib	21449034
project-ilm/chakra, page	Temporal Cycle Observatory: offline multi-tradition astronomical/calendrical observatory, tested library, and URL → JSON API. GPL-3.0.	none recorded
zistgah/jyotish, page	A Sanatan panchang and jyotisha, computed and explained; the Kaivalyik projection of CHAKRA, offline	none recorded
project-ilm/research-kundali	Research Kundali · ریسرچ کنڈلی	none recorded
zistgah/janapad, page	Humanesque Mahajanapad	none recorded
zistgah/mrd, page	Market Requirements for the Zistgah / AyeAI estate	none recorded
project-ilm/foundry	Canonical engineering substrate for Project ILM.	none recorded

Cyclers, desk and publishing

Repository	What it is	DOI
zistgah/mathba, page	The press. Posters in, book out, sealed, pushed, minted.	21948734
zistgah/khwab, page	A creativity cycle for turning a reel into a cued, published book	22005507
zistgah/awaz, page	The listening room	22005575
zistgah/tilasm, page	The immersive cycler. AR, XR, VR.	22010869
zistgah/pench, page	The embodied cycler. Robotics, cyberphysical, sim2real, real2sim.	22010962
zistgah/yadein, page	The record. A multimodal diary, staggered, toward TransEg.	22011073
zistgah/genie, page	Generalized Emotive-Narrative Interaction Engine	none recorded
zistgah/cycles, page	Cycles	none recorded
zistgah/alam, page	Small work, cycled with any AI. Part of the Zistgah ecosystem.	none recorded
zistgah/aab, page	The AI · Human Co-Development Studio	none recorded
zistgah/fiza, page	Verification-Gated Human-AI Co-Development for the Environment (Scaffold)	none recorded
zistgah/mez, page	The desk.  . Local-first, no account, no vendor.	none recorded
zistgah/kitab, page	A config-driven book template	none recorded
zistgah/genie-book, page	Kitab: a config-driven book template	none recorded
zistgah/varzish	Varzish	none recorded
zistgah/vakil	Litigation simulation	none recorded
zistgah/wake, page	Wake	none recorded
zistgah/marham	Marham	none recorded

Written corpora and frameworks

Repository	What it is	DOI
zistgah/tabdili, page	The Geometry of Becoming: Motion, Information, and Thermodynamics Across Scales	21764943
zistgah/zindagi, page	Geometry of Living (poster collection)	none recorded
zistgah/qaem, page	Geometry of Persisting (Beyond Life)	none recorded
zistgah/tarbiyat, page	Topological Bildung	none recorded
zistgah/dukedom, page	The Dukedom of Humanesque	none recorded
zistgah/duke2, page	The Dukedom of Humanesque, Volume II	none recorded
zistgah/paradox, page	The Paradox of Innovation	none recorded
zistgah/poie, page	Proclamation of Individual Equity	21397274
zistgah/copa, page	Cost of Perceived Authority: a black-box behavioural methodology (measurement framework + open hypothesis)	none recorded
project-ilm/vgc-notes	Notes on verification-gated human-AI co-development	none recorded
pvjournal/pilla, page	Parametric Information Layering Logic (Anatemnein). © 1993-2026 Abhishek Choudhary.	21764570
pvjournal/pajr, page	Index of PajR case reports · CC BY-NC-SA 4.0 · not medical advice	none recorded

Provenance, governance and operations

Repository	What it is	DOI
zistgah/governance	The ecosystem's centralized contracts, contexts & contribution gate	none recorded
zistgah/estate	IP estate: inventory, licensing, offerings, workspace sync	none recorded
zistgah/ztools	Zseed + mistyx	21312769
project-ilm/misty-doi, page	Automation-first DOI minting & research publication packaging. Muh Mitha Kijiye!	20719388
project-ilm/tok-doi, page	Tok DoI	21402745
project-ilm/spi-scan	Spi-scan	none recorded
project-ilm/ops	Project ILM operations	none recorded
project-ilm/ai-scratch, page	No description recorded on the forge.	none recorded
zistgah/zistgah.github.io, page	Zistgah	none recorded

Earlier platforms and field work

Repository	What it is	DOI
ayelai/ayelai.github.io	Aye AI ... the inclusive AI (www repo)	none recorded
ayelai/ayelai.ayeam.com	AyeAI ∴ AyeAM (a fork)	none recorded
ayelai/upload	Staggered Upload (TM)	none recorded
ayelai/ayevidi	Virtual Deployment IaaS (tools) for Scientific & Cognitive Computing	none recorded
ayelai/spl	AyeAI Singularity Public License	none recorded
ayelai/anusaaraka	Imported from https://web.archive.org/web/20200218013006/https://sourceforge.net/p/anusaaraka/git/ci/master/tree/	none recorded
ayelai/chuha	Chat Hosting Utility with Hypelink Automation	none recorded
ayelai/ayeq	No description recorded on the forge.	none recorded
pvjournal/pvjournal.github.io, page	GSF VIKRAM - An inclusive rural digitalization platform by GramSheel Foundation (a fork)	none recorded
pvjournal/journal	Journal	none recorded
obonac/obonac.github.io	No description recorded on the forge.	none recorded
obonac/research	No description recorded on the forge.	none recorded
obonac/ner_spec	No description recorded on the forge.	none recorded
obonac/sillyscope	Oscilloscope and spectrogram that reads from stdin	none recorded

Repository	What it is	DOI
obonac/webLinks	No description recorded on the forge.	none recorded

Records that stand on their own

Record	DOI
PEDLER, Point Event-Driven Learner (November 2001): the authoritative record	10.5281/zenodo.17497559
NI2A2, Natural-Interfaced Intelligent Adaptive Agents: the 2003 thesis built on PEDLER	10.5281/zenodo.17254607
VGC, Verification-Gated Human-AI Co-Development: the methodology paper	10.5281/zenodo.21264248
VGC-Health: the research proposal	10.5281/zenodo.21303401
CHAKRA: concept DOI of the observatory from which Jyotish is derived	10.5281/zenodo.21253628

Part III: Setting up

13. Choosing a platform

In plain words. Which computer set-up to use: Linux works best, Windows works through WSL, and a Mac works for most things.

Platform	What works	What is limited
Linux, native: Ubuntu 24.04 LTS or 26.04 LTS	Everything in the estate: shell scripts, compilers, Docker, GPU work, serial devices	Nothing; this is the reference platform
Windows 10 or 11 with WSL 2	Nearly everything, inside a real Ubuntu	USB devices need usbipd-win; the GPU is reached through the Windows driver; files under <code>/mnt/c</code> are slow
Windows, native	Reading, the browser-based tools, editing, Python, Node	The estate's bash scripts, Docker Engine, OpenTimestamps and systemd services do not run natively
macOS	Most command-line work, through Homebrew	No CUDA; BSD versions of the core tools unless GNU versions are installed

The estate's own workstation is a Ryzen 7 laptop with 16 GB of memory and an RTX 3050 with 4 GB, running Ubuntu with Docker, and TransEg names the same target. Any machine with 8 GB of memory and 30 GB of free disk runs every example in this guide.

Installing Ubuntu on a PC

1. Download the desktop image and the `SHA256SUMS` file from ubuntu.com/download/desktop into one folder
2. Check the download. The command must print `OK` against the image's name:

```
sha256sum -c SHA256SUMS --ignore-missing
```

3. Write the image to a USB stick of 8 GB or more. On Windows, use Rufus or balenaEtcher. On Linux, find the stick first; the commands below assume it is `sdb`, and they erase it:

```
lsblk -d -o NAME,SIZE,MODEL
ISO=$(ls ubuntu-*-desktop-amd64.iso)
sudo dd if="$ISO" of=/dev/sdb bs=4M status=progress oflag=sync
```

4. Boot from the stick (the boot-menu key is usually F12, F10, F2 or Esc), choose *Install Ubuntu*, and either erase the disk or install alongside Windows. Before installing alongside Windows, turn off Windows Fast Startup, and if BitLocker is on, save its recovery key or suspend it. Secure Boot can stay on; Ubuntu's boot loader is signed.
5. After the first login:

```
sudo apt update && sudo apt full-upgrade -y
```

Windows with WSL 2

In a PowerShell window run as Administrator:

```
wsl --install
```

Restart, then:

```
wsl --list --online
wsl --install -d Ubuntu-24.04
wsl --set-default-version 2
wsl --list --verbose
```

The last command must show `VERSION 2` beside Ubuntu. Open Ubuntu from the Start menu, create your Linux user and update it with `sudo apt update && sudo apt full-upgrade -y`. Then switch on systemd, which Docker Engine and several estate services expect:

```
printf '[boot]\nsystemd=true\n' | sudo tee /etc/wsl.conf
```

Run `wsl --shutdown` in PowerShell and open Ubuntu again; `systemctl is-system-running` now answers `running` or `degraded`.

Five rules that save hours:

- Keep your work in the Linux file system, under `~/work`, never under `/mnt/c`; file access across that boundary is many times slower. Windows reaches your Linux files at `\\wsl$`.
- Use VS Code on Windows with its WSL extension, and type `code` inside a Linux folder.
- For an NVIDIA GPU install only the Windows driver; do not install a Linux display driver inside WSL. `nvidia-smi` inside Ubuntu then lists the card.
- To hand a USB device such as an ESP32 board to WSL, install `usbipd-win` with `winget install usbipd`, then in an Administrator PowerShell run `usbipd list`, `usbipd bind --busid 1-4` and `usbipd attach --wsl --busid 1-4`, using the bus identifier that `usbipd list` printed for your device in place of `1-4`.
- To cap WSL's memory, create a file named `.wslconfig` in your Windows user folder containing the two lines `[wsl2]` and `memory=8GB`.

Windows, native

```
winget install --id Git.Git -e
winget install --id Python.Python.3.12 -e
winget install --id OpenJS.NodeJS.LTS -e
winget install --id Microsoft.VisualStudioCode -e
git config --global core.autocrlf input
```

The estate's scripts are bash scripts written for GNU tools. Git Bash runs simple ones, but not systemd services, Docker Engine or OpenTimestamps workflows. `flex` and `bison` come through MSYS2, with `pacman -S --needed base-devel flex bison mingw-w64-ucrt-x86_64-gcc`. Windows file systems ignore case, lock open files and end lines with CRLE, and each of these breaks scripts written for Linux; the `autocrlf input` setting above prevents the third. Native Windows is right for reading, for the browser-based tools (the PANINI studio, CHAKRA, Jyotish, the Zamin simulator) and for Python notebooks. Build and run everything else in WSL.

macOS

```
xcode-select --install
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
brew install git python@3.12 node flex bison gcc make
echo 'export PATH="$(brew --prefix bison)/bin:$(brew --prefix flex)/bin:$PATH"' >> ~/.zshrc
```

macOS ships a bison from 2006 and BSD versions of `sed`, `grep` and `date`. Homebrew installs `flex` and `bison` “keg-only”, that is, not on the PATH, and the last line puts them first. There is no CUDA on current Macs; Docker runs through Docker Desktop or Colima.

14. The baseline environment and a first run

In plain words. How to install the basic tools and run your first real program from this work.

Packages

On Ubuntu, native or in WSL:

```
sudo apt update
sudo apt install -y build-essential git curl wget unzip make cmake pkg-config \
    flex libfl-dev bison gawk gdb python3 python3-venv python3-pip pipx nodejs npm jq \
    fonts-noto-core fonts-indic
gcc --version | head -1; flex --version; bison --version | head -1
python3 --version; node --version
```

This brings GCC and binutils, the lexer and parser generators of Chapter 21 with flex's library and the gawk that Hindawi's driver needs, Python with virtual environments, Node for the JavaScript engines of PANINI, CHAKRA and Jyotish, and fonts for the Indic scripts.

Git and GitHub

Use the name and email of your own GitHub account in place of the example ones:

```
git config --global user.name "Ada Lovelace"
git config --global user.email "ada@example.org"
git config --global init.defaultBranch main
ssh-keygen -t ed25519 -C "ada@example.org"
cat ~/.ssh/id_ed25519.pub
```

Paste the printed key at GitHub under *Settings, SSH and GPG keys*, and check it with `ssh -T git@github.com`. The GitHub CLI turns forks and pull requests into single commands:

```
sudo apt install -y gh
gh auth login
```

Docker Engine

These are Docker's own commands for Ubuntu, from docs.docker.com/engine/install/ubuntu:

```
sudo apt install -y ca-certificates curl
sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc
sudo chmod a+r /etc/apt/keyrings/docker.asc
echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc] \
https://download.docker.com/linux/ubuntu $(. /etc/os-release && echo "$VERSION_CODENAME") stable" \
| sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt update
sudo apt install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin
sudo usermod -aG docker "$USER"
```

Log out and back in, then `docker run hello-world` must greet you.

A first run: PANINIq

PANINIq, the quantum-simulation backend of PANINI, is a good first repository: small, tested, and with its own gate. Its README gives the install; the lines below follow it, with the virtual environment kept outside the clone:

```
mkdir -p ~/work && cd ~/work
git clone https://github.com/zistgah/paniniq.git
python3 -m venv ~/work/venv-paniniq
. ~/work/venv-paniniq/bin/activate
cd paniniq
export PYTHONDONTWRITEBYTECODE=1
pip install -r requirements.txt qiskit
python3 -m unittest discover -s tests
python3 demo.py
bash ops/verify.sh
```

The environment variable and the outside virtual environment matter because the repository's gate counts any `__pycache__` folder or compiled `.pyc` file inside the clone as a build artefact (its clause P08). Python writes those caches when it runs the tests, and a virtual environment inside the clone holds thousands of them, so either would turn the gate red for a reason that has nothing to do with the code.

This is the run made while this guide was built, on the repository at commit 82a5cf2:

Recorded run: `data/logs/paniniq-run-2026-09-29.txt`

```
$ git log -1 --format='%h %ad' --date=short
82a5cf2 2026-09-28

$ export PYTHONDONTWRITEBYTECODE=1

$ python3 -c 'import numpy, scipy, qiskit; print(numpy.__version__, scipy.__version__, qiskit.__version__)'
2.5.3 1.18.1 2.5.2

$ python3 -m unittest discover -s tests
-----
Ran 23 tests in 0.217s

OK

$ bash ops/verify.sh
Panini Q: contract verification
PASS    P01  every source carries the copyright and SPDX lines
PASS    P02  no affiliation other than AyeAI is claimed
PASS    P03  the test suite passes
PASS    P04  the suite catches the anti-ferromagnetic zeroing bug
PASS    P05  demo.py runs to completion
PASS    P06  LICENSE is the FSF's verbatim GPL-3.0 text
PASS    P07  every row the README marks Tested names its test
PASS    P08  no build artefacts
PASS    P09  no path outside the folder is named
PASS    P10  misty.json and CITATION.cff agree; no placeholder DOI
10 passed, 0 failed, 0 unjudged
```

And the demonstration, as it printed:

Recorded run: `data/logs/paniniq-demo-2026-09-29.txt`

```
$ python3 demo.py

=====
1. Point events -> inclination field -> relaxation -> spawn
=====
event 0: point=[0.85, 0.12, 0.44]
        order parameter R = 0.3874   |S| = 5   -> SPAWNED a node
event 1: point=[0.1, 0.9, 0.3]
        order parameter R = 0.5859   |S| = 5
event 2: point=[0.83, 0.14, 0.46]
```

```

order parameter R = 0.4573  |S| = 6  -> SPAWNED a node

=====
2. Lifecycle maintenance (decay priors, then prune/consolidate)
=====
|S| before maintenance: 8
pruned=1, merged=0
|S| after maintenance: 7

=====
3. Compile to Qiskit + statevector validation
=====

q_0: {0
q_1: {1 CTQW(K,t)
q_2: {2

meas: 3/
0 1 2

statevector dim = 8, sums to 1.000000 (unitarity check)

=====
4. Max-Cut benchmark: oscillator vs. brute force vs. QAOA
=====
oscillator  cut_value=4.0  time=2.90 ms
bruteforce  cut_value=4.0  time=0.22 ms
qaoa        cut_value=4.0  time=52.09 ms

=====
5. Hardware bridge (mocked -- no real serial device exists here)
=====
21 frame(s) 'sent' to MOCK
sample frame (hex): aa55000100b3
This only proves the K-matrix -> byte-frame packing is correct. No physical oscillator was touched -- there's no serial port
or USB device available in this sandbox.

=====
6. ClassicalSubstrate (PyTorch) scaling demo
=====
SKIPPED: torch is not installed in this sandbox (no network access to install it here).
pedler/substrate.py::ClassicalSubstrate is written and ready -- `pip install torch` and re-run this script; it will
automatically use a GPU if `torch.cuda.is_available()`.

```

Chapter 23 explains every number in it: the order parameter R , the spawned states, the continuous-time quantum walk on three qubits, and why the oscillator, brute force and QAOA agree on a cut of 4.

Every other component

Chapter 15 gives the commands for every component of the estate: the ones executed for this guide, and the rest as their own READMEs state them.

15. Running every component

In plain words. Step-by-step commands to start each part on your own computer, with what you should see when it works and what to do when it does not.

Every component below was cloned and run on a clean machine on 3 October 2026, with the commands its own README gives, and each card says exactly what happened. Where a step needed something the README does not mention, the card adds it; where a component did not pass, the card says why. Set up the machine first (Chapters 13 and 14).

Each card has the same shape: what the component is, what it needs, the commands to type in order, what you should see, and its status. Every card starts with a clone of the repository into a folder of the same name; the commands that follow are typed inside that folder:

Machine: Ubuntu 24.04, Python 3.12, Node 22, GCC 13. *Run on:* 3 October 2026.

Language and identity

Hindawi and Romenagri

What it is. The Hindawi Programming System's C front end, its driver hinc and the Romenagri script layer. *Needs:* gcc, make, flex, libfl-dev, gawk.

```
git clone https://github.com/hindawiai/chintamani.git && cd chintamani
mkdir -p .root/usr/bin .root/usr/lib .root/usr/include
make -C Romenagri all && make -C Romenagri install INSTRROOT="$PWD/.root"
export PATH="$PWD/.root/usr/bin:$PATH"
make -C Hindawi/guru all && make -C Hindawi/guru install INSTRROOT="$PWD/.root"
make -C Hindawi/hindr install INSTRROOT="$PWD/.root"
cp Hindawi/samples/HindiC.uhin . && hinc HindiC.uhin
printf 'रमनगरम\n' | ./hin.exe
```

You should see: संकलन के परिणाम, then the program asks आपका नाम क्या है? and counts to ten. *Status:* Executed for this guide. The compilers chapter prints the run.

The Urdu edition

What it is. Urdu shailis and the driver urducc, which sends an Urdu-script program through the Devanagari hub and Romenagri. *Needs:* the Hindawi build above, on the PATH.

```
git clone https://github.com/zistgah/urdu-ilm.git && cd urdu-ilm
mkdir -p .root/usr/bin .root/usr/lib .root/usr/include
make -C Romenagri all && make -C Romenagri install INSTRROOT="$PWD/.root"
export PATH="$PWD/.root/usr/bin:$PATH:<chintamani clone>/.root/usr/bin"
make -C ILM/guru all
bash ILM/urducc -s ILM/samples/UrduC.uhin
bash ILM/urducc ILM/samples/UrduC.uhin && ./hin.exe
```

You should see: C with the names na and ka, then the prompt in the hub script. *Status:* Executed for this guide. Replace with the folder of the Hindawi clone; the compilers chapter prints the run.

tajziya

What it is. The parser federation for the classical languages, by family. *Needs:* python3, make.

```
git clone https://github.com/zistgah/tajziya.git && cd tajziya
PYTHONPATH=src python3 -m tajziya commonalities
make check
```

You should see: make check: 19 passed, 0 failed, 1 unjudged. *Status:* Passed when run on 3 October 2026. Run the module with PYTHONPATH=src, from the clone's root.

The PANINI family

PANINI, the prompt-cycle language

What it is. The language the cyclers are written in, with its engine and studio. *Needs:* python3, node.

```
git clone https://github.com/zistgah/panini.git && cd panini
python3 panini.py check cyclers/pench.pni
python3 panini.py stages cyclers/yadein.pni
node tests/panini.test.mjs
python3 panini.py serve
```

You should see: 0 problems; 45 pass, 0 fail; the studio answers on 127.0.0.1:8717. *Status:* Passed when run on 3 October 2026.

PANINI, parser and interpreter

What it is. A general-purpose PANINI implementation with its conformance report. *Needs:* node.

```
git clone https://github.com/zistgah/panini_by_claude.git && cd panini_by_claude
node bin/panini.mjs conformance
npm test
```

You should see: 145 passed, 0 failed. *Status:* Passed when run on 3 October 2026.

PANINI, the stage-0 bootstrap

What it is. PANINI with a bootstrap that records its own evidence. *Needs:* node.

```
git clone https://github.com/zistgah/panini_by_grok.git && cd panini_by_grok
node src/cli.js run examples/hello.pni
node tests/run.mjs
node scripts/bootstrap.mjs
```

You should see: Hello, PANINI; 32 passed, 0 failed; stage-0 evidence written. *Status:* Passed when run on 3 October 2026. Its wider report, npm test, passes 81 of 83 here; the two others need tools this machine lacks.

The Humanesque merged release

What it is. PANINI, ILM, Hindawi, the cyclers and Mez in one tree. *Needs:* bash, node; optionally zig.

```
git clone https://github.com/zistgah/humanesque.git && cd humanesque
bash VERIFY.sh
npm test
```

You should see: 79 of 83 checks pass. *Status:* Partly passes, as run on 3 October 2026. The four that do not pass here need two downloadable model files (wllama.wasm and stories15M.Q4_0.gguf), the zig toolchain, and the Bengali flattening round trip.

PANINIq

What it is. The oscillator substrate, the PEDLER engine and quantum validation. *Needs:* python3, numpy, scipy, qiskit.

```
git clone https://github.com/zistgah/paniniq.git && cd paniniq
python3 -m venv ../venv-paniniq && . ../venv-paniniq/bin/activate
export PYTHONDONTWRITEBYTECODE=1
pip install -r requirements.txt qiskit
python3 -m unittest discover -s tests
python3 demo.py
bash ops/verify.sh
```

You should see: Ran 23 tests, OK; its gate 10 passed. *Status:* Executed for this guide. The first-run chapter prints the run.

PANINIB, biological realization

What it is. Biological intent lowered through typed IRs to a realization plan. *Needs:* python3.

```
git clone https://github.com/zistgah/paninib.git && cd paninib
python3 panini compile examples/replace.panini -o out
python3 panini inspect out --stage sequence-ir
python3 panini verify out
```

You should see: each step prints its JSON record. *Status:* Passed when run on 3 October 2026.

PANINIphy, physical realization

What it is. Physical intent compiled through a physical IR, with parts chosen by requirement. *Needs:* python3.

```
git clone https://github.com/zistgah/paniniphy.git && cd paniniphy
python3 panini compile examples/modular_block.panini -o out
python3 panini parts --category actuator --force 50N --stroke 20mm --voltage 12V
```

You should see: physical resolution=RESOLVED verify=True. *Status:* Passed when run on 3 October 2026.

Cyclers, desk and studio

matba, the press

What it is. Posters in, book out, sealed, pushed, minted. *Needs:* python3.

```
git clone https://github.com/zistgah/matba.git && cd matba
python3 matba.py serve
```

You should see: the press answers on 127.0.0.1:8710. *Status:* Passed when run on 3 October 2026.

khwab, the cutting room

What it is. The visual cyler. *Needs:* python3.

```
git clone https://github.com/zistgah/khwab.git && cd khwab
python3 khwab.py serve
```

You should see: the cutting room answers on 127.0.0.1:8711. *Status:* Passed when run on 3 October 2026.

awaz, the listening room

What it is. The audio cyler. *Needs:* python3.

```
git clone https://github.com/zistgah/awaz.git && cd awaz
python3 awaz.py serve
```

You should see: the listening room answers on 127.0.0.1:8712. *Status:* Passed when run on 3 October 2026.

tilasm, pench and yadein

What it is. The immersive, embodied and record cyclers; each repository has the same tests and its own page. *Needs:* node.

```
git clone https://github.com/zistgah/tilasm.git && cd tilasm
node tests/cycler.test.mjs
node tests/inbox.test.mjs
node tests/models.test.mjs
```

You should see: 20, 23 and 22 pass, 0 fail. *Status:* Partly passes, as run on 3 October 2026. Run the same three in zistgah/pench and zistgah/yadein. Two further test files, genie.test.mjs and leak.test.mjs, import modules the repositories do not carry.

GENIE

What it is. The Generalized Emotive-Narrative Interaction Engine. *Needs:* node.

```
git clone https://github.com/zistgah/genie.git && cd genie
node test/run.js
```

You should see: 191 passed, 0 failed. *Status:* Passed when run on 3 October 2026. Open index.html, or its page, for the engine itself.

alam

What it is. Small work, cycled with any AI. *Needs:* node, npm.

```
git clone https://github.com/zistgah/alam.git && cd alam
npm install jsdom
node test/alam-test.js
node test/dom-test.js
```

You should see: alam dom: 9 passed, 0 failed. *Status:* Passed when run on 3 October 2026. The page test needs jsdom, installed by the first line.

Mez, the Cognitive Workbench

What it is. The local-first desk. *Needs:* python3.

```
git clone https://github.com/zistgah/mez.git && cd mez
./mez doctor
./mez bearings
./mez serve
```

You should see: no account, no API key, no network required; the desk answers on 127.0.0.1:7373. *Status:* Passed when run on 3 October 2026.

Knowledge, domains and instruments

FAKIR

What it is. The knowledge kernel and its explorer. *Needs:* bash.

```
git clone https://github.com/zistgah/fakir.git && cd fakir
bash ops/verify.sh
```

You should see: CONTRACT SATISFIED, 6 checks. *Status:* Passed when run on 3 October 2026. The explorer is its page: zistgah.org/fakir

Dhancha

What it is. The domain spine every FAKIR domain is built on. *Needs:* bash, python3, make.

```
git clone https://github.com/zistgah/dhancha.git && cd dhancha
make check
```

You should see: site matches the registries; pr_gate bites on a claim and passes a clean change. *Status:* Passed when run on 3 October 2026.

ertabat

What it is. The communication-engineering domain. *Needs:* bash.

```
git clone https://github.com/zistgah/ertabat.git && cd ertabat
bash ops/verify.sh
```

You should see: CONTRACT OK. *Status:* Passed when run on 3 October 2026.

jugaad28

What it is. The 28 nm tape-out catalogue. *Needs:* bash.

```
git clone https://github.com/zistgah/jugaad28.git && cd jugaad28
bash ops/verify.sh
```

You should see: VERIFY GREEN. *Status:* Passed when run on 3 October 2026.

TransEg

What it is. The local digital-twin reference implementation. *Needs:* bash, python3; Docker for the stand-up.

```
git clone https://github.com/zistgah/transeg.git && cd transeg
./scripts/test.sh
./scripts/bootstrap.sh --dry-run
```

You should see: 10 passed; the stand-up plan. *Status:* Passed when run on 3 October 2026. The full stand-up, ./scripts/bootstrap.sh and docker compose, needs Docker and a GPU; it was not run here.

TransEg identity governance

What it is. Typed layers, envelope, staged policies and audit. *Needs:* python3, jsonschema, pytest.

```
git clone https://github.com/zistgah/transeg-idgov.git && cd transeg-idgov
pip install jsonschema pytest
python3 tools/idgov_validate.py policies
python3 -m pytest -q tests/
```

You should see: VALID: 8 stage policies; 14 passed. *Status:* Passed when run on 3 October 2026.

CHAKRA

What it is. The offline observatory and its calendars. *Needs:* node.

```
git clone https://github.com/project-ilm/chakra.git && cd chakra
node test/run.js
```

You should see: ALL SUITES PASSED. *Status:* Passed when run on 3 October 2026. Open index.html for the observatory.

Jyotish

What it is. The panchang, the Kaivalyik projection of CHAKRA. *Needs:* bash, node.

```
git clone https://github.com/zistgah/jyotish.git && cd jyotish
bash ops/verify.sh
```

You should see: PASSED, with 1 check skipped and named. *Status:* Passed when run on 3 October 2026. Open index.html for the panchang.

Cognition and hardware

PRATIK kernel

What it is. The balanced-ternary, event-driven kernel. *Needs:* g++; nvcc for CUDA.

```
git clone https://github.com/zistgah/pratik_core_mvp.git && cd pratik_core_mvp
bash build.sh
bash build.sh cuda
```

You should see: ALL CHECKS PASSED. *Status:* Passed when run on 3 October 2026. The CPU build passed here; the CUDA build needs nvcc and an NVIDIA GPU. The script is not marked executable, so call it with bash.

Zamin

What it is. Ternary silicon: SPICE cells, layout and a simulator. *Needs:* ngspice, python3.

```
git clone https://github.com/zistgah/zamin.git && cd zamin
cd spice && ngspice -b sign_frustration_tb.sp
```

You should see: a SPICE run of the ternary cells. *Status:* Does not run as is, as run on 3 October 2026. On ngspice 42 the testbench stops: the source v_gnd is reported as a shorted voltage source. The interactive simulator is its page.

QEDLER

What it is. The event-hypergraph physics programme. *Needs:* bash.

```
git clone https://github.com/project-ilm/qedler.git && cd qedler
bash ops/verify.sh
bash ops/resume.sh && cat RESUME.md
```

You should see: 5 of 6 clauses pass. *Status:* Partly passes, as run on 3 October 2026. Clause C2, manifest completeness, fails: 15 files are not in MANIFEST.sha256. Resealing the manifest clears it.

Provenance, tools and labs

Misty DOI

What it is. DOI minting and packaging for Zenodo. *Needs:* python3, pip.

```
git clone https://github.com/project-ilm/misty-doi.git && cd misty-doi
pip install misty-doi
misty --help
misty init
misty validate -m misty.json
```

You should see: usage: misty [-h] [--version] {init,validate,transform,package,ots,publish} . *Status:* Passed when run on 3 October 2026. Publishing needs a Zenodo token in ZENODO_TOKEN.

Research Kundali

What it is. A research map of a person, a lab or a paper. *Needs:* python3 and network access to the public records it reads.

```
git clone https://github.com/project-ilm/research-kundali.git && cd research-kundali
python3 kundali/kundali.py <ORCID, lab name or DOI> --out out
```

You should see: a research map in out/. *Status:* Not run here: needs network access. It reads public records such as ORCID; it could not reach them from this build machine, and stops rather than write an empty chart.

The linguistics lab

What it is. A starter lab in a container. *Needs:* Docker

```
git clone https://github.com/project-ilm/linguistics-labs.git && cd linguistics-labs
docker build -t ilm-lab .
docker run --rm -p 8888:8888 -v "$PWD":/lab ilm-lab
```

You should see: the lab on localhost:8888. *Status:* Not run here: needs Docker. Not run here: this build machine has no Docker

This guide

What it is. The site, the PDF and the Markdown, with their gate. *Needs:* python3, gcc, make, flex, libfl-dev, gawk, gdb; for a rebuild also pandoc, node, graphviz.

```
git clone https://github.com/zistgah/rahnuma.git && cd rahnuma
bash ops/verify.sh
bash examples/hindawi.sh words
python3 tools/build.py
```

You should see: 12 passed, 0 failed, 0 unjudged. *Status:* Passed when run on 3 October 2026.

Components you open in a browser

Component	Page
FAKIR's explorer	zistgah.org/fakir
The shared dome	zistgah.github.io/dome
AAB, the factory process	zistgah.org/aab
GENIE	zistgah.github.io/genie
tilasm, pench and yadein	zistgah.github.io/tilasm
Kitab	zistgah.github.io/kitab
Tok DOI	project-ilm.github.io/tok-doi
CHAKRA	project-ilm.github.io/chakra
Jyotish	zistgah.github.io/jyotish
ILM	ilm.codes
PANINI's studio	zistgah.github.io/panini
The estate	zistgah.org

16. Labs: containers now, launch scripts next

In plain words. How a container packs a whole laboratory into one box that runs the same on every computer, and what comes next.

What a container is

A container is an ordinary process given its own view of the file system, the process table and the network through Linux namespaces, and held to a share of processor and memory by control groups. It starts from an image built in layers, each layer a set of file changes identified by its hash. The same image gives the same environment on every laptop, which is exactly what a teaching or research lab needs: the lab stops depending on what happens to be installed on the machine in front of you.

Checking that your machine is ready

```
docker run --rm hello-world
docker compose version
```

For GPU work, install the NVIDIA Container Toolkit by NVIDIA's [installation guide](#), then:

```
sudo nvidia-ctk runtime configure --runtime=docker
sudo systemctl restart docker
sudo docker run --rm --runtime=nvidia --gpus all ubuntu nvidia-smi
```

The last command prints the same table as `nvidia-smi` on the host, from inside a container:

What exists, and what comes next

Three things run in containers or isolated environments today: project-ilm/linguistics-labs, a starter container for linguistics collaborators whose full specification the author is to provide; TransEg's compose stand-up, which brings up the local digital twin; and PANINIq, which needs only a Python virtual environment.

The next release of this package carries the lab launch scripts: one script per lab, the software dockerised and retrieved from the existing repositories rather than rebuilt. The labs follow the pattern of the reference lab: most of the research behind current AI, robotics and automation papers, at modest cost, aligned with courses at the level of GATE in computer science and in robotics and automation.

Part IV: Foundations

17. Numbers inside the machine

In plain words. How a computer stores numbers using only ones and zeros, including negative numbers and fractions.

A processor holds voltages that stand for bits. Everything else on this page, including the hexadecimal digits, is notation for people.

Positional notation

A number written with digits d_i in base b is

$$n = \sum_{i=0}^{k-1} d_i b^i, \quad 0 \leq d_i < b.$$

So $42 = 101010_2 = 2A_{16}$: the letters A to F are simply names for the values ten to fifteen, chosen so that one hexadecimal digit covers exactly four bits. The machine never sees the letter A, any more than it sees the letters of a keyword in a program.

Signed integers: two's complement

With n bits $b_{n-1} \dots b_0$, the value is

$$v = -b_{n-1} 2^{n-1} + \sum_{i=0}^{n-2} b_i 2^i, \quad -2^{n-1} \leq v \leq 2^{n-1} - 1.$$

Negative x is therefore stored as $2^n - |x|$. The payoff is that one adder serves signed and unsigned arithmetic alike, which is why every processor in use works this way.

Reals: IEEE 754

A double-precision number has a sign bit s , eleven exponent bits e and fifty-two fraction bits f , and for normal numbers

$$v = (-1)^s \times (1.f)_2 \times 2^{e-1023}.$$

One tenth has no finite binary expansion, so it is stored rounded, and sums of rounded values are rounded again.

Byte order and bit order

A 32-bit value occupies four bytes. A little-endian machine (x86, and ARM as usually configured) stores the least significant byte at the lowest address; a big-endian machine stores the most significant byte first, and network protocols use big-endian, which is why C has `htonl` and `ntohl`. That is endianness, and it concerns bytes.

Bit numbering is a separate convention: whether bit 0 is the least or the most significant bit. The hexadecimal digit E is **1110** when written most significant bit first and **0111** when written least significant bit first, and a serial line such as a UART sends the least significant bit first. Both conventions are representations. The circuit has only the wires.

Executed while this guide was built: `python3 examples/numbers_demo.py`

```
Integers and reals
0x0A0B0C0D little-endian bytes: 0d 0c 0b 0a    big-endian: 0a 0b 0c 0d
-5 in 8-bit two's complement: 11111011    (256 - 5 = 251)
```

```
0.1 as an IEEE 754 double: 3fb99999999999a = 0x1.999999999999ap-4  
0.1 + 0.2 == 0.3 is False; the sum is 0.3000000000000004
```

These are the facts the GATE syllabus asks for under number representation and computer arithmetic, fixed and floating point (Chapter 26).

18. Characters inside the machine

In plain words. How letters from any script become numbers inside a computer, and why a name must survive every tool a program passes through.

ASCII

Seven bits, 128 codes: 0 to 31 and 127 are control codes, 32 to 126 are printable, and **A** is 65, or 0x41. It was designed in the 1960s for English teleprinters, and nearly every later character code keeps it as its first 128 positions.

ISCII

ISCII, the Indian Script Code for Information Interchange (IS 13194:1991), is an 8-bit code. Codes 0 to 127 are ASCII. The upper half holds one layout shared by the Brahmi-derived scripts of India: the same code means KA in Devanagari, Bengali, Gurmukhi, Gujarati, Oriya, Tamil, Telugu, Kannada and Malayalam, and an attribute code selects the script for display. The shared layout records the common phonological organisation of these scripts. The price is that the script, and with it the language, has to be known from context.

Unicode

Unicode has room for 1,114,112 code points, U+0000 to U+10FFFF, in 17 planes of 65,536. Plane 0, the Basic Multilingual Plane, holds every major Indic block. Each script has its own block, so KA has a different code point in each. The blocks were laid out from ISCII's order; however, so KA sits at offset 0x15 and the virama at offset 0x4D in every one of the nine: the commonality ISCII carried in a single code, Unicode carries as a fixed offset.

UTF-8

A code point U is written in one to four bytes:

$$\begin{aligned} U < 2^7 : & \quad 0xxxxxxx \\ U < 2^{11} : & \quad 110xxxxx \ 10xxxxxx \\ U < 2^{16} : & \quad 1110xxxx \ 10xxxxxx \ 10xxxxxx \\ U < 2^{21} : & \quad 11110xxx \ 10xxxxxx \ 10xxxxxx \ 10xxxxxx \end{aligned}$$

The bits of U fill the **x** positions from the left. Every Devanagari code point lies between 2^{11} and 2^{16} , so it takes three bytes in UTF-8 against two in UTF-16. Because no continuation byte (**10xxxxxx**) can be mistaken for a leading byte, a reader can find character boundaries from any point in a stream, and an ASCII byte never appears inside a multi-byte character, which is why UTF-8 passes safely through tools written for ASCII.

Executed while this guide was built: `python3 examples/chars_demo.py`

1. The same letter KA in nine Unicode blocks: every block keeps ISCII's order, offset 0x15
Devanagari block U+0900 KA U+0915 क virama U+094D
Bengali block U+0980 KA U+0995 ক virama U+09CD
Gurmukhi block U+0A00 KA U+0A15 ਕ virama U+0A4D
Gujarati block U+0A80 KA U+0A95 ક virama U+0ACD
Oriya block U+0B00 KA U+0B15 ଳ virama U+0B4D
Tamil block U+0B80 KA U+0B95 க virama U+0BCD
Telugu block U+0C00 KA U+0C15 క virama U+0C4D
Kannada block U+0C80 KA U+0C95 ಕ virama U+0CCD
Malayalam block U+0D00 KA U+0D15 ക virama U+0D4D
2. UTF-8 by hand for U+0915 (KA): 16 bits split 4 + 6 + 6 into 1110xxxx 10xxxxxx 10xxxxxx
0000 100100 010101 -> 11100000 10100100 10010101 = E0 A4 95
Python agrees: E0 A4 95
the Devanagari block U+0900 to U+097F is E0 A4 80 to E0 A5 BF
3. Words are sequences of code points; aksharas are clusters of them

```
प्रकृति: 7 code points, 21 UTF-8 bytes: U+092A U+094D U+0930 U+0915 U+0943 U+0924 U+093F
पितृ: 4 code points, 12 UTF-8 bytes: U+092A U+093F U+0924 U+0943
क्षि: 4 code points, 12 UTF-8 bytes: U+0915 U+094D U+0937 U+093F
categories in क्षि: U+0915 Lo, U+094D Mn, U+0937 Lo, U+093F Mc
```

```
4. Normalisation: QA U+0958 decomposes, and composition does not rebuild it
NFD(U+0958) = U+0915 U+093C
NFC(that)    = U+0915 U+093C
equal as strings before normalising: False; after NFC on both: True
```

What the output shows

Clusters. A reader sees क्षि as one unit, a grapheme cluster in the terms of Unicode's UAX #29, but it is four code points: KA (a letter, category Lo), VIRAMA (a nonspacing mark, Mn), SSA (Lo) and VOWEL SIGN I (a spacing mark, Mc). The vowel sign is stored after the consonant and drawn before it. Reordering and conjunct formation happen when text is rendered, in a shaping engine such as HarfBuzz working from the font's OpenType tables. Storage order is phonetic; display order is visual. Cursor movement, deletion, search and string length therefore have to work on clusters, not on code points.

Normalisation. QA has its own code point, U+0958, and also a decomposed spelling, KA followed by NUKTA (U+0915 U+093C). The two look identical and compare unequal until normalised, and because U+0958 is a composition exclusion, normalising to NFC yields the decomposed form. A lexer that accepts identifiers in Indic scripts has to normalise them, or two names that look the same will be different names.

Joiners. ZERO WIDTH JOINER and ZERO WIDTH NON-JOINER (U+200D, U+200C) ask for, or prevent, a joined form: KA, VIRAMA, ZWJ, SSA requests a half-form of KA, and KA, VIRAMA, ZWNJ, SSA requests a visible virama. They are invisible, and invisible characters are a classic route for impersonation.

Confusables. Characters from different scripts that look alike, such as Latin `a` and Cyrillic `а`, let one name pose as another. Unicode's UTS #39 defines how to detect confusable strings, with its data in `confusables.txt`, and sets rules for mixing scripts in one identifier. Any system that accepts identifiers in several scripts should apply them.

A name has to survive the whole toolchain

Python 3 accepts identifiers in any script under PEP 3131, and GCC has accepted UTF-8 identifiers in C since version 10. That settles only the first step, and for systems work it settles the wrong one.

A keyword disappears during compilation. By the time a program is machine code, a loop written with `for` and the same loop written with `क्रम` have both become a comparison and a branch (Chapter 21). A name does not disappear. Every function, global variable, parameter and register mapping lives on: in the object file's symbol table, `.symtab` with its strings in `.strtab`; in the DWARF debugging information; in linker scripts and map files; in the kernel's own symbol table, from which a panic trace is printed; and in every tool a developer points at the binary: `nm`, `objdump`, `addr2line`, GDB, OpenOCD driving a JTAG or SWD probe on a microcontroller; a vendor's IDE. Each of these has its own rules about which bytes a symbol may contain, and the one alphabet every one of them accepts is letters, digits and the underscore.

So the question for firmware, kernels and microcontrollers is not whether a compiler accepts a name written in Devanagari. It is whether that name comes back, byte for byte, at every one of those tools, and can still be read against the source. Hindawi's answer, since 2003, is Romenagri: every name travels as a word over A to Z, a to z and the underscore, which every tool in the chain accepts, and the mapping is bijective, so the name in the script can be recovered from the binary at any point. Chapter 19 sets out the scheme and Chapter 21 runs it through GCC, the ELF symbol table, DWARF and GDB.

Composing conjuncts

Hindawi's APCISR composed half-consonant conjuncts of any length in software, in the years when display hardware for Indian scripts had fixed limits on them. Shaping engines solve the same problem today. It still matters wherever text must be shown without a full shaping stack, for example on a boot console before an operating system has loaded its

fonts.

19. Writing systems, sound, and reversible transliteration

In plain words. How Romenagri turns Indian scripts into plain English letters and back again without losing anything, and why Urdu written without its vowel marks is harder

Kinds of writing system

- An **alphabet** writes consonants and vowels as separate letters of equal standing: Latin, Greek, Cyrillic.
- An **abjad** writes consonants, with vowels left out or added as optional marks: Arabic, Hebrew, and the Perso-Arabic scripts of Urdu, Sindhi, Kashmiri and Punjabi in Shahmukhi.
- An **abugida** writes a consonant with an inherent vowel; other vowels are marks on the consonant, and a virama removes the vowel: Brahmi and all its descendants, Devanagari, Bengali, Gurmukhi, Gujarati, Odia, Tamil, Telugu, Kannada, Malayalam and more, and also Ethiopic.
- A **syllabary** has one sign per syllable: Japanese kana.
- A **logography** writes words and morphemes: Chinese characters. Egyptian hieroglyphs mix logograms with signs for sounds.

The Brahmi abugidas record the phonology of their languages closely, place of articulation by place of articulation, and that is why they can carry a whole computing stack: a mapping between their letters and a small ASCII alphabet can be made reversible, and the reversibility can be proved (below). In this work, full reversibility has so far been reached for the Brahmi family; for Latin written with diacritics, and for abjads that leave vowels unwritten, it has not yet been reached.

Script, language, standard: the three axes

ILM keeps three things apart.

Axis	What it is	Example
Script	How the text is written; handled by transliteration tables	Punjabi in Gurmukhi and in Shahmukhi: one language, two scripts
Language	The words: the vocabulary of keywords and names	Devanagari for Hindi, Marathi, Nepali and Sanskrit: one script, four languages
Standard	The computational construct and its host realisation	the C11 and C17 standards; a <code>for</code> loop in C, Python and Rust

A programming system that confuses these axes cannot grow past its first language. The pipeline in Chapter 21 applies them in order: script first, then language, then standard.

One spelling, several pronunciations

The consonant letter carries an inherent vowel, but languages pronounce it differently and drop it in different places.

Language	Inherent vowel	कमल (lotus) is said	Notes
Hindi, as spoken in Uttar Pradesh	a schwa, [ə]	kamal	the final inherent vowel is deleted, and some medial ones
Marathi	a schwa, [ə]	kamal	final deletion, as in Hindi
Bengali	an open o, [ɔ], raised to [o] before a high vowel in the next syllable	komol	the inherent vowel is usually kept inside a word and dropped at its end
Kannada	[a]	kamala	every vowel is pronounced; no deletion

The vowel written ऋ, ṛ in IAST, is a second example. Hindi as spoken in Uttar Pradesh says it as *ri*; Marathi, Gujarati, Kannada and Telugu say it as *ru*. So प्रकृति is *prakṛiti* in Hindi and *prakṛuti* in Marathi and Kannada, and Bengali, which also says *ri* but rounds the inherent vowel, has *prokṛiti*. The three spellings pitri, pitru and pitra seen for पितृ record three spoken realisations of the same written vowel.

The consequence is exact: a Unicode string records the spelling, not the sound. A text-to-speech system, a transliterator or a keyword table has to know the language, not only the script, and the same Devanagari text read in Lucknow and in Pune is a different sequence of sounds.

Transliteration, transcription and romanisation

Transliteration maps letters to letters and can be made reversible. **Transcription** maps sounds to symbols, usually IPA, and is not reversible to a spelling. **Romanisation** is any rendering in Latin letters. IAST and ISO 15919 transliterate with diacritics (ṛ, ṣ, ṇ); Harvard-Kyoto and ITRANS use plain ASCII with capitals and letter pairs.

Romenagri

The Romenagri Transliteration System has been written since 2003 and released under the GNU GPL. Its current tree is the `Romenagri` folder of [hindawiai/chintamani](#), with a Perso-Arabic seed library in [project-ilm/romenagri](#). It is ASCII-7 and uses no diacritics. Its alphabet is exactly the characters a C identifier may contain apart from digits: A to Z, a to z and the underscore. The underscore is its one marker. It prefixes the independent vowels (अ is `_a`, आ is `_aa`); it separates the dental stops from the retroflex ones (त is `_ta`, ट is `ta`); it marks ह, the anusvara and the visarga (`_ha`, `ka_m`, `ka_hh`); and it keeps apart the readings that a plain Latin spelling would merge: ख is `kha`, क followed by ह is `ka_ha`, and क joined to ह by a virama is `k_ha`. A consonant written with no vowel after it joins the next consonant, as the virama does: क्ष is `kx`, and प्र begins `pr`.

Hindawi uses it through four filters. `uni2acii` turns Unicode into ACII, the 8-bit code, laid out after ISCII, in which the tools work; `acii2cf` writes ACII in Romenagri's compiler form, in which every word in the script becomes its Romenagri word while English words and the contents of string literals pass through untouched; `rmn2acii` and `acii2uni` go back. The output below was produced by those filters, built from the retrieved sources by `examples/hindawi.sh`, not typed in, and every entry was converted back and compared.

Executed while this guide was built: `bash examples/hindawi.sh words`

```
Letters, as Devanagari = Romenagri:
अ = _a      आ = _aa     इ = _i      ई = _ee     उ = _u      ऊ = _oo
ऋ = _ri    ए = _aee    ऐ = _ai     ओ = _oa     औ = _ou    क = ka
ख = kha     ग = ga      घ = gha     ङ = _nga    च = cha     छ = chha
ज = ja      झ = jha     ञ = _yna    ट = ta      ठ = tha     ड = da
ढ = dha     ण = _nna    त = _ta     थ = _tha    द = _da     ध = _dha
न = na      प = pa      फ = pha     ब = ba      भ = bha     म = ma
य = ya      र = ra      ल = la      व = wa      श = sha     ष = xa
स = sa      ह = _ha     ळ = lvra    ऋ = kza    ज्ञ = jza    फ़ = phza
ड़ = _rda   ढ़ = _rdha  कं = ka_m   कः = ka_hh
```

Names, to Romenagri and back through the inverse:

```
योग -> yoaga -> योग
जोड़ो -> joa_rdoa -> जोड़ो
सीमा -> seemaa -> सीमा
गिनती -> gina_tee -> गिनती
प्रकृति -> prak_ri_ti -> प्रकृति
पितृ -> pi_t_ri -> पितृ
क्षि -> kxi -> क्षि
ज्ञान -> j_ynaana -> ज्ञान
हृदय -> _h_ri_daya -> हृदय
उत्तर -> _u_t_tara -> उत्तर
क्षेत्रफल -> kxae_traphala -> क्षेत्रफल
त्रिकोण -> _trikoa_nna -> त्रिकोण
ख -> kha -> ख
कह -> ka_ha -> कह
वह -> k_ha -> वह
```

67 of 67 come back identical; every Romenagri form uses only A-Z, a-z and _: yes

ISO 15919, the international standard for transliterating Indic scripts, has a 7-bit fallback of its own, and Romenagri's tree records the correspondence in `iso15919_map.csv`. That fallback spells with punctuation:

Letter	ISO 15919	ISO 15919, 7-bit	Romenagri
ट	ṭ	.t	ta
त	t	t	_ta
ण	ṇ	.n	_nna
ज	ñ	~n	_yna
श	ś	;s	sha
ष	ṣ	.s	xa
ऋ	ṛ	,r	_ri

The first three columns are from `iso15919_map.csv`; the Romenagri column is the filters' output above. A full stop, a semicolon, a comma or a tilde cannot appear in a C identifier, an assembler symbol or an unquoted linker-script name, so the standard's own ASCII form cannot carry a program's names through a toolchain. Romenagri can, and that is why it exists.

One identity across scripts

ISCII gave every Brahmi script one layout, and Hindawi builds on that. Its `flatten_uni_dev`, generated from `iscii_map.csv`, maps each letter of Bengali, Assamese, Gurmukhi, Gujarati, Oriya, Tamil, Telugu, Kannada and Malayalam to its Devanagari counterpart. Devanagari is the hub; Romenagri runs on the hub; filters such as `fltr_hi_te` carry the hub back out to a script. The same word written in three scripts therefore has one identity, `prak_ri_ti`, which the cover of this guide shows.




Figure 10. One identity across scripts: Brahmi scripts meet Devanagari in one hub; Perso-Arabic meets it at its written, canonical form.

This is the second pillar of linguistic equity at work. Because the script layer is independent of the language layer, any language can be written in any script, Hindi in Telugu letters or Telugu in Devanagari, and a program written either way is the same program to the compiler and the same names to the debugger.

Executed while this guide was built: `bash examples/hindawi.sh scripts`

1. The Brahmi hub: a word in Telugu or Bengali goes to Devanagari, to Romenagri, and back

```
ప్రకృతి -> ప్రకృతి -> prak_rī_tī -> ప్రకృతి
తెలుగు -> తెలుగు -> _taelugu -> తెలుగు
ఇకృతి -> ప్రకృతి -> prak_rī_tī -> ప్రకృతి
వారణా -> బాలా -> baa_mlaa -> బాలా
```

2. Any language in any script

```
Hindi written in Telugu script:   हिंदी भाषा -> హిందీ భాషా
Telugu written in Devanagari:   తెలుగు భాష -> తెలుగు భాష
Urdu written in Devanagari:      ہندوی کتاب -> हनदवी कताब
```

the same, with zer and sukun: ہندوی کتاب -> ہندوی کتاب

3. Measured on the corpora in this tree

Hindi, words of letters and vowel signs: 984 distinct, 983 identical after Romenagri and back

Telugu, through the hub and back: 1416 distinct, 1347 identical; of the 69 others, 69 carry the long-u sign U+0C42, which fltr_hi_te returns as U+0942

Urdu: 1 of 794 words carry any zabar, zer, pesh, sukun, shadda or tanwin

Of the 984 distinct Hindi words made of letters and vowel signs, 983 come back identical; the exception is a malformed word in the corpus, with three u signs in a row, which comes back in its canonical form. Of the 1,416 Telugu words taken through the hub and back, 1,347 are identical, and every one of the other 69 carries the long-u sign ు, for which the reverse filter in this tree has no rule, so it returns the Devanagari sign: a one-line gap in fltr_hi_te, not in Romenagri. The author's own measurements in [project-ilm/romenagri](http://project-ilm.com/romenagri) report 99.47 per cent of 5,807 words across the Hindi, Bengali and Telugu corpora exact. From the other side, over every string of up to four ASCII letters, his paper reports the 2003 to 2004 kernel as 98.68 per cent reversible or canonicalisable, with an irreducible floor of 1.31 per cent; the figure depends on the variant of the kernel measured, and his runbook says so.

The Perso-Arabic frontier

An abjad writes consonants and long vowels. The short vowels, zabar, zer and pesh, the absence of a vowel, sukun, and doubling, shadda, are marks that everyday Urdu leaves out: in the Urdu corpus in this tree, one word in 794 carries any of them. A written Urdu word therefore reaches the hub as what is written, its canonical form, and ہندوی arrives as हन्दवी, not हिंदवी. The mapping is exact and reversible to that form. What the script does not carry is which short vowel a speaker says.

With the marks written, the hub form is exact: ہندوی arrives as हन्दवी, the author's own example in [hindawi_tashkil.txt](#). So the difficulty is phonetic rather than a failure of the mapping. Which vowel a word carries has to come from knowledge of the language, its lexicon and its phonology, and that is precisely why ILM keeps the language axis apart from the script axis: the script layer maps what is written, exactly, and the language layer supplies what speakers know. The author estimates the words that need it at some 15 to 20 per cent.

Urdu programs already compile through the hub, as Chapter 21 shows. The reverse filter that renders hub text back into Urdu script, fltr_hi_ur, used by urducc -r, is a first cut in these trees; hardening the Perso-Arabic direction is the next step his seed release of Romenagri names. The Northwest Semitic abjads follow the same route to the Arabic hub: Hebrew final letter forms merge there, a documented many-to-one, and are reversible to the canonical form.

When is a transliteration reversible?

Let Σ be the source alphabet, the letters and marks of a script, and Γ the target alphabet, for example ASCII letters. A transliteration assigns each $\sigma \in \Sigma$ a non-empty target word $T(\sigma) \in \Gamma^+$ and extends to strings letter by letter:

$$T(\sigma_1\sigma_2\cdots\sigma_n) = T(\sigma_1)T(\sigma_2)\cdots T(\sigma_n).$$

It is reversible exactly when this extended map is injective, that is, when $T(s) = T(s')$ implies $s = s'$ for all strings $s, s' \in \Sigma^*$. Then there is an inverse with $T^{-1}(T(s)) = s$ for every string, which is the round-trip property a test can check. Injectivity holds if and only if T is one-to-one on single letters and the set of target words $C = T(\Sigma)$ is a *uniquely decodable code*.

Prefix codes. If no word of C is a prefix of another, every target string can be split into words from the left without looking ahead, so C is uniquely decodable.

Kraft and McMillan. If C is uniquely decodable over an alphabet of D symbols, with word lengths $\ell_1, \dots, \ell_m$, then

$$\sum_{i=1}^m D^{-\ell_i} \leq 1,$$

and whenever lengths satisfy this inequality, a prefix code with those lengths exists. A code whose sum exceeds 1 cannot be reversible, whatever else is done to it.

Sardinas and Patterson. For sets of words A and B write $A^{-1}B = \{w : aw \in B \text{ for some } a \in A\}$. Put $S_1 = C^{-1}C \setminus \{\varepsilon\}$ and $S_{k+1} = C^{-1}S_k \cup S_k^{-1}C$. The code is uniquely decodable if and only if no S_k contains a word of C . The sets are built from finitely many suffixes, so the test ends.

The example below runs both tests on four codes. The first is the mistake every naive scheme makes: k for क, h for ह and kh for ख. The string kh then reads both as ख and as क followed by ह, and the two tests catch it independently: the Kraft sum is $\frac{1}{2} + \frac{1}{2} + \frac{1}{4} = \frac{5}{4} > 1$, and the dangling suffix h is itself a codeword. The second is Romenagri's choice, as the output above shows: ह carries the marker, $_h$, so the letter h never stands alone as a codeword and kh has exactly one reading.

Executed while this guide was built: `python3 examples/ud_check.py`

```
naive: k, h, kh          D=2  Kraft sum=5/4    (1.250)  uniquely decodable: False
Romenagri's choice: k, _h, kh  D=3  Kraft sum=5/9    (0.556)  uniquely decodable: True
not prefix-free: a, ab, bb    D=2  Kraft sum=1      (1.000)  uniquely decodable: True
prefix code: 0, 10, 110, 111  D=2  Kraft sum=1      (1.000)  uniquely decodable: True

"kh" under the naive code has 2 readings: k + h; kh
"khk" under the naive code has 2 readings: k + h + k; kh + k
"abbb" under a, ab, bb has 1 reading: ab + bb
```

The third code shows that a code can be uniquely decodable without being a prefix code; the price is lookahead, since a followed by b cannot be settled until the next symbols arrive. An abugida adds one more demand. KA with its inherent vowel, KA with a virama, and KA in a conjunct must all remain distinguishable in the target. Romenagri meets it by writing the inherent vowel (ka), by letting a consonant with no vowel after it stand for the virama (k_ha , kx), and by marking what would otherwise merge; and it keeps its whole alphabet inside the identifier alphabet, so that every transliterated name is at once a valid C identifier, a valid assembler symbol and a valid linker name. The domain of reversibility, whether all Unicode strings or only well-formed syllables, still has to be stated rather than assumed; the 67 entries above are the check this guide ran.

20. Pāṇini's grammar as a formal system

In plain words. Pāṇini's grammar of Sanskrit, written well over two thousand years ago, works like a computer program, and this chapter runs part of it.

The Aṣṭādhyāyī, “eight chapters”, is a generative grammar of Sanskrit, usually dated between the sixth and the fourth centuries BCE. It has close to four thousand sūtras, 3,959 in the usual count, arranged in eight *adhyāyas* of four *pādas* each and cited by chapter, quarter and number, so that 6.1.77 is the seventy-seventh sūtra of the first quarter of the sixth chapter. It works with auxiliary lists: the Śiva Sūtras, which order the sounds; the Dhātupāṭha, the list of verbal roots; and the Gaṇapāṭha, lists of words that rules refer to as classes. Around it grew a review tradition: Kātyāyana's *vārttikas* on individual rules (around the third century BCE), Patañjali's *Mahābhāṣya* (around the second century BCE), the *Kāśikā* (seventh century CE), and Bhaṭṭoji Dīkṣita's *Siddhāntakaumudī* (seventeenth century), which rearranged the rules by topic for learners.

What follows treats the grammar the way a computer scientist would: as data structures and a rewriting system. Every claim below can be checked by running the program at the end of the chapter.

The Śiva Sūtras: a sound inventory with tags

Fourteen short lines list the sounds of the language. Each line ends in an *it*, a marker that is not pronounced and not part of the inventory; it is a tag.

No.	Devanagari	Sounds (IAST)	Marker
1	अ इ उ ण्	a i u	ṇ
2	ऋ लृ क्	ṛ ṝ	k
3	ए ओ ङ्	e o	ṅ
4	ऐ औ च्	ai au	c
5	ह य व र ढ्	h y v r	ṭ
6	ल ण्	l	ṇ
7	ज म ङ ण न म्	ñ m ṇ n	m
8	झ भ ञ्	jh bh	ṇ
9	घ ढ ध ष्	gh ḍh dh	ṣ
10	ज ब ग ड द श्	j b g ḍ d	ś
11	ख फ छ ठ थ च ट त व्	kh ph ch ṭh th c ṭ t	v
12	क प य्	k p	y
13	श ष स र्	ś ṣ s	r
14	ह ल्	h	l

Pratyāhāra: character classes by interval

Sūtra 1.1.71, *ādir antyena sahetā*, says that a first sound together with a later marker names every sound from the first up to that marker. So *ac* runs from *a* in line 1 to the marker *c* at the end of line 4 and names all the vowels; *hal* runs from *h* in line 5 to the final *l* and names all the consonants; *ik* is *i u ṛ ṝ*; *yaṇ* is *y v r l*. Each such abbreviation, a *pratyāhāra*, is exactly a character class defined by an interval over an ordered list, the same device as `[a-z]` in a regular expression, but over a list arranged so that the classes the grammar needs come out contiguous. The marker *ṇ* occurs twice, in lines 1 and 6, so *aṇ* is ambiguous between *a i u* and the longer run to line 6; the grammar knows from context which is meant. The sound *h* appears twice, in lines 5 and 14, so that it falls into the classes it must belong to.

How well the fourteen lines solve that arrangement problem has been studied formally: Kiparsky (1991) explained the arrangement by principles of economy, and Petersen (2004) analysed it mathematically, including why one sound has to be listed twice for every class the grammar needs to come out as an interval.

Kinds of sūtra

- *Samjñā*, a definition: 1.1.1 *vṛddhir ādaic* defines *vṛddhi* as ā, ai, au; 1.1.2 *adeṇ guṇaḥ* defines *guṇa* as a, e, o.
- *Paribhāṣā*, a meta-rule about how other rules are read and applied.
- *Vidhi*, an operation that changes a form.
- *Niyama*, a restriction on another rule.
- *Atideśa*, an extension that treats one thing as another.
- *Adhikāra*, a heading whose scope runs over the rules after it.

Compression and scope

A word stated in one sūtra continues into the following sūtras until something cancels it. This *anuvṛtti* lets a rule be written with only what is new in it; the rest is inherited from context, the way a variable declared in an enclosing block is visible inside it, or the way a default argument is supplied unless overridden.

Deciding which rule applies

Real grammars have rules that compete. The Aṣṭādhyāyī resolves conflicts with stated principles, and they are the part most worth studying.

- **General rule and exception.** An *apavāda*, the specific rule, blocks the *utsarga*, the general rule, where both could apply: the principle a pattern matcher follows when it prefers the more specific pattern.
- **1.4.2 *vipratīṣedhe paraṃ kāryam*.** In a conflict between rules of equal standing, the later rule in the text applies.
- **8.2.1 *pūrvatrāsiddham*.** The rules of the last three quarters of the eighth chapter, the *tripādī*, are treated as not having applied, as far as every earlier rule is concerned. It is an ordered pipeline: a late stage whose output the earlier stages are not allowed to see, as a compiler's later passes are invisible to its earlier ones.
- **1.1.56 *sthānivad ādeśo 'nalvidhau*.** A substitute behaves like what it replaced, except for rules about individual sounds: substitution through a shared interface.
- **1.3.2, 1.3.3 and 1.3.9.** Nasalised vowels and final consonants in teaching forms are markers, and markers are deleted: tags are stripped once they have done their work, as annotations are removed after parsing.

Four sandhi rules, computed

When words meet, their sounds change by rule. The program below builds the pratyāhāras from the table above by sūtra 1.1.71 and applies four rules of external sandhi to eight pairs.

- 6.1.101 *akaḥ savarṇe dīrghaḥ*: a simple vowel followed by a similar vowel becomes one long vowel.
- 6.1.77 *iko yaṇ aci*: a sound of class *ik* becomes the corresponding sound of class *yaṇ* before a vowel, *ac*. As a rewrite rule: $\{i, u, r, l\} \rightarrow \{y, v, r, l\} / _ ac$.
- 6.1.87 *ād guṇaḥ*: *a* or *ā* followed by a vowel of class *ik* gives the *guṇa* vowel.
- 6.1.88 *vṛddhir eci*: *a* or *ā* followed by *e*, *o*, *ai* or *au* gives the *vṛddhi* vowel.

File: `examples/panini_demo.py`

```
# panini_demo.py: the Śiva Sūtras, pratyāhāra formation (1.1.71) and three sandhi rules.
# © 1993–2026 Abhishek Choudhary. All rights reserved. AyeAI.
# SPDX-License-Identifier: GPL-3.0-or-later
# A teaching model in IAST. It is not tajziya's parser and makes no claim to cover the grammar.

SUTRAS = [ # (sounds, it-marker), the fourteen Maheshvara (Shiva) Sūtras
```

```

([ "a", "i", "u"], "ṇ"), ([ "ṛ", "ḷ"], "k"), ([ "e", "o"], "ṇ"), ([ "ai", "au"], "c"),
([ "h", "y", "v", "ṛ", "ṭ"], ([ "l"], "ṇ"), ([ "ñ", "m", "ṇ", "n"], "m"),
([ "jh", "bh"], "ṇ"), ([ "gh", "ḍh", "dh"], "ṣ"), ([ "j", "b", "g", "d", "d"], "ś"),
([ "kh", "ph", "ch", "ṭh", "th", "c", "ṭ", "t"], "v"), ([ "k", "p"], "y"),
([ "ś", "ṣ", "s"], "r"), ([ "h", "l"],
]

def pratyahara(first, marker, occurrence=1):
    """1.1.71 ādir antyena sahetā: a first sound with a final it-marker names every sound
    from the first up to the marker. Where a marker occurs twice (ṇ), say which one."""
    out, started, seen = [], False, 0
    for sounds, it in SUTRAS:
        for s in sounds:
            if s == first and not started:
                started = True
            if started:
                out.append(s)
            if started and it == marker:
                seen += 1
            if seen == occurrence:
                return list(dict.fromkeys(out)) # a sound listed twice (h) counts once
    raise ValueError(f"no pratyāhāra {first}{marker}")

AC = pratyahara("a", "c") # all vowels
HAL = pratyahara("h", "l") # all consonants
IK = pratyahara("i", "k")
YAN = pratyahara("y", "ṇ", 1) # the only ṇ after y is the second one in the list

LONG = {"a": "ā", "i": "ī", "u": "ū", "ṛ": "ṛ"}
SAVARNA = {k: k for k in LONG} | {v: k for k, v in LONG.items()}
YAN_OF = {"i": "y", "ī": "y", "u": "v", "ū": "v", "ṛ": "ṛ", "ṛ": "ṛ", "ḷ": "l"}
GUNA = {"i": "e", "ī": "e", "u": "o", "ū": "o", "ṛ": "ar", "ṛ": "ar"}
VRDDHI = {"e": "ai", "ai": "ai", "o": "au", "au": "au"}

def first_vowel(w):
    return w[:2] if w[:2] in ("ai", "au") else w[0]

def join(left, right):
    x, y = left[-1], first_vowel(right)
    if x in SAVARNA and y in SAVARNA and SAVARNA[x] == SAVARNA[y]: # 6.1.101 akaḥ savarṇe dīrghaḥ
        return left[:-1] + LONG[SAVARNA[x]] + right[len(y):], "6.1.101 akaḥ savarṇe dīrghaḥ"
    if x in YAN_OF and y in AC: # 6.1.77 iko yaṇ aci
        return left[:-1] + YAN_OF[x] + right, "6.1.77 iko yaṇ aci"
    if x in ("a", "ā") and y in VRDDHI: # 6.1.88 vṛddhir eci
        return left[:-1] + VRDDHI[y] + right[len(y):], "6.1.88 vṛddhir eci"
    if x in ("a", "ā") and y in GUNA: # 6.1.87 ād guṇaḥ
        return left[:-1] + GUNA[y] + right[len(y):], "6.1.87 ād guṇaḥ"
    return left + right, "no vowel sandhi"

if __name__ == "__main__":
    print("ac =", " ".join(AC))
    print("hal =", " ".join(HAL), f"({len(HAL)} consonants)")
    print("ik =", " ".join(IK), " yaṇ =", " ".join(YAN))
    print("aṇ (first ṇ) =", " ".join(pratyahara("a", "ṇ", 1)),
          " aṇ (second ṇ) =", " ".join(pratyahara("a", "ṇ", 2)))
    for a, b in [ ("dadhi", "atra"), ("madhu", "ari"), ("deva", "ālaya"), ("guru", "upadeśa"),
                  ("rāma", "iti"), ("mahā", "indra"), ("deva", "ṛṣi"), ("tava", "eva") ]:
        w, rule = join(a, b)
        print(f"{a} + {b} -> {w:12s} by {rule}")

```

Executed while this guide was built: `python3 examples/panini_demo.py`

```

ac = a i u ṛ ḷ e o ai au
hal = h y v r l ñ m ṇ ṇ n jh bh gh ḍh dh j b g ḍ d kh ph ch ṭh th c ṭ t k p ś ṣ s (33 consonants)
ik = i u ṛ ḷ yaṇ = y v r l
aṇ (first ṇ) = a i u aṇ (second ṇ) = a i u ṛ ḷ e o ai au h y v r l
dadhi + atra -> dadhyatra by 6.1.77 iko yaṇ aci
madhu + ari -> madhvāri by 6.1.77 iko yaṇ aci
deva + ālaya -> devālaya by 6.1.101 akaḥ savarṇe dīrghaḥ
guru + upadeśa -> gurūpadeśa by 6.1.101 akaḥ savarṇe dīrghaḥ
rāma + iti -> rāmeti by 6.1.87 ād guṇaḥ
mahā + indra -> mahendra by 6.1.87 ād guṇaḥ

```

deva + ṛṣi -> devarṣi	by 6.1.87 ād guṇaḥ
tava + eva -> tavaiva	by 6.1.88 vṛddhir eci

The program tries the rules in an order chosen by hand. Pāṇini's system does not need a hand-chosen order: the meta-rules above decide it, and that is the difference between a script that produces the right answers for eight examples and a grammar that produces them for the language.

Pāṇini and formal language theory

A modern reader will recognise context-sensitive rewrite rules, ordered rule application, meta-rules for conflict, and a generative system that derives surface forms from roots and affixes. In 1967 Peter Ingerman proposed in *Communications of the ACM* that the notation used to define programming languages, Backus-Naur Form, be called Pāṇini-Backus Form, since Pāṇini's notation is equivalent in power. Pāṇinian derivation has been implemented in code, for example in vidyut-prakriya, whose sūtra text tajziya uses under the MIT licence. Tamil has its own grammatical tradition beginning with the Tolkāppiyam, and ILM's parsers are organised by language family so that such traditions sit side by side.

Questions to work through

1. Compute the pratyāhāras *jhal* and *yar* from the table by hand. Which sounds does each contain?
2. Why must *h* appear twice in the Śiva Sūtras?
3. Derive *madhu + ari* by hand. Which rule applies, and why does 6.1.101 not apply?
4. What would go wrong in a derivation if 8.2.1 were removed?
5. Compare 1.4.2 with the way a parser generator settles a shift-reduce conflict by precedence (Chapter 21).
6. Pāṇini's inventory lists sounds; Unicode's Devanagari block lists written signs. Which of the two is a phonology and which a script, and what does each leave out?

21. Compilers, from lexer to microcode

In plain words. How a program written in Hindi or Urdu becomes instructions a machine can run, with every name kept intact all the way to the debugger.

A compiler is a chain of translations, each into a representation closer to the machine: source text, tokens, a syntax tree, a checked tree, an intermediate representation, optimised intermediate code, assembly, an object file, an executable, a running process, and finally, inside the processor, micro-operations.

Lexing: regular languages and automata

Tokens are described by regular expressions. Thompson's construction turns a regular expression of length n into a nondeterministic finite automaton with at most $2n$ states; the subset construction turns that into a deterministic automaton whose states are sets of the original states, up to 2^n in the worst case and few in practice; Hopcroft's algorithm minimises the result in $O(n \log n)$ time.

lex (Lesk and Schmidt, Bell Labs, 1975) generates such a scanner in C from a list of patterns and actions; flex is the fast, free lex that Linux systems carry. flex works on bytes, so a keyword written in Devanagari is simply a pattern of UTF-8 bytes, and the whole Devanagari block is the pattern `\xE0\xA4[\x80-\xBF]|\xE0\xA5[\x80-\xBF]`, whose end points the output in Chapter 18 printed.

Parsing: context-free grammars

A grammar $G = (N, \Sigma, P, S)$ has nonterminals, terminals, productions and a start symbol. The classic expression grammar is

$$\begin{aligned} E &\rightarrow T E' & E' &\rightarrow + T E' \mid \epsilon \\ T &\rightarrow F T' & T' &\rightarrow * F T' \mid \epsilon \\ F &\rightarrow (E) \mid \text{id} \end{aligned}$$

A predictive, LL(1), parser chooses a production by looking at one token, guided by two families of sets. With $\neg$ marking the end of input:

$$\begin{aligned} \text{FIRST}(E) &= \text{FIRST}(T) = \text{FIRST}(F) = \{ (, \text{id} \} \\ \text{FIRST}(E') &= \{ +, \epsilon \}, \quad \text{FIRST}(T') = \{ *, \epsilon \} \\ \text{FOLLOW}(E) &= \text{FOLLOW}(E') = \{), \neg \} \\ \text{FOLLOW}(T) &= \text{FOLLOW}(T') = \{ +,), \neg \} \\ \text{FOLLOW}(F) &= \{ +, *,), \neg \} \end{aligned}$$

LR parsing (Knuth, 1965) reads left to right, shifting tokens onto a stack and reducing them when a production's right side is on top, which builds a rightmost derivation in reverse. LALR(1) (DeRemer, 1969) merges LR(1) states that share a core, which keeps the tables small enough for practice. yacc (Johnson, Bell Labs) generates LALR(1) parsers in C, and GNU Bison is its compatible successor; the name is a joke on yacc, since a yak and a bison are cousins. When a grammar allows two moves, the generator reports a shift-reduce or reduce-reduce conflict, and declarations such as `%left '+' '-'` and `%left '*' '/'` settle it by precedence and associativity: a rule-ordering device of the kind Pāṇini's 1.4.2 states for his grammar. The standard example is the dangling else, `if a then if b then s1 else s2`, which has two parse trees; yacc resolves it by shifting, so the `else` joins the nearest `if`.

Meaning: symbol tables and syntax-directed translation

After parsing, names are resolved through symbol tables, one per scope, and types are checked. An attribute grammar attaches values to the nodes of the tree: synthesised attributes flow up from children, inherited attributes flow down from parents and siblings, and definitions that use only synthesised attributes, or inherited attributes from the left, can be evaluated during parsing itself. This is what the GATE syllabus calls syntax-directed translation.

Intermediate representation

Three-address code has at most one operator per instruction:

```
t1 = x * 4
t2 = t1 + 6
return t2
```

Instructions without internal jumps form basic blocks, and the blocks form a control-flow graph. In static single assignment form every variable is assigned exactly once; where control paths join, a φ -function chooses the value from whichever predecessor was taken, $x_3 = \varphi(x_1, x_2)$. SSA is built by placing φ -functions at the dominance frontiers of the definitions (Cytron and colleagues, 1991). GCC's middle end works on GIMPLE in SSA form; LLVM's IR is SSA throughout.

Dataflow analysis as a fixed point

A variable is live at a point if its current value may still be read. For each instruction n ,

$$in[n] = use[n] \cup (out[n] \setminus def[n]), \quad out[n] = \bigcup_{s \in succ(n)} in[s].$$

Start with every set empty and apply the equations until nothing changes. The sets form a lattice ordered by inclusion, the equations are monotone, and the lattice has finite height, so the iteration stops at the least fixed point. Reaching definitions, available expressions and constant propagation are the same pattern over different lattices.

Optimisation, watched on a real compiler

The usual transformations are constant folding and propagation, common-subexpression elimination, dead-code elimination, moving loop-invariant code out of loops, strength reduction (a multiplication by four becomes a shift or an address computation), inlining, unrolling and vectorisation.

File: `examples/fold.c`

```
/* examples/fold.c: two small functions, to watch an optimising back end at work.
   © 1993-2026 Abhishek Choudhary. All rights reserved. AyeAI.
   SPDX-License-Identifier: GPL-3.0-or-later */
int scale(int x) { int k = 2 * 3; return x * 4 + k; }
int sum_to(int n) { int s = 0; for (int i = 1; i <= n; i++) s += i; return s; }
```

Captured while this guide was built: `gcc -O2 -S -masm=intel -fno-asynchronous-unwind-tables -fcf-protection=none -fno-ident -o - examples/fold.c | sed -n '/^scale:/,/size.sum_to/p'`

```
scale:
    lea eax, 6[0+rdi*4]
    ret
.size    scale, .-scale
.p2align 4
.globl  sum_to
.type   sum_to, @function
sum_to:
    test edi, edi
    jle .L6
    lea ecx, 1[rdi]
    xor edx, edx
    and edi, 1
    mov eax, 1
    je .L5
    mov eax, 2
    mov edx, 1
    cmp eax, ecx
    je .L3
.p2align 4,,10
.p2align 3
```

```

.L5:
    lea edx, 1[rdx+rax*2]
    add eax, 2
    cmp eax, ecx
    jne .L5
.L3:
    mov eax, edx
    ret
    .p2align 4,,10
    .p2align 3
.L6:
    xor edx, edx
    mov eax, edx
    ret
    .size    sum_to, .-sum_to

```

`scale` became a single instruction. `lea eax, 6[0+rdi*4]` computes $4x + 6$ in the address unit: the constant 2×3 was folded, and the multiplication by four was reduced to scaled addressing. In `sum_to`, GCC peeled one iteration when n is odd and then unrolled the loop by two, adding i and $i + 1$ in one step as $2i + 1$ with another `lea`. The exact output depends on the compiler's version and target; this one was produced by the GCC recorded in the package's `docs/BUILD.json`.

Registers, instructions and scheduling

Register allocation is graph colouring. Each variable's live range is a node; two nodes are joined when the variables are live at the same time; colouring the graph with k colours assigns k registers. Deciding whether a graph is k -colourable is NP-complete for $k \geq 3$, so compilers use heuristics: Chaitin's simplify-and-spill (1982), Briggs's optimistic colouring, and linear scan in just-in-time compilers where speed of compilation matters most. Instruction selection tiles the IR tree with machine instructions, instruction scheduling reorders them to hide latencies, and a peephole pass cleans up short sequences.

Assembling, linking, loading

The assembler writes an object file, ELF on Linux, with sections (`.text`, `.data`, `.rodata`, `.bss`), a symbol table, and relocations marking the places whose final addresses are not yet known. The linker resolves symbols across objects and libraries and applies the relocations; the dynamic loader maps shared libraries into the process when it starts.

Below the instruction set: micro-operations and microcode

On current high-performance processors, architectural instructions are decoded into internal micro-operations, and complex instructions are sequenced from microcode, the lowest layer at which code is still written by people; microcode can be updated after the chip is made. GATE's computer organisation syllabus asks for the design of hardwired and microprogrammed control units, and a microprogrammed control unit is exactly a small engine that runs microcode to drive the datapath.

Front ends and back ends: why a shared back end is not a dependency

With m source languages and n target machines, separate compilers would number mn ; with a shared intermediate representation, m front ends and n back ends suffice, $m + n$ in all. GCC's front ends for C, C++, Objective-C, Fortran, Ada, Go, D, Modula-2 and others all lower to GENERIC and GIMPLE and then RTL, and its back ends serve dozens of processors. LLVM's IR is shared by Clang, Rust, Swift, Julia and others. Nobody says that Rust depends on C++ because LLVM is written in C++; a new front end reusing a mature back end is the ordinary path.

The back end never sees keywords. By the time a program is GIMPLE or RTL, a loop written with `for` and the same loop written with `क्रम` have become the same structure: basic blocks, a condition and a branch. The human language of the keywords is settled and discarded at the front end. That is the precise answer to the objection that a system which uses GCC as its back end depends on English; the same holds for PANINIq's use of Qiskit, where what crosses the boundary is a Hamiltonian, a matrix of complex numbers. What the back end does see is the program's names, which it must carry into the object file intact; that is the script layer's job, and the last sections of this chapter run it.

Bootstrapping and the self-hosting fixed point

A compiler has three languages: the source *S* it reads, the target *T* it writes, and the language *I* it is written in, drawn as a T-shaped tombstone diagram. Diagrams chain: a compiler for *X* written in C, built by an existing C compiler, runs on the machine; with it you compile a compiler for *X* written in *X*, and from then on the language compiles itself. The test of self-hosting is a fixed point: stage 1, built by the old compiler, builds stage 2; stage 2 builds stage 3; stages 2 and 3 must be identical. GCC's own `make bootstrap` performs this three-stage build and compares the last two stages.

A reversible transliteration (Chapter 19) plays the part of the first compiler in such a chain. It lets the whole existing toolchain process programs written in the script, and its inverse restores the script in the output and the diagnostics, so the native-script tools can then be brought up the same way. Hindawi's own keyword lexer was brought up exactly so, through Romenagri, as the next sections show.

Trust: the precise meaning of a sovereign stack

Ken Thompson showed in 1984 that a compiler can be altered to insert a back door into the login program, and to insert the same alteration into every compiler it later compiles, so that the back door survives after both source files have been cleaned. Reading the source is not enough. David Wheeler's diverse double-compiling (2005, and his dissertation of 2009) answers it: compile the compiler's source with a second, independently developed compiler; use the result to compile the source again, and compare. If the output is bit-for-bit identical to the original self-compiled binary, that binary corresponds to its source, under stated assumptions.

Together with reproducible builds, in which the same source always yields the same bits, this gives the phrase “sovereign stack” an exact technical meaning: a stack that its users can rebuild from source and verify with independent tools, with people able to do both. Which alphabet the keywords use has nothing to do with it.

Hindawi's pipeline, run

The Hindawi Programming System has nine *shailis*, each a front end over one host language:

Shaili	Host
गुरु guru	C
श्रेणी shraeni	C++
यंत्र yantra	assembly
कृत्रिम kritrima	Java
प्राथमिक praatha	BASIC
सूची soochee	Python
व्याकरण wyaaka	yacc
रोबोट robot	LOGO
शब्द shabda	lex

A program begins with a line naming its shaili, `<शैली गुरु>` for C, and one driver, `hincc`, compiles every shaili:

```
cat $1 | fixuninum | iconv -f utf-8 -t utf-16 | uni2acii | acii2cf | hincc.awk
```

`fixuninum` turns Devanagari numerals into ASCII digits; `iconv` and `uni2acii` turn Unicode into ACII; `acii2cf` writes the whole program in Romenagri's compiler form, every keyword and every name, leaving string literals as they are; and `hincc.awk` reads the shaili line, which is now `<shailee guru>`, and hands the rest of the program to that shaili's driver. For C that is `gurucc`:

```
cat $1 | acii2uni | iconv -f UTF-16 -t UTF-8 | h2c > tempfil0123.tmphin.c
gcc tempfil0123.tmphin.c -o hin.exe -lm
```

Its `acii2uni` restores the string literals to Unicode, so that the program prints in the script, and `h2c` maps each Romenagri keyword to its C keyword: `poor_nnaa_mka` (पूर्णक) to `int`, `krama` (क्रम) to `for`, `ma_likhoa` (म_लिखो) to `printf`, `mukhya` (मुख्य) to `main`. Every other word is a name, and it reaches GCC in Romenagri. The stages run in the order of the three axes: the script first, in `hincc`; then the language, in `h2c`, and for lex and yacc sources in `h2l` and its companions before it; then the standard, in flex and GCC.

The keyword lexer is itself written in Hindi. Its source, `h2c.uhin`, passes through the same `uni2acii | acii2cf` before flex compiles it, so the tool that reads Hindi programs is brought up through its own script layer, the bootstrapping step described above. The transducers also compose: a lex or yacc action block is C, so `h2l` carries only what lex adds beyond C, mapping `श_शब्द` to `yyllex`, `श_ब_मान` to `yylval`, `श_माला` to `yytext` and `श_पंक्ति` to `yylineno`, and hands the rest to `h2c`. Counted that way, shabda has 6 rules of its own and 321 through `h2c`, 327 in all, and wyaaka, for yacc, has 341. Reading a composing shaili on its own would report a working system as broken.

The run below builds Romenagri, the guru shaili and `hincc` from the retrieved sources in `vendor/chintamani` with their own Makefiles, then compiles and runs Hindawi's own C sample:

Executed while this guide was built: `bash examples/hindawi.sh sample`

```
$ hincc HindiC.uhin
संकलन के परिणाम
=====

The C that gurucc handed to gcc:
#include <stdio.h>

int main()
{
    char _a[80];
    int ka;
    printf("आपका नाम क्या है?\n");
    scanf("%s", _a);
    printf("नमस्ते %s.\n", _a);
    for(ka=1; ka<=10; ka++)
        printf("%d\n", ka);
    scanf("%s", _a);
    return 0;
}

$ ./hin.exe, with राम typed at both prompts
आपका नाम क्या है?
नमस्ते राम.
1
2
3
4
5
6
7
8
9
10
```

अ has become `_a` and क has become `ka`; the strings still print in Devanagari.

The symbol bridge, from source to debugger

This program, in the guru shaili, has a global, a function, a parameter and a local, all named in Hindi:




Figure 11. The symbol bridge: a Hindi program through hincc, GCC and the debugger; its names carried in Romenagri and rendered back.

File: `examples/yog.uhin`

```

<शैली गुरु>
#समावेश <मानकपन. स>

पूर्णांक योग;

पूर्णांक जोड़ो(पूर्णांक सीमा)
{
    पूर्णांक गिनती;

```

```

क्रम(गिनती=1; गिनती<=सीमा; गिनती++)
योग = योग + गिनती;
वापस योग;
}

पूर्णांक मुख्य()
{
म_लिखो("१ से १० तक का योग = %d\n", जोड़ो(10));
वापस 0;
}

```

Executed while this guide was built: `bash examples/hindawi.sh yog`

```

$ hinc yog.uhin
संकलन के परिणाम
=====

The C that gurucc handed to gcc:
#include <stdio.h>

int yoaga;

int joa_rdoa(int seema)
{
    int gina_tee;
    for(gina_tee=1; gina_tee<=seema; gina_tee++)
        yoaga = yoaga + gina_tee;
    return yoaga;
}

int main()
{
    printf("1 से 10 तक का योग = %d\n", joa_rdoa(10));
    return 0;
}

$ ./hin.exe
1 से 10 तक का योग = 55

```

योग is `yoaga`, जोड़ो is `joa_rdoa`, सीमा is `seema` and गिनती is `gina_tee`: names over A to Z, a to z and the underscore, which GCC, the assembler, the linker and every later tool accept as they are. Below, the same program is built with debugging information and examined with the ordinary tools, and then the same transcript is rendered through the Romenagri inverse. The binary is not touched; only what a person reads changes. The C names that Hindawi has words for, `main` and `printf`, come back through `c2h` as they do in Hindawi's own `std2hin`; the program's own names come back through `rmn2acii` and `acii2uni`.

Captured while this guide was built: `bash examples/hindawi.sh debug`

```

$ nm --defined-only yog-g, the program's own symbols
0000000000001149 T joa_rdoa
0000000000001182 T main
0000000000004014 B yoaga

$ readelf --debug-dump=info yog-g, the names DWARF records
variable yoaga
subprogram printf
subprogram main
subprogram joa_rdoa
formal_parameter seema
variable gina_tee

$ gdb -batch: break joa_rdoa, run, backtrace, info args, finish, print yoaga
Breakpoint 1 at 0x1154: file ./tempfil0123.tmphin.c, line 8.

Breakpoint 1, joa_rdoa (seema=10) at ./tempfil0123.tmphin.c:8
8      for(gina_tee=1; gina_tee<=seema; gina_tee++)
#0  joa_rdoa (seema=10) at ./tempfil0123.tmphin.c:8
#1  0x000055555555194 in main () at ./tempfil0123.tmphin.c:15
seema = 10
0x000055555555194 in main () at ./tempfil0123.tmphin.c:15

```

```

15     printf("1 से 10 तक का योग = %d\n", joa_rdoa(10));
Value returned is $1 = 55
$2 = 55

$ addr2line -f -s -e yog-g 0x0000000000001149
joa_rdoa
tempfil0123.tmphin.c:6

The same, rendered through the Romenagri inverse; the binary is unchanged:
$ nm --defined-only yog-g, the program's own symbols
0000000000001149 T जोड़ो
0000000000001182 T मुख्य
0000000000004014 B योग

$ readelf --debug-dump=info yog-g, the names DWARF records
variable योग
subprogram म_लिखो
subprogram मुख्य
subprogram जोड़ो
formal_parameter सीमा
variable गिनती

$ gdb -batch: break जोड़ो, run, backtrace, info args, finish, print योग
Breakpoint 1 at 0x1154: file ./tempfil0123.tmphin.c, line 8.

Breakpoint 1, जोड़ो (सीमा=10) at ./tempfil0123.tmphin.c:8
8     for(गिनती=1; गिनती<=सीमा; गिनती++)
#0 जोड़ो (सीमा=10) at ./tempfil0123.tmphin.c:8
#1 0x0000555555555194 in मुख्य () at ./tempfil0123.tmphin.c:15
सीमा = 10
0x0000555555555194 in मुख्य () at ./tempfil0123.tmphin.c:15
15     म_लिखो("1 से 10 तक का योग = %d\n", जोड़ो(10));
Value returned is $1 = 55
$2 = 55

$ addr2line -f -s -e yog-g 0x0000000000001149
जोड़ो
tempfil0123.tmphin.c:6

```

Three layers do three jobs here. The front end lowers keywords to C, where they vanish into control flow. The symbol table and the DWARF information carry every name in Romenagri, byte for byte, through the assembler and the linker into the ELF file, and from there to `nm`, `addr2line`, GDB, OpenOCD on a JTAG or SWD probe, or a kernel's panic trace. A presentation layer applies the Romenagri inverse wherever a person reads: a GDB front end, an IDE panel, a filter on `addr2line` or on a panic trace. No link in the chain has to accept anything but letters, digits and the underscore, and nothing is lost, because the map is bijective. Drop the script layer and pass names as raw UTF-8, or flatten them into names that cannot be mapped back, and some link in that chain either rejects them or hands the developer names that no longer match the source.

Two details are worth knowing. GDB lists the generated C because the guru shaili does not emit `#line` directives; with them, the same session would list the Hindi source line by line. And Romenagri writes an independent vowel with a leading underscore, so a vowel-initial name at file scope, `_aapakaa` for आपका, begins with an underscore, which the C standard reserves for the implementation at that scope; GCC accepts it, and a project that exports such symbols can fix a prefix for them.

The Urdu edition, run

`zistgah/urdu-ilm` carries the same shailis for Urdu. The guru shaili's lexer is written in Urdu, `h2c_urdu.uhin`, and the driver `urducc` sends an Urdu program through `fixuninum` and `fixarabnum` for the digits, `fltr_ur_hi` to the Devanagari hub, `uni2acii` and `acii2cf` to Romenagri, then `h2c_urdu` and the standard Hindawi driver. The run below compiles the Urdu edition's own C sample:

File: `vendor/urdu-ilm/ILM/samples/UrduC.uhin`

```

<شیلی گرو>
# شامل <معیاری.س>
( ) صحیح مرکزی
{

```

```

; حرف ن [80]
; صحیح ک
; لکھو ("آپ کا نام کیا ہے؟" );
; لکھو ("ن", "%s" );
; لکھو ("ن", "السلام علیکم %s.\n" );
; برائے (ک=1 ; ک>=10 ; ++ک)
; لکھو ("د\n", "ک");
; واپس 0
}

```

Executed while this guide was built: `bash examples/hindawi.sh urdu`

```

$ urducc UrduC.uhin
संकलन के परिणाम
=====

The C that reached gcc:
#include <stdio.h>
int main()
{
    char na[80];
    int ka;
    printf("आप का नाम क्या है?\n");
    scanf("%s",na);
    printf("इलसलाम एलीकम %s.\n",na);
    for(ka=1; ka<=10; ka++)
        printf("%d\n",ka);
    return 0;
}

$ ./hin.exe, with Ali typed at the prompt
आप का नाम क्या है?
इलसलाम एलीकम Ali.
1
2
3
4
5
6
7
8
9
10

```

The keywords are C and the names are Romenagri, reached through the hub: ن is `na` and ک is `ka`. The strings print in the hub script and show the frontier of Chapter 19 at work: ک, written without its short vowel, arrives as कया rather than क्या, and ے as हे rather than है. Rendering them back into Urdu is the reverse filter's job.

ILM's construct model

In the merged release, `ilm/constructs.csv` holds 39 constructs in 27 human languages and `ilm/decorators.csv` holds 216 realisations of them in 8 host languages; a keyword table of 1,801 constructs in 27 languages generates the Hindawi distributions. The construct is primary: a counted loop is translated once per human language, not once per host language, and a decorator records what a host requires, such as `stdio.h` for printing in C. The construct model is the keyword layer; the layer `h2c` implements for C. Names still pass through the script layer, Romenagri, exactly as above.

22. Below the software: logic, chips, ternary and qubits

In plain words. What happens inside chips: switches, circuits, three-valued logic, and the very cold machines that hold qubits.

From logic to CMOS

Every digital function reduces to Boolean algebra, and NAND alone is enough to build all of it. In CMOS, a network of p-type transistors pulls the output up and a dual network of n-type transistors pulls it down, so that in a settled state no current flows from supply to ground. Power is spent mainly in switching, $P \approx \alpha C V^2 f$, with α the fraction of nodes switching per cycle, C the capacitance switched, V the supply voltage and f the clock frequency; the V^2 is why every generation of chips has lowered its voltage.

How a chip is designed

1. Describe the hardware at register-transfer level in Verilog, VHDL or SystemVerilog, or generate it from Chisel or Amaranth.
2. Simulate it, with Icarus Verilog or Verilator.
3. Synthesise it into a netlist of standard cells, with Yosys.
4. Floorplan, place, build the clock tree and route, with OpenROAD.
5. Check design rules and compare layout against schematic, with Magic, KLayout and Netgen.
6. Write the layout as GDSII or OASIS and send it to the foundry: the tapeout.
7. Fabricate, test and package.

Open process design kits let anyone take this path to real silicon: SkyWater SKY130 (130 nm), GlobalFoundries GF180MCU (180 nm) and IHP SG13G2 (130 nm, SiGe BiCMOS). zistgah/jugaad28 catalogues integrated circuits taped out on 28 nm, with evidence-gated descriptors and a low-cost open laboratory. A hardware description language is a language like any other: its front end can be given native-language keywords by the three-axis method of Chapter 21, and what reaches the synthesiser is a netlist, not words.

Instruction sets and the boot chain

The instruction set is the contract between software and hardware. RISC-V is an open one: a small base, RV32I or RV64I, with standard extensions for multiplication (M), atomics (A), floating point (F and D), compressed instructions (C) and vectors (V), which anyone may implement without a licence fee. An open instruction set, open design tools and reproducible toolchains carry the argument of Chapter 21 down to the silicon.

A computer starts through a chain of programs: a boot ROM, the firmware (BIOS or UEFI), a boot loader such as GRUB, the kernel, the first user process such as systemd, and then everything else. Each is a program in some language, and each can be given a native-language front end without changing what the hardware executes; the hindawiai organisation keeps a fork of the Linux kernel tree, [hinlin](https://hinlin.org), for exactly this work.

Balanced ternary: Zamin and PRATIK

Balanced ternary uses the digits -1 , 0 and $+1$, written T, 0 and 1:

$$n = \sum_i t_i 3^i, \quad t_i \in \{-1, 0, +1\}.$$

It needs no sign: a number is negated by flipping every digit, and rounding is simply truncation. It is also the most economical integer base. Representing numbers up to N in base b takes about $\log_b N$ digits of b states each, a cost proportional to

$$b \log_b N = \frac{b}{\ln b} \ln N,$$

which is least at $b = e \approx 2.718$; among whole numbers, 3 wins over 2 and 4. The Setun computer, built at Moscow State University in 1958, ran on balanced ternary.

[zistgah/zamin](#) is the physical realisation of the PRATIK balanced-ternary substrate for the PEDLER event algebra. Its poised zero is a physical high-impedance ground, so that the electrical ground reference and the computational zero are the same node, and its spawning operator creates new physical dimensions on a memristive three-dimensional crossbar. The repository holds SPICE cells, an interactive circuit simulator and a paper; [zistgah/pratik_core_mvp](#) is the kernel in C++ and CUDA.

Executed while this guide was built: `python3 examples/ternary.py`

```

5 =      1TT   negated by flipping every digit:      T11 = -5
-5 =      T11   negated by flipping every digit:      1TT = 5
13 =     111   negated by flipping every digit:      TTT = -13
40 =     1111  negated by flipping every digit:     TTTT = -40
-121 =    TTTT  negated by flipping every digit:    11111 = 121
2026 = 10T10001 negated by flipping every digit: T01T000T = -2026

```

```

radix economy b/ln b for b = 2.000: 2.8854
radix economy b/ln b for b = 2.718: 2.7183
radix economy b/ln b for b = 3.000: 2.7307
radix economy b/ln b for b = 4.000: 2.8854
radix economy b/ln b for b = 10.000: 4.3429

```

Superconducting qubits

A Josephson junction is two superconductors separated by a thin insulator, usually aluminium, aluminium oxide and aluminium. Its current and voltage obey

$$I = I_c \sin \varphi, \quad V = \frac{\hbar}{2e} \frac{d\varphi}{dt},$$

where φ is the phase difference across it. It behaves as a nonlinear inductor, and the nonlinearity is what makes the lowest two energy levels of a circuit usable as a qubit. The transmon, the most common design, has the Hamiltonian

$$H = 4E_C (\hat{n} - n_g)^2 - E_J \cos \hat{\varphi},$$

and is operated with E_J/E_C around 50, which makes it nearly insensitive to stray charge (Koch and colleagues, 2007). Its transition frequency lies around 4 to 8 GHz.

The cold follows from that frequency. Thermal energy has to be far below the energy of one microwave photon at the qubit frequency. At 10 mK, which is -273.14°C , $k_B T/h$ is about 0.21 GHz, well below 5 GHz; at the 4 K of liquid helium it would be about 83 GHz, and the qubit would be swamped by thermal excitation. Hence the dilution refrigerator; the large gold-coloured structure in photographs of quantum computers; the processor itself is a small chip at its bottom.

That chip is an integrated circuit: aluminium or niobium films on silicon or sapphire, patterned by lithography from a layout file like any other chip, through processes related to, but distinct from, CMOS logic processes. Qiskit Metal is IBM's open tool for laying out and analysing such chips. Other ways to build qubits are trapped ions, neutral atoms, photons and spins in silicon; nitrogen-vacancy centres in diamond keep their spin coherence at room temperature and are used today as sensors.

PANINIq's oscillator substrate, Chapter 23, is different in kind: a classical network of coupled oscillators that can be built from ordinary electronics and behaves as an Ising machine. Its hardware bridge is mocked today; its mathematics is not.

23. Quantum computing and PANINIq, with the mathematics

In plain words. How quantum computers work, and how a network of swinging oscillators can solve the same kind of puzzle.

States, gates and measurement

A qubit's state is a unit vector in $\mathbb{C}^2$,

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle, \quad |\alpha|^2 + |\beta|^2 = 1,$$

which, up to a global phase, is a point on the Bloch sphere: $|\psi\rangle = \cos \frac{\theta}{2} |0\rangle + e^{i\phi} \sin \frac{\theta}{2} |1\rangle$. The state of n qubits lives in the tensor product of n copies, a space of 2^n complex amplitudes. Gates are unitary matrices:

$$X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad Z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}, \quad H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \quad \text{CNOT} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}.$$

A Hadamard followed by a CNOT entangles two qubits: $\text{CNOT}(H \otimes I)|00\rangle = \frac{1}{\sqrt{2}}(|00\rangle + |11\rangle)$. Measuring in the computational basis gives outcome x with probability $|\langle x|\psi\rangle|^2$, the Born rule. Between measurements the state evolves under a Hamiltonian H by $i\hbar \frac{d}{dt}|\psi\rangle = H|\psi\rangle$, so that, in units with $\hbar = 1$,

$$|\psi(t)\rangle = e^{-iHt} |\psi(0)\rangle.$$

Simulating n qubits exactly means storing 2^n amplitudes of 16 bytes each: 30 qubits need 16 GiB. That is why a laptop simulates up to about thirty qubits, and why the encoding of a problem into qubits matters so much.

Ising, QUBO and Max-Cut

An Ising problem assigns spins $s_i \in \{-1, +1\}$ to minimise

$$E(s) = -\sum_{i<j} J_{ij} s_i s_j - \sum_i h_i s_i.$$

With $x_i = (1 - s_i)/2 \in \{0, 1\}$ the same problem becomes a quadratic unconstrained binary optimisation (QUBO). Max-Cut asks for the split of a weighted graph's vertices into two sides that maximises the weight of the edges crossing between them:

$$C(s) = \sum_{(i,j) \in E} w_{ij} \frac{1 - s_i s_j}{2}.$$

Each term is w_{ij} when the edge is cut and 0 when it is not, so maximising C is minimising $\sum w_{ij} s_i s_j$, the energy of an anti-ferromagnet with $J_{ij} = -w_{ij}$. Max-Cut is NP-hard; the best classical guarantee in polynomial time, Goemans and Williamson's semidefinite relaxation with random rounding, reaches 0.878 of the optimum.

QAOA

The Quantum Approximate Optimisation Algorithm prepares

$$|\gamma, \beta\rangle = \prod_{k=1}^p e^{-i\beta_k H_M} e^{-i\gamma_k H_C} |+\rangle^{\otimes n}, \quad H_C = \sum_{(i,j) \in E} w_{ij} \frac{1 - Z_i Z_j}{2}, \quad H_M = \sum_i X_i,$$

and a classical optimiser adjusts the $2p$ angles to maximise $\langle \gamma, \beta | H_C | \gamma, \beta \rangle$; measuring the final state gives candidate cuts. With $p = 1$ on 3-regular graphs, the expected cut is at least 0.6924 of the optimum (Farhi, Goldstone and Gutmann, 2014). QAOA uses one qubit per vertex, so exact simulation costs 2^N amplitudes; PANINIq's QAOA baseline therefore stops at $N \leq 12$ and uses the COBYLA optimiser.

The continuous-time quantum walk, compactly encoded

A continuous-time quantum walk (Farhi and Gutmann, 1998) evolves a state over the vertices of a graph under a Hamiltonian built from the graph: $|\psi(t)\rangle = e^{-iHt}|\psi(0)\rangle$. PANINIq encodes the N vertices as the basis states of $n = \lceil \log_2 N \rceil$ qubits, a binary address, so its Hilbert space has $2^n < 2N$ dimensions rather than 2^N . Its Hamiltonian is the coupling matrix placed in the top-left block of a $2^n \times 2^n$ zero matrix, and the evolution is one arbitrary unitary gate:

```
# pedler/quantum_validator.py in zistgah/paniniq, commit 82a5cf2 (excerpt)
num_qubits = max(1, int(np.ceil(np.log2(max(N, 2)))))
dim = 2 ** num_qubits
H = np.zeros((dim, dim), dtype=complex)
H[:N, :N] = K_matrix
...
U_t = expm(-1j * H * evolution_time)
qc.unitary(Operator(U_t), list(range(num_qubits)), label="CTQW(K,t)")
```

The test suite checks that the circuit's statevector equals e^{-iHt} applied directly. A general unitary on n qubits compiles into $O(4^n)$ elementary gates, so this is a construction for validating the classical dynamics, not an efficient circuit for hardware; the repository says the same.

The Jelly Substrate: coupled phase oscillators

The substrate is a network of phase oscillators obeying the Kuramoto equations, with coupling that decays with distance:

$$\frac{d\theta_i}{dt} = \omega_i + \sum_j K_{ij} \sin(\theta_j - \theta_i), \quad K_{ij} = K_0 e^{-\alpha r_{ij}}, \quad K_{ii} = 0.$$

```
# pedler/substrate.py in zistgah/paniniq, commit 82a5cf2 (excerpt)
def kuramoto_rhs(self, t, theta):
    phase_diffs = theta[None, :] - theta[:, None]
    return self.omegas + np.sum(self.K * np.sin(phase_diffs), axis=1)
```

How synchronised the network is shows in the order parameter,

$$R e^{i\psi} = \frac{1}{N} \sum_{j=1}^N e^{i\theta_j}, \quad 0 \leq R \leq 1,$$

near 1 when the phases lock together and near 0 when they are scattered. For identical all-to-all coupling K/N , each oscillator feels only the mean field, $\dot{\theta}_i = \omega_i + KR \sin(\psi - \theta_i)$, and synchrony appears once K exceeds $K_c = 2/(\pi g(0))$, where g is the symmetric, single-peaked distribution of natural frequencies; for a Lorentzian of half-width γ , $K_c = 2\gamma$.

Oscillators as an Ising machine

Coupled oscillators can minimise an Ising energy directly (Wang and Roychowdhury, 2019). With an added second-harmonic injection of strength K_s ,

$$\frac{d\theta_i}{dt} = -K \sum_j J_{ij} \sin(\theta_i - \theta_j) - K_s \sin(2\theta_i),$$

the function

$$\mathcal{E}(\theta) = -K \sum_{i,j} J_{ij} \cos(\theta_i - \theta_j) - K_s \sum_i \cos(2\theta_i)$$

never increases along the motion, so the network settles into its minima. The injection pushes each phase to 0 or π , where $\cos(\theta_i - \theta_j) = s_i s_j$ with $s_i = \cos \theta_i = \pm 1$, and $\mathcal{E}$ becomes the Ising energy. PANINIq's Max-Cut solver uses the simplest form of this: it sets the coupling to minus the adjacency, so that neighbours repel in phase, relaxes, and reads each spin as the sign of $\cos \theta_i$. Its tests check the result against brute force on six graphs: never above the optimum, and exact on the bipartite ones.

```
# pedler/maxcut.py in zistgah/paniniq, commit 82a5cf2 (excerpt)
sub.K = -1.0 * adjacency          # connected nodes repel in phase
...
spins = np.sign(np.cos(final_phases))
```

The PEDLER automaton on the substrate

PANINIq implements PEDLER as a six-tuple $P = (I, G, U, S, F, *)$. U is the stream of point events, vectors in $\mathbb{R}^3$. S is the set of states, one per oscillator; and F maps each state to a point in $\mathbb{R}^3$. G is the resolution-match kernel, $\exp(-\gamma \|F(v_i) - r_t\|)$, the similarity between an event r_t and each state. I is the inclination field over the states, a historical prior plus the weighted resolution match, computed afresh for each event. The operator $*$ governs growth: when an event is presented and the relaxed network fails to lock, its order parameter falling below the novelty threshold, a new state is spawned; on a separate schedule, faded states are pruned and strongly locked ones consolidated. “Inclination field”, “resolution match” and “point event” are the framework’s own terms; the code implements the equations as given and claims nothing more about cognition. The repository also records one open question plainly: PEDLER’s canonical inclination is signed, lives on edges and lies in $[-1, +1]$, while this code’s is non-negative and lives on nodes, and which one the substrate should carry is the author’s decision.

What is tested, and what is open

From the repository’s own status table: the Jelly Substrate, the PEDLER engine, the Max-Cut solvers, the frame packing of the hardware bridge, the Hamiltonian and unitary mathematics, the Qiskit circuit and QAOA are **tested**; the PyTorch substrate is **written** and not executed; the hardware port is **mocked**; a Duffing extension with amplitude as well as phase, a real firmware protocol, and the larger system running from brain signals through symbolic orchestration to a physical instruction set and manufacturing are **planned**. The run in Chapter 14 shows all of it: states spawning as the order parameter drops, a three-qubit walk whose statevector sums to one, and the oscillator, brute force and QAOA agreeing on a cut of 4.

Exercises

1. Compute $H|0\rangle$, $H|1\rangle$ and HZH .
2. Show that $\frac{1}{2}(1 - s_i s_j)$ is 1 exactly when an edge is cut.
3. For a graph with 12 vertices, how many qubits do the walk and QAOA each need, and how many amplitudes does each simulation store?
4. Differentiate $\mathcal{E}$ along the motion and show that $d\mathcal{E}/dt \leq 0$ when J is symmetric.
5. In PANINIq, change the novelty threshold and predict, before running, how the number of spawned states changes.

Part V: Working with the estate

24. AI, from MYCIN to transformers, and how to use any AI here

In plain words. How AI grew from rule books to today's language models, and how to use any AI, safely, to help with this work.

Two traditions

The first tradition in AI wrote knowledge down as rules. MYCIN, built at Stanford in the 1970s (Shortliffe, 1976), identified the bacteria behind severe infections and recommended antibiotics with about six hundred production rules of the form *if these findings, then this conclusion, with this certainty*. Each rule carried a certainty factor between -1 and $+1$, and two positive factors for the same conclusion combined as

$$CF = CF_1 + CF_2(1 - CF_1).$$

In evaluations its recommendations stood comparison with those of infectious-disease specialists, yet it was never used on patients, for reasons of integration, accountability and trust rather than accuracy: the first lesson of applied AI, and still the current one. Its name comes from the suffix of the antibiotics it prescribed, streptomycin, erythromycin and the rest, which in turn comes from the Greek *mykēs*, fungus, because the soil organisms that yield many of them were once taken for fungi. Take the medical rules out and the inference engine that remains, EMYCIN, became the pattern for a generation of expert systems.

The second tradition learns from data: perceptrons, then multilayer networks trained by backpropagation, then deep learning. The estate uses both. PANINI and ILM are symbolic, with grammars, constructs and transducers; PEDLER is adaptive and event-driven.

Transformers

A language model reads tokens, not letters. Text is cut into subword units by byte-pair encoding or a similar scheme, learned from a training corpus. Tokenisers trained mostly on English cut Indic text into many more tokens for the same content, and since every Devanagari letter is already three bytes in UTF-8, the same sentence costs more to process, fills the context window sooner and is modelled less well. Script-aware normalisation and tokenisation, of the kind ILM provides, address exactly this.

Each layer of a transformer mixes information between positions by attention,

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V,$$

in several heads at once, with the queries Q , keys K and values V computed from the tokens (Vaswani and colleagues, 2017). The model is trained to predict the next token, maximising $\sum_t \log p_\theta(x_t | x_{<t})$ over a very large corpus, and then tuned on human preferences.

Does AI only pull and pool data?

Partly. A model is trained on existing text, and it can state errors fluently, so everything it says is a claim to be checked, not a fact. But it does not look documents up and paste them: it computes a probability distribution over continuations, and it can apply a method to material it has never seen. The stance that follows is the estate's verification-gated co-development: nothing is accepted without an oracle, whether a test suite, a second implementation, a measurement or a review. The examples in this guide were written with an AI assistant and accepted only because they ran, and their outputs were checked against the sources.

Choosing an AI

Any of them will do for the method below. Chat assistants include Claude, ChatGPT, Gemini, Grok, DeepSeek and Mistral's Le Chat. They differ in their models, context lengths and tools, and they change every few months.

Local models keep everything on your own machine. With Ollama:

```
curl -fsSL https://ollama.com/install.sh | sh
ollama run llama3.2
```

A GPU with 4 GB of memory runs models of a few billion parameters quantised to four bits; llama.cpp and LM Studio are alternatives.

Coding agents work in the terminal, inside a repository, reading files and running commands with your permission. The install commands are each vendor's documented ones at the time of writing:

Agent	Install	Start it inside a repository
Claude Code	<code>npm install -g @anthropic-ai/claude-code</code>	<code>claude</code>
Codex CLI	<code>npm install -g @openai/codex</code>	<code>codex</code>
Gemini CLI	<code>npm install -g @google/gemini-cli</code>	<code>gemini</code>
Aider, with many providers and local models	<code>python3 -m pip install aider-install && aider-install</code>	<code>aider</code>

The method: a cycle with any AI

The estate's cyler protocol works with every one of them.

1. **Intent.** One sentence saying what you want.
2. **Context.** Give the AI the repository's own files before anything else: `README.md`, `CONTEXT.md`, `CONTRACT.md`, `AGENTS.md` where present, and `ops/verify.sh`. Many estate repositories carry these files precisely so that any person or any AI can start cold. For this guide, the whole text is one file, `docs/rahnuma.md`, and `docs/llms.txt` lists its chapters.
3. **A meaningful prompt.** Name the files, the exact change, and the check that will decide whether it is accepted.
4. **Inspect.** Run the tests or the gate yourself and read the diff. The AI's report that something passed is not evidence that it passed.
5. **Keep the record.** Commit the prompt, the response and the result together.
6. **Next prompt,** until the artifact is one you authored with intention.

A prompt that works, for PANINIq:

Here are README.md, CONTRACT.md, tests/test_core.py and ops/verify.sh from zistgah/paniniq. Add one test to tests/test_core.py that checks the order parameter R lies in [0, 1] for 200 random phase vectors. Show the new test, then the output of `python3 -m unittest discover -s tests` and of `bash ops/verify.sh`.

Four rules keep this safe. Never paste a password, token or private key into any AI. Do not let an agent push, publish or mint; those steps stay with a person. Do not accept code you have not run. And for anything current, such as versions, dates or prices, make the AI search, or check it yourself.

25. Open source, provenance and contributing

In plain words. How to share, copy and improve this work in the open, and how each piece is stamped with its date and a permanent reference.

Git as a data structure

A Git repository is a directed acyclic graph of commits. Each commit names a tree, the hash of the directory's contents, its parent commits, its author and its message, and every object is addressed by the hash of its content. Change one byte anywhere and every hash above it changes, which is why a commit hash pins an exact state of every file. A branch is a movable name for a commit; a merge commit has two parents.

Contributing on GitHub

```
gh repo fork zistgah/paniniq --clone
cd paniniq
git switch -c order-parameter-bounds
bash ops/verify.sh
git commit -am "Test that the order parameter stays within [0, 1]"
git push -u origin order-parameter-bounds
gh pr create --fill
```

Between switching branches and verifying comes your change. `gh repo fork --clone` also adds the original repository as a remote named `upstream`, so `git fetch upstream && git rebase upstream/main` brings your branch up to date. To report a problem instead of fixing it, `gh issue create --repo zistgah/paniniq` opens an issue.

Licences

The GNU GPL, version 3, is copyleft: anyone may use, study, change and share the software, and whoever distributes a changed version must distribute its source under the same terms. Permissive licences such as MIT and Apache 2.0 allow changes to be kept closed; Apache adds an explicit patent licence. Creative Commons Attribution-ShareAlike 4.0 does for text and media what the GPL does for code. Among the `zistgah` and `project-ilm` repositories, every licence that GitHub recognises is a version of the GPL; others carry licence files that GitHub does not recognise, or none that it detects. In every case the repository's own licence file, where it has one, is the authority. This guide's text is CC BY-SA 4.0 and its scripts are GPL-3.0-or-later.

White-labelling, and the commons

White-labelling is selling a product made by one party under another's brand. Under a permissive licence a company can take open code, change it and sell it closed; under the GPL it may still sell it, but it must pass on the source and the same freedoms. Contributing upstream is the opposite of white-labelling: the improvement returns to everyone.

The Aṣṭādhyāyī's review tradition is an old form of the same practice. Kātyāyana's *vārtikas* proposed amendments to specific rules, Patañjali's *Mahābhāṣya* discussed and adjudicated them, and the rules were kept intact alongside the commentary, the way a pull request and its review thread are kept alongside the code.

DOIs

A DOI is a persistent identifier that `doi.org` resolves to a landing page. Zenodo, run by CERN, issues DOIs through DataCite and keeps two kinds: a concept DOI that always points to the latest version, and a version DOI for each release. `project-ilm/misty-doi` writes the metadata and makes the deposit; its commands are in Chapter 14.

Timestamps with OpenTimestamps

A timestamp proves that a file existed before a certain time without revealing the file. The file's SHA-256 hash is one leaf among many; pairs of hashes are hashed together up a Merkle tree,

$$h_{\text{parent}} = \text{SHA256}(h_{\text{left}} \parallel h_{\text{right}}),$$

and the root is written into a Bitcoin transaction. The proof for one file is the path of sibling hashes from its leaf to the root, $\lceil \log_2 N \rceil$ of them for N files, together with the block that contains the root. Anyone can recompute the path and check that the file existed before that block was mined. The commands are `ots stamp FILE`, then `ots upgrade FILE.ots` once the block is confirmed, and `ots verify FILE.ots`.

Attestations

An in-toto statement is a signed JSON document that binds a subject, file names with their digests, to a predicate saying what was done to them: built, tested, reviewed. It travels in a DSSE envelope, and the SLSA framework defines levels of supply-chain assurance built from such statements. Candor records the estate's typed intents in this form, under each repository's `attest/` folder.

How the estate releases

Every repository carries a `CONTRACT.md` of numbered clauses and an `ops/verify.sh` that checks them and reports PASS, FAIL or UNJUDGED for each. A push happens only when the gate is green, and it writes a receipt, a timestamp, a manifest of hashes in `MANIFEST.sha256` and a seal; a mint then records the DOI in `CITATION.cff` and in the README. The governance repository's contract sets twelve clauses for all of this: sole authorship and AyeAI as the only affiliation; the copyright line; no fabrication; verification by execution; reading before writing; typed gates before anything irreversible, whether a mint, a push, a deletion, an issue or a pull request; staying inside the working folder; provenance travelling with every publication; fork, pull request and merge for shared repositories; small reviewable increments; an agreed split of roles between AIs; and declared repository classes. Clauses C31 to C37 add retrieval before assumption, no vendor lock-in, experience that scales with the user, certification in isolation and in integration, context that travels with every surface, routing as a preference rather than a property, and never bluffing. This package follows the same pattern; its contract is in `CONTRACT.md` and its gate is `bash ops/verify.sh`.

26. GATE: sit it, whatever your age

In plain words. India's national engineering examination: what it covers, how to apply, and why it is worth sitting at any age.

The facts for 2027

The Graduate Aptitude Test in Engineering 2027 is organised by IIT Madras. There are 30 test papers, including Robotics and Automation (RA) for the first time; a candidate may sit one paper or an allowed pair of two. The regular registration, first due to close on 27 September 2026, has been extended to 5 October 2026 without a late fee, and the extended registration, with a late fee, closes on 12 October 2026, all on the portal gate2027.iitm.ac.in. Registration goes through DigiLocker for Indian nationals. The application correction window runs from 14 to 21 October 2026, the examination is held in February 2027, and results are due on 19 March 2027. There is no upper age limit: a candidate needs to be in the third or a later year of an undergraduate degree, or to have completed one. A GATE score is valid for three years and opens postgraduate admissions and recruitment in public-sector undertakings.

The official syllabi

- Computer Science and Information Technology (CS): [CS_GATE2027_Syllabus.pdf](#)
- Robotics and Automation (RA): [RA_GATE2027_Syllabus.pdf](#)
- Instrumentation Engineering (IN): [IN_GATE2027_Syllabus.pdf](#)
- General Aptitude, common to every paper (GA): [GA_GATE2027_Syllabus.pdf](#)

Where the CS syllabus meets this guide

CS syllabus section	Chapters
Engineering mathematics: discrete mathematics, linear algebra, probability	11 (codes and the Kraft inequality), 13 (lattices and fixed points), 15 (vectors, unitary matrices, eigenvalues)
Digital logic: Boolean algebra, circuits, number representation and computer arithmetic	9, 14
Computer organisation and architecture: instructions and addressing, data path and control unit, pipelining	13 (microprogrammed control), 14
Programming and data structures	the examples throughout, in Python and C
Algorithms: graph algorithms and complexity, P and NP	13 (graph colouring), 15 (Max-Cut)
Theory of computation: regular and context-free languages, automata, Turing machines	12 (rewrite systems), 13 (automata and grammars)
Compiler design: lexing, parsing, syntax-directed translation, intermediate code, local optimisation, dataflow analysis	13, directly
Operating systems	8 (namespaces and control groups), 14 (the boot chain)
Computer networks	9 (network byte order); not otherwise covered
Databases	not covered

For Robotics and Automation, read the RA syllabus beside the estate repositories that meet it in practice: [zistgah/pench](#), the embodied cyler; [zistgah/transeg](#), the local digital twin; and [zistgah/jugaad28](#) and [zistgah/zamin](#) on the hardware side.

Courses and the reference lab

Courses at GATE level for both CS and RA, aligned with the reference lab and its dockerised launch scripts, are **planned** as the next step after this guide. They will be announced in the zistgah organisation.

27. Syllabus links, and our syllabus

In plain words. Links from your school or exam syllabus to the chapters that teach the same ideas, worked out on your own device, and a draft syllabus of our own.

A student's syllabus and this guide cover many of the same ideas under different names. This chapter links them without copying a word of any syllabus: the guide stores only the official address of each syllabus and its own list of keywords, and the matching happens on the reader's own device.

How the link works. Choose a board or exam; the page opens its official syllabus, from the publisher's own site. If the browser may read it directly, the page reads it; if the publisher does not allow that, which most do not, save the syllabus from the official page and drop the file on this page. Either way, the text is read in your browser, matched against the guide's own keywords, and every match becomes a link into the chapter that teaches it. Nothing is uploaded, stored or sent anywhere, and the syllabus remains the property of its publisher.

The boards and exams linked

Board or exam	Body	Levels	Region	Official page
GATE 2027, Computer Science and Information Technology	IIT Madras	Graduate	India	gate2027.iitm.ac.in
GATE 2027, Robotics and Automation	IIT Madras	Graduate	India	gate2027.iitm.ac.in
GATE 2027, Instrumentation Engineering	IIT Madras	Graduate	India	gate2027.iitm.ac.in
GATE 2027, General Aptitude	IIT Madras	Graduate	India	gate2027.iitm.ac.in
GATE 2027, every paper	IIT Madras	Graduate	India	gate2027.iitm.ac.in
CBSE, classes 9 to 12	Central Board of Secondary Education	Secondary, senior secondary	India	cbseacademic.nic.in
NCERT textbooks	National Council of Educational Research and Training	Classes 1 to 12	India	ncert.nic.in
ICSE, class 10	Council for the Indian School Certificate Examinations	Secondary	India	cisce.org
ISC, class 12	Council for the Indian School Certificate Examinations	Senior secondary	India	cisce.org
NIOS	National Institute of Open Schooling	Secondary, senior secondary	India	www.nios.ac.in
Cambridge IGCSE	Cambridge International Education	Secondary	International	www.cambridgeinternational.org
Cambridge International AS and A Level	Cambridge International Education	Senior secondary	International	www.cambridgeinternational.org
IB Diploma Programme	International Baccalaureate	Senior secondary	International	www.ibo.org
IB Middle Years Programme	International Baccalaureate	Middle, secondary	International	www.ibo.org

Board or exam	Body	Levels	Region	Official page
Advanced Placement courses	College Board	Senior secondary	United States	apcentral.collegeboard.org
Common Core State Standards	National Governors Association and CCSSO	Kindergarten to grade 12	United States	www.thecorestandards.org
AQA GCSE and A level	AQA	Secondary, senior secondary	United Kingdom	www.aqa.org.uk
Pearson Edexcel qualifications	Pearson	Secondary, senior secondary	United Kingdom	qualifications.pearson.com
OCR qualifications	OCR	Secondary, senior secondary	United Kingdom	www.ocr.org.uk
Australian Curriculum, version 9	ACARA	Foundation to year 10	Australia	v9.australiancurriculum.edu.au
Singapore national examinations	SEAB	Secondary, pre-university	Singapore	www.seab.gov.sg
Telangana Board of Secondary Education	Government of Telangana	Secondary	India, state board	bse.telangana.gov.in
Andhra Pradesh Board of Secondary Education	Government of Andhra Pradesh	Secondary	India, state board	bse.ap.gov.in
Maharashtra State Board	Government of Maharashtra	Secondary, higher secondary	India, state board	mahahsscboard.in
Karnataka School Examination and Assessment Board	Government of Karnataka	Secondary	India, state board	kseab.karnataka.gov.in
Tamil Nadu Directorate of Government Examinations	Government of Tamil Nadu	Secondary, higher secondary	India, state board	dge.tn.gov.in
Kerala SCERT	Government of Kerala	School	India, state board	scert.kerala.gov.in
West Bengal Board of Secondary Education	Government of West Bengal	Secondary	India, state board	wbbse.wb.gov.in
Uttar Pradesh Madhyamik Shiksha Parishad	Government of Uttar Pradesh	Secondary, intermediate	India, state board	upmsp.edu.in
Board of Secondary Education, Rajasthan	Government of Rajasthan	Secondary, senior secondary	India, state board	rajeduboard.rajasthan.gov.in
Gujarat Secondary and Higher Secondary Education Board	Government of Gujarat	Secondary, higher secondary	India, state board	www.gseb.org
Punjab School Education Board	Government of Punjab	Secondary, senior secondary	India, state board	www.pseb.ac.in

The GATE papers' addresses were retrieved for this guide on 29 September 2026; the other addresses are the publishers' own sites, where the current syllabus is published. Each carries its review mark in the register.

Our syllabus, a draft for review

A standard syllabus built from this guide, level by level, with its own outcomes, chapters and laboratory work. It is a draft for the author's review, and each module carries its review mark in the register at the end of the guide.

Foundation, classes 6 to 8

Module	Title	Hours	Outcomes	In this guide	Laboratory
F1	Numbers inside the machine	6	Count in binary and convert to decimal; Explain why a computer needs only two symbols	Chapter 17	
F2	Letters, scripts and your language	6	Show that every letter is stored as a number; Write a word in two scripts and say what stays the same	Chapter 18, Chapter 19	
F3	The sky and the calendar	4	Read sunrise, moonrise and the tithi for a day; Explain why calendars differ	Chapter 10	chakra, jyotish
F4	Making with a cyclus, safely	4	Make a short story or picture with a cyclus; Say what you checked before keeping the result	Chapter 24	genie

Secondary, classes 9 and 10

Module	Title	Hours	Outcomes	In this guide	Laboratory
S1	Binary, logic and circuits	10	Build truth tables for basic gates; Relate a gate to transistors	Chapter 17, Chapter 22	
S2	Encoding and reversible transliteration	8	Explain UTF-8; Run a Romenagri round trip and judge whether it is exact	Chapter 18, Chapter 19	hindawi
S3	Programming in your own language	10	Compile and run a Hindi program; Find the same names in the debugger	Chapter 21	hindawi, urdu
S4	Using AI and checking it	6	Use any AI on a task and record each step; Tell an AI's claim from a verified result	Chapter 24, Chapter 8	alam

Senior secondary, classes 11 and 12

Module	Title	Hours	Outcomes	In this guide	Laboratory
H1	Grammar as a formal system	10	Write a regular expression and a small grammar; Read a Pāṇinian rule as a rewrite rule	Chapter 20, Chapter 21	tajziya
H2	From source code to machine	12	Trace a program from source through compiler, assembler and linker; Read symbols in an ELF file	Chapter 21, Chapter 22	hindawi

Module	Title	Hours	Outcomes	In this guide	Laboratory
H3	Open source, provenance and contributing	6	Make a commit, open a pull request and pass a gate; Explain a hash, a manifest and a DOI	Chapter 25	rahnuma
H4	A project seeded from FAKIR	15	Choose a domain, a language and a layer and state a gate; Deliver the work with its evidence	Chapter 29	fakir

Undergraduate and GATE

Module	Title	Hours	Outcomes	In this guide	Laboratory
U1	Theory of computation and compilers	30	Automata, grammars and parsing; Build and debug a compiler pipeline	Chapter 20, Chapter 21, Chapter 26	hindawi , panini_by_claude
U2	Digital logic and architecture	30	From gates to a processor; Ternary logic and its trade-offs	Chapter 22, Chapter 26	pratik , zamin
U3	Quantum and physical computing	20	Qubits, gates and QAOA; Oscillator networks and Ising problems	Chapter 23	paniniq

Postgraduate and research

Module	Title	Hours	Outcomes	In this guide	Laboratory
P1	Research on the open questions	60	Choose an open question from Appendix C or D; Publish a measured, reproducible result	Appendix C, Appendix D, Appendix F	paniniq , tajziya

Professional

Module	Title	Hours	Outcomes	In this guide	Laboratory
Q1	Adopting a component	12	Run a component's gate; Integrate it with its contract and provenance intact	Chapter 15, Chapter 25, Appendix E	dhancha , humanesque

28. Questions that come up, answered technically

In plain words. Short technical answers to questions people often ask, such as whether AI needs English, and who controls it.

These questions are asked whenever the work is discussed. Each answer points to the chapter that carries the evidence.

If it uses GCC, or Qiskit, does it not still depend on foreign technology?

A compiler is front ends and back ends joined by an intermediate representation (Chapter 21). By the time a program reaches GCC's GIMPLE, a loop written with `for` and the same loop written with `क्रम` are the same basic blocks; the language of the keywords is settled and discarded at the front end. What does reach the back end is the program's names, and Hindawi hands those over in Romenagri, so that they come out of the binary intact and map back to the script (Chapter 21). Reusing a mature back end is how Rust, Swift and Julia were all built on LLVM, and nobody calls them dependent on C++. PANINiq hands Qiskit a Hamiltonian, which is a matrix of complex numbers. The dependency that does matter is on source code and on people who can rebuild it, which is the next answer.

What would a sovereign stack actually require?

Source code for every layer; builds that reproduce bit for bit, independent verification of the compilers by diverse double-compiling, and enough trained people to rebuild, verify and modify each layer (Chapter 21). None of this depends on the alphabet of the keywords, and all of it depends on skills. That is why the next steps here are courses and laboratories, not a logo.

Is this not white-labelling borrowed technology?

White-labelling sells someone else's product under a new brand, often closed. Everything here is the opposite: public source under the GPL, dated DOIs, and forks and pull requests invited (Chapter 25). What is new is stated precisely: the language front ends, the construct model with its decorators, the reversible transliteration, the transducers and the substrates. What is reused is reused openly and named, as every serious compiler project reuses back ends, libraries and standards.

Are Unicode and InScript not enough?

Unicode encodes characters and InScript is a keyboard layout. Together they let you type and store Indic text (Chapter 18). Neither translates keywords, libraries, compiler messages or documentation; neither tells a program which language a Devanagari string is in; and neither makes a name written in the script survive an assembler, a linker script, an ELF symbol table, DWARF, a debugger and a JTAG probe and come back readable. The first two gaps are the language front ends; the third is the script layer's, Romenagri's (Chapters 19 and 21).

Is a script not the same as a language?

No. Punjabi is written in Gurmukhi and in Shahmukhi; Devanagari writes Hindi, Marathi, Nepali and Sanskrit, and pronounces the same letters differently in each (Chapter 19). A system that ties language to script cannot grow past its first language, which is why ILM keeps script, language and standard on separate axes.

Should there not be a conference first?

A conference is where results are presented and checked against each other, and results begin with running code. The repositories are open now, so the most useful contribution anyone can make today is an issue or a pull request against something that runs. A conference with running code behind it has something to test.

Is 2047 the horizon, or 2030?

The number of people who can build decides the date. Courses at GATE level and reference labs are how that number grows, which is why they come next (Chapters 16 and 26). The estate's own schedule is in Chapter 31.

Is quantum hardware not out of reach?

Much less than it looks. A laptop simulates about thirty qubits exactly (Chapter 23), and PANINIq runs its quantum walk and QAOA baseline on one. Cloud services give access to real quantum processors, some with free tiers. Oscillator Ising machines can be built from ordinary electronics, and chips can be designed and taped out on open process design kits (Chapter 22). What is out of reach is building a dilution refrigerator from scratch, and that is not where the work starts.

What has the Indian knowledge tradition contributed that is technical?

Specific, checkable things. Pāṇini's grammar is a formal generative system with character classes, ordered rewrite rules and meta-rules for conflict (Chapter 20). Piṅgala's *Chandaḥśāstra* enumerates the patterns of long and short syllables in metres with rules equivalent to counting in binary, and its *meru-prastāra* is the triangle of binomial coefficients later named after Pascal. Brahmagupta stated rules for arithmetic with zero in 628 CE. Āryabhaṭa's *kuṭṭaka* solves linear equations in integers, the problem the extended Euclidean algorithm solves. Mādhava of Saṅgamagrāma, in the fourteenth century, gave infinite series for the arctangent, sine and cosine, including

$$\frac{\pi}{4} = 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots,$$

about two centuries before they appeared in Europe. Where a tradition is said to be not yet understood, the useful response is the one Chapter 20 takes: encode it, run it, and let the output speak.

Does AI depend on ASCII, and would it stop working without it?

There is no separate “ASCII foundation” to cut off. ASCII is the first 128 code points of Unicode and the one-byte range of UTF-8 (Chapter 18). A current language model reads text as tokens built over UTF-8 bytes or Unicode characters, and it processes Devanagari, Urdu or Chinese through the same machinery as English.

What does exist is an asymmetry, in three places. The training data are mostly English. The tokenisers, trained mostly on English, cut Indic text into more tokens for the same content: every Devanagari letter is already three bytes in UTF-8, so the same sentence costs more to process and fills the context window sooner. And the toolchains beneath the models keep their keywords, libraries and symbol alphabets in ASCII (Chapters 18 and 21). None of the three is a law to surrender to. Each has a technical remedy: more native text, tokenisers with script-aware merges and, at the toolchain, Romenagri, which uses ASCII-7 as a carrier while the script's identity stays reversible.

Romenagri's own measurements are recorded in [project-ilm/romenagri](#) and should be quoted as they stand. Written in Romenagri, the test corpus takes 1.76 times fewer bytes, and its alphabet shrinks from 51 symbols to 24, with nothing lost. In the recorded tokeniser sweep, however, Romenagri did not reduce the number of tokens at matched numbers of merges: the ratio ran from 0.62 to 1.0. The three- to five-fold token gain predicted in the author's tokenisation paper therefore stands as a falsifiable prediction still to be shown, and his runbook asks for exactly this result to be reported, not hidden.

Who owns the leading AI engines, and does ownership decide which languages and nations are sovereign?

Ownership of the leading models changes by the quarter and is not an engineering property; summaries of it, including AI-generated ones, need checking against primary sources before they are repeated. What owners decide is what their own tools do. They do not decide what a community can build. In the technical sense of Chapter 21, sovereignty belongs to whoever can rebuild, verify, modify and run the stack from source. The estate is built to that standard and to

be independent of any one model: the cyclers run with any AI, Romenagri, ILM and PANINI are open and run on a laptop, and open-weight models can be run locally (Chapter 24). A community that holds its scripts' path through the toolchain, its tokenisers and its corpora does not depend on who owns the most popular engine.

What is the state of the art for Indic and Chinese text in AI?

The field moves monthly, so this answer gives the mechanisms rather than model names and scores, which should be taken from published model cards and papers. Chinese and the Indic scripts pose different problems. A Chinese character is a single code point, three bytes in UTF-8, and usually carries a morpheme, so the question is how large a vocabulary of characters and words a tokeniser can afford. An Indic akshara is a cluster of several code points (Chapter 18), so the question is whether the tokeniser keeps clusters and morphemes whole. The approaches in use are larger vocabularies with merges learned from native text, pre-tokenisation that respects grapheme clusters, byte-level and character-level models, more native training data, and evaluation written in the languages themselves rather than translated from English. Reversible transliteration such as Romenagri is a further route, with the measured gains and limits given above.

What is a software factory, and what is the estate's?

The idea is older than most of the tools that now revive it. Douglas McIlroy proposed mass-produced software components at the NATO software engineering conference of 1968, and Hitachi organised software production as a factory in Japan in 1969; Cusumano's *Japan's Software Factories* (1991) is the history. The core is a universal set of software components that can be synthesised architecturally into a working product, each component keeping its provenance, which in present terms means its characterisation and its contracts. University courses and industry writing have taken the idea up again, for example themodernsoftware.dev and [factory.com's account](https://factory.com).

The estate builds its own version. **FAKIR** is the component registry: its categories follow the international classifications of the UN and the ILO, ISIC for economic activity, ISCO for occupations and ISCED for fields of education and learning, so that every lawful human economic activity has a place. **AAB** is the factory process: a paint program for systems, from painted intent to verified, sealed artifact. In the author's words, AAB and FAKIR together constitute the working definition of AGI used here. What most current work leaves out is the provenance of the components; here the registry records identity, provenance records ancestry and transformation, and evidence establishes what may be claimed. A release is quality-gated; beta versions are open on GitHub. The landing page at zistgah.org/fakir shows the component categories as a cloud near the ceiling of the dome, some 11,000 of them by his count, with a quadrotor to fly through them; passing through a component lets you start building it and send the work back as a pull request. The process has its own page at zistgah.org/aab.

How should an AI's confident answer be judged?

By what it has run, not by how certain it sounds. **COPA**, the Cost of Perceived Authority (doi:10.5281/zenodo.21782217), is the estate's protocol for the pattern: a black-box behavioural method for inferring the effective operational objective of a deployed AI system. Its metric, the Authority Projection Index, is the number of explanatory tokens before the first empirical action divided by the total number of tokens. COPA states its hypothesis together with a fair null and does not treat a single reading as proof of intent. What it establishes regardless of cause is the cost of assertion before verification: automation bias, overtrust, anchoring, delayed checking and reduced scrutiny. The target is calibrated trust, and the working rule follows from it: ask the AI to run the thing before you accept its explanation.

Can I point any AI at Romenagri and have it check the claims?

Yes, provided the AI can execute code: Claude with code execution, Claude Code, Codex CLI, Gemini CLI or Aider on your own machine (Chapter 24). Give it this:

Clone <https://github.com/hindawiai/chintamani>. Read [Romenagri/](#) and [Hindawi/hindrv/hincc](#). Build Romenagri with `make`, then the guru shaili and the driver, installing into a folder inside the clone. Run a few Devanagari words through `uni2acii | acii2cf` and back through `xmn2acii | acii2uni`, and compare. Compile [Hindawi/samples/HindiC.uhin](#)

with `hinc` and run it. Then build it with `-g` and show the names in `nm` and in GDB. Show every command and its output, and report nothing you did not run.

Chapter 15 gives the same steps as commands, and Chapters 19 and 21 show what they print.

29. The roadmap, for every collaborator and every domain

In plain words. The plan to finish the releases and open the work to everyone: what happens when, which lane each kind of person can join, and every field of work, from farming to medicine to software.

The plan to complete the releases and to open the estate to collaborators is kept as data and run as a GitHub project: one board, an epic for each phase, and an issue for every piece of work, each with its finding, its task, the exact gate that proves it done, the evidence to attach, and what it depends on. The same data gives every collaborator a lane and every domain a lane, so anyone can find where they fit and what is open to them. The roadmap page filters it all by phase, by collaborator lane, by domain family and by owner, at zistgah.org/roadmap.




Figure 12. The plan's phases; the lanes run alongside from P3.

Phases

Phase	Window	Goal	Issues
P0 Baseline	Weekend of 3 to 4 October 2026	One measured inventory, the board, tokens with the right scopes, and the list of rulings.	4
P1 Close the releases	10 to 25 October 2026	Everything built is released, and every defect found by execution is fixed.	19
P2 Licences and DOIs	24 October to 8 November 2026	Every repository carries a licence file and correct metadata; minting covers the estate by family.	5
P3 Collaboration infrastructure	31 October to 15 November 2026	A newcomer can go from Rahnuma to a merged pull request in one sitting.	7
P4 Reference lab and courses	November to December 2026	The reference lab runs from launch scripts, with GATE-aligned modules and a pilot.	5
P5 Papers and IPR	Rolling, from October 2026	Disclosure ledger first, then preprints of the four papers with results.	7
P6 Monetisation	From November 2026	Two offers defined and priced, on the dossier standard.	3
Lanes: every collaborator and every domain	Continuous, from October 2026	Every collaborator type and every domain family has a visible lane on the board, with an entry point, open work and a way to propose more.	62

Collaborator lanes

Lane	Start here	Contribute	Issues
School students, 10 to 17	the Rahnuma view for school students; CHAKRA, PANINI's studio, GENIE and the FAKIR dome	learning paths, translations of plain-words summaries, questions	1
Undergraduates: BE, BTech, BSc, BS	the GATE view; Running every component	good first issues, component fixes, capstones seeded from FAKIR	52
Postgraduates and doctoral researchers	the researcher views; Appendices C and D	theses and papers on the open questions, measured results, reviews	52
Teachers, lecturers and lab builders	the teacher view; the reference lab	course modules, lab sheets, pilots, reviews of plain-words text	9
Software engineers	the engineer view; the compilers chapter	code, tests, gates, tooling, pull requests	22
Hardware, embedded and chip engineers	Below the software; zamin, pratik_core_mvp, jugaad28	SPICE, layouts, firmware, device control, ESP32 work	8
Linguists, translators and language workers	the linguist view; writing systems and Pāṇini	language tables, corpora, transliteration rules, measurements	7
Researchers in AI, quantum and cognition	the researcher view; PANINIQ; PEDLER	models, experiments, validation, papers	7
Researchers in the humanities and social sciences	the humanities view; Appendix C	studies of agency, language, knowledge and community	4
Clinicians, health workers and medical researchers	VIKRAM's health section; CEMb and QNBCI	clinical questions, partnerships, review, health-domain tasks	3
Community workers, NGOs and village institutions	GramSheel and Project VIKRAM	local needs, pilots, language and field knowledge	30
Farmers, artisans and local enterprises	FAKIR's domain lanes; VIKRAM's livelihood and agriculture sections	domain knowledge, problems worth solving, field testing	24
Artists, writers, filmmakers and musicians	the maker view; the cyclers	works made with the cyclers, recorded as method	3
Makers, roboticists and manufacturers	PANINIphy, pench, transeg	physical realizations, parts, embodiment	6
Policy makers, administrators and public bodies	Appendices C and E; sovereignty in Chapter 19	policy questions, adoption, public infrastructure	5
Lawyers, IPR and standards professionals	provenance and licences; the disclosure ledger	licence review, IPR clearance, standards work	14
Industry partners and companies	Appendix E's adoption path; the component cards	integration, pilots, paid offers, support	9
Funders, sponsors and investors	this roadmap and its measures	funding of lanes, labs and pilots	4
Institutions: colleges, universities and labs	the reference lab and the GATE modules	pilots, courses, joint research	13
AI agents working under a human	any issue's URL; Chapter 22, how to use any AI	issue work through pull requests, with gate output attached	65

Domain families

From FAKIR's lattice: zistgah/fakir data/lattice.json (ISIC Rev.4, ISCO-08, ISCED-F 2013), 1,505 nodes.

Family	Name	Domains within	Issues
ISIC A	Agriculture, forestry and fishing	54	1
ISIC B	Mining and quarrying	29	1
ISIC C	Manufacturing	232	4
ISIC D	Electricity, gas, steam and air conditioning supply	7	1
ISIC E	Water supply; sewerage, waste management and remediation activities	18	1
ISIC F	Construction	22	1
ISIC G	Wholesale and retail trade; repair of motor vehicles and motorcycles	66	1
ISIC H	Transportation and storage	36	1
ISIC I	Accommodation and food service activities	15	1
ISIC J	Information and communication	42	22
ISIC K	Financial and insurance activities	31	1
ISIC L	Real estate activities	5	1
ISIC M	Professional, scientific and technical activities	35	11
ISIC N	Administrative and support service activities	51	1
ISIC O	Public administration and defence; compulsory social security	11	1
ISIC P	Education	14	7
ISIC Q	Human health and social work activities	21	1
ISIC R	Arts, entertainment and recreation	19	1
ISIC S	Other service activities	26	1
ISIC T	Activities of households as employers; undifferentiated goods- and services-producing activities of households for own use	8	1
ISIC U	Activities of extraterritorial organizations and bodies	3	1
ISCO 1	Managers	46	1
ISCO 2	Professionals	125	1
ISCO 3	Technicians and Associate Professionals	109	1

Family	Name	Domains within	Issues
ISCO 4	Clerical Support Workers	41	1
ISCO 5	Service and Sales Workers	57	1
ISCO 6	Skilled Agricultural, Forestry and Fishery Workers	30	1
ISCO 7	Craft and Related Trades Workers	85	1
ISCO 8	Plant and Machine Operators, and Assemblers	57	1
ISCO 9	Elementary Occupations	50	1
ISCO 0	Armed Forces Occupations	9	1
ISCED 00	Generic programmes and qualifications	6	1
ISCED 01	Education	5	4
ISCED 02	Arts and humanities	13	5
ISCED 03	Social sciences, journalism and information	8	5
ISCED 04	Business, administration and law	10	7
ISCED 05	Natural sciences, mathematics and statistics	13	4
ISCED 06	Information and Communication Technologies (ICTs)	4	16
ISCED 07	Engineering, manufacturing and construction	15	4
ISCED 08	Agriculture, forestry, fisheries and veterinary	9	1
ISCED 09	Health and welfare	12	1
ISCED 10	Services	14	1

The AGI layers are L0 Hardware Definition, L1 Verification, L2 Synthesis/Compilation, L3 Firmware/Embedded, L4 Systems, L5 Parallel, L6 Distributed/Cloud, L7 AI/Deep Learning, L8 Robotics/Embodiment, L9 AGI/ASI. Every domain node, crossed with every layer and every language in ILM's registry, is a possible seed:

$$|S| = |D| \times |L| \times |\Lambda| = 1,505 \times 10 \times 7,867 = 118,398,350.$$

How a collaborator comes in




Figure 13. How any collaborator, in any domain, goes from the guide to a sealed contribution.

Four issue forms in [zistgah/governance](#) are the ways in: **Join as a collaborator**, with your lane, domains and languages; **Propose a domain task**, seeded the way FAKIR seeds it, with a domain code, an AGI layer, a language, an intent and a gate; **Report a finding**, with the commands run and what they printed; and **Request a ruling**, for the questions only the author can settle. Every issue says which owner class takes it: a Claude Code session, ChatGPT, Gemini or Fable as workers, or the author himself for push, Pages, mint, merges and rulings.

The work flows the same way for everyone. A worker, a person or an AI, opens a pull request with the gate's output attached; the author merges what is green; anything pushed or minted carries its Candor receipt and goes through seal, clear, attest and mint. Each Saturday the status digest shows every issue as closed, ready, or blocked and by what.

30. What the work would cost: a model-based estimate

In plain words. A measured estimate of how much work the estate holds and what it would cost to rebuild, with and without AI.

How much work does the estate represent, and what would it cost a team to rebuild it, with AI in the process or without? The estimate kit in this guide's repository, under `estimate/`, answers with a recognised model rather than an opinion. It counts every tracked file of every repository once, leaving out vendored, generated, minified and third-party material, measures cyclomatic complexity, and applies COCOMO II.2000, then prices the effort with published salary data and the AI scenario with published token prices and the two controlled trials that bound AI's effect on developer speed. Every rate names its source in `estimate/rates.json`.

Demonstration on 38 repositories cloned for Rahnuma, about a fifth of the estate, most without their full history:

Measure	Value
Unique source lines, each file counted once	243,719
Lines repeated across repositories, counted once	341,227
Words of prose	368,689
Functions measured; above complexity 10; above 20	9,288; 546; 191

Case	Effort, person-months	Schedule, months	Cost, India	Cost, United States
low	701	27.1	₹12.9 crore	\$7.5 million
likely	1,239	35.3	₹41.0 crore	\$18.6 million
high	2,279	47.3	₹130.7 crore	\$49.9 million

AI in the process	Time ratio	Human effort, person-months	AI usage
low	0.44	554	\$140
likely	0.75	944	\$578
high	1.19	1,498	\$5,861

The AI usage costs a few hundred to a few thousand dollars at list prices; the outcome turns on the time ratio, which the trials put between 0.44 and 1.19.

To run it on the whole estate: `bash estimate/estimate.sh --clone`, which clones every repository with full history, or `bash estimate/estimate.sh --root <folder of clones>`. The report also sets the model against the estate as it was actually built, from the distinct days in its git history, with hours per day and AI spend that the author sets in `rates.json`.

COCOMO prices the construction of what is written down. It does not price the research, the invention and the judgement behind it, so every figure here is a floor, not a valuation.

Part VI: Status

31. This release and the next

In plain words. What this edition contains, what changed in it, and what comes next.

This release

Rahnuma 2.0.0 is this guide. It opens with four steps that find each reader's part of it, from a school student to an investor; offers a PDF edition for every kind of reader and every field of study, links school and exam syllabi to the chapters that teach the same ideas on the reader's own device, proposes names for the projections and the procedure for raising a component into Humanesque, publishes the estimate of the work, and marks every object with the author's review. Its policy for versions: when a release changes how the guide is entered or organised, it is a new major version; the previous major version moves, unchanged, to `v<N>/` beside it, and is linked from the cover [Version 1](#), the edition for advanced readers and senior professionals, stays there as it was. Rahnuma 1.3.0, dated 3 October 2026, was version 1's last release. It adds fourteen diagrams drawn from the recorded relations, an appendix of formal definitions, and the roadmap, with a lane for every collaborator and every domain family (Chapter 29). Version 1.2.0, of the same date, added appendices for researchers in the humanities and social sciences, in science, technology, engineering and mathematics, and for practitioners, an appendix on reading the architecture without flattening it, and the questions behind GramSheel and VIKRAM as the author put them on Gandhi Jayanti. Version 1.1.0, of the same date, supersedes both printings of 1.0.0, one of 53 pages and one of 67, which carried the same version number. It adds the dependency graph of the estate (Chapter 7), the AyeAI Triad, GramSheel and Project VIKRAM, and Humanesque with its projections and their names, each set apart (Chapter 10), the dated entries for GramSheel's beginning in 1993 and VIKRAM's launch in 2020, the entries for Kitab and Research Kundali, and the rule that a record not yet retrieved is not thereby false. It also runs every component of the estate on a clean machine and reports each one as it ran, with the set-up a README leaves out and the reason when a component does not pass (Chapter 15); it gives every chapter a summary in plain words; and on the site it lets a reader choose a background and an age, and shows a reading path, the components to try first, the depth of text, a larger type and a light or dark page. Version 1.0.0 is deposited at [doi:10.5281/zenodo.23062059](https://doi.org/10.5281/zenodo.23062059). Like 1.0.0, it is this guide, in three forms: the site, the PDF, and one Markdown file for reading or for giving to an AI. It comes with its examples, its build and its gate. Every example shown as executed was executed, and `bash ops/verify.sh` runs them again. That includes Hindawi's C shaili, built with Romenagri from its retrieved sources and run through GCC and GDB in Chapter 21.

The estate today

Chapter 12 lists every public repository and every DOI a repository records about itself. Most repositories are pushed and public without a DOI of their own. One case is worth naming because it is often asked about: *zistgah/jyotish*, the offline panchang, is pushed and its site is live, and it has not been minted. A repository already pushed can be minted later with the estate's seeder; minting is the author's decision, taken repository by repository.

Next

The next release of this package carries the lab launch scripts: one script per lab, the software dockerised and retrieved from the existing repositories, not rebuilt (Chapter 16). The GATE-level courses in computer science and in robotics and automation follow, aligned with the reference lab (Chapter 26).

The consolidated release

This guide marks a plateau: the estate consolidated in one place, with the mathematics written out and the entry points documented. After it, papers from the work are to appear over the coming weeks and months, and the consolidated body over one to two years. In the author's view, ASI and his laboratory fall within that horizon.

Taking part

Open an issue or a pull request on any repository, cite the DOIs, fork freely, and build in your own language.

Colophon

Written and compiled by Abhishek Choudhary, AyeAI, with AI assistance under verification: every example was executed and every statement about the repositories was retrieved from them. ORCID [0009-0002-0684-8320](https://orcid.org/0009-0002-0684-8320). Some address him as Dr; he holds no doctorate. The author is immaterial; the work is what matters.

Typeset in Tiro Devanagari Sanskrit, Noto Serif and its script families, and Noto Sans Mono, all under the SIL Open Font Licence; mathematics rendered by MathJax. Text © 1993–2026 Abhishek Choudhary. All rights reserved. AyeAI. Licensed CC BY-SA 4.0; scripts GPL-3.0-or-later.

Appendices

A. The estate's terms, and the nearest common terms

In plain words. A table matching the special terms used here with the closest common terms.

This table is retrieved from `GLOSSARY.md` in `zistgah/paniniq` at commit `82a5cf2`. The terms are the author's and are kept; the right-hand columns exist so that a reader, or another AI, who knows the market vocabulary can find their footing, and where the two differ the difference is stated.

His term	What it means here	Nearest common terms	Where they differ
Cycler	An AI-agnostic protocol for human and AI working together: intent, context, a meaningful prompt, any AI answers, the human inspects, the artifacts and responses are kept under configuration management, the next prompt follows, until a final artifact the human authored, carrying its intention. The family is named by what it produces: print (matba), video (khwab), audio (awaz), immersive (tilasm), embodied (pench), memory (yadein), research (genie).	agent loop, agentic workflow, human-in-the-loop harness, prompt chain	A cycler is never the author and never an AI provider. Every prompt must create, verify, execute, measure, falsify or integrate an artifact, and it stays useful if every commercial provider vanishes.
Agent, harness	His placement: the cyclers and genie are the agents and harnesses of Act I, AGI completeness.	agent, agent harness, evaluation harness	The market term usually names the software; his names the protocol the software keeps.
Configuration management (step 3 of every cycle)	Every prompt, response, human decision and artifact version is identified, hashed and appended; nothing is overwritten; the current state is derived from the record.	configuration management (IEEE 828, ISO 10007), provenance tracking, experiment tracking, audit log	Responses are configuration items in their own right, not chat history.
AAB (آب, water)	The paint program for systems: the human paints actors, components, stores, gates, interfaces and environments, wires the flows, and the painting is the manifest every prompt carries. Quests are the eight VGC stages; a node glows verified only with a recorded oracle. Process cyclers, such as building one language layer, run here.	low-code system designer, spec-driven development, model-based systems engineering canvas	Gamification never outruns grounding: no oracle, no glow.
VGC	Verification-Gated Human-AI Co-Development: nothing is accepted without an oracle.	eval-gated development, CI gating, test-driven development with AI	Acceptance belongs to an oracle, never to the human or the AI.
Oracle	The check that decides acceptance: a test suite, a second implementation, a measurement or a review.	test oracle, acceptance test, conformance suite	
Quest	One VGC stage, completed only with typed evidence.	stage gate, milestone with acceptance criteria	
PANINI	His language for prompt cycles and his construct compiler: common semantics plus decorators, with many front ends and realization backends.	domain-specific language, orchestration language, compiler toolchain	The construct is primary; syntaxes are views of it.

His term	What it means here	Nearest common terms	Where they differ
Front end, middleware, realization backend	His layering of 27 Sep 2026: ILM and the language parsers (tajziya) are the front end; PANINI's own front ends and its fine-tuned language are the middleware; the realization backends are where programs meet a substrate, Panini Q among them.	compiler front end, intermediate representation and passes, compiler back end or target	
ILM	Integrative Linguistic Multiscript: one language core for all languages, with script, language and standard kept apart.	multilingual NLP, script-agnostic text processing, transliteration, Unicode normalisation	Script, language and standard are separate axes, never conflated.
CEM (CEM ^b , CEM ^s , Eco-CEM)	Cognitive enablement modules, in bioform, synthetiform and ecological families; his Act II, with the accelerators.	AI accelerators, neuromorphic hardware, cognitive architectures, brain-computer interfaces	
The three Acts	His note “3 Act ASI $\wedge$ Panini”: Act I, AGI complete; Act II, post-AGI to ASI; Act III, post-ASI. PANINI runs through all three.	AGI, superintelligence (ASI) roadmaps	Placement says when a piece matters most, not when work on it starts.
Candor, Tok DOI, Misty DOI	Signed intent receipts; atomic timestamps; DOI minting, in the order seal, clear, attest, mint.	in-toto and SLSA attestations, OpenTimestamps, Zenodo DOI deposit	
Reference lab	His laboratory pattern for doing most of the research behind high-end AI, robotics and automation papers at modest cost, alongside the courses built on it.	AI research lab, robotics and automation lab	Public name pending his choice.

B. Reading the architecture without flattening it

In plain words. How to read the architecture without squashing it into a straight line: what an arrow can mean, and the loops that run at different scales.

The architecture is dense because what it addresses is dense: language, cognition, knowledge, participation, embodiment, verification, provenance, community capability, identity and continuity. This appendix does not simplify it by removing parts. It shows how to read it so that its structure stays open to questioning. Statements marked † follow the author's working notes and write-ups as supplied on 29 September and 3 October 2026.

A graph read one route at a time

Any sequence printed on a page, A then B then C, is one route through the graph of Chapter 7, chosen because a sequence is easier to read. Before trusting an arrow, ask which of seven things it means:

The arrow is	Example in this guide
temporal: one thing came before another	the dated record; PEDLER (2001) before Romenagri (2003)
causal: one thing brings about another	an observation changing a knowledge state
a dependency: one cannot be built or understood without the other	the cyclers are written in the PANINI language
an interface: two things meet at a defined boundary	Romenagri names crossing into the host toolchain
a composition: things used together without merging	Mez composing GENIE, Kitab and the cyclers
an expository order: the order of explanation only	the reading paths of the reader views
an interpretive mapping: a reader's bridge between vocabularies	the nearest analogue in each term entry

These are not interchangeable. The estate itself keeps three of its orderings apart: the Acts are temporal, the projections are cultural, and the edges of the graph are architectural; none of them is a ladder of maturity.

Loops at different scales

The architecture has several loops, and they run at different scales; collapsing them into one “AI feedback loop” loses what each is for




Figure 14. Loops at four scales: PEDLER's, GENIE's, the realization spine's, and continuity as the write-ups read it (†).

The cognitive loop, PEDLER's. State, observation, counter-observation, evaluation, inclination, selection, intent, Act, execution, event, and a new state. Inclination is a signed directional gradient, not yet an intention; the Act is executable intent; Natural Justice gates the passage from intent to Act.

The research loop, GENIE's. Create, verify, execute, measure, falsify, integrate, with nine epistemic tags on every statement and a gate that a simulation cannot pass in place of an experiment.

The realization loop, PANINI's. Language, semantic IR, realization requirement, domain resolution, domain IR, then a common layer of geometry, trajectory, animation, observation, verification and an ArtifactGraph. The realization repositories keep desired, predicted, observed and validated results apart, and compilation success, execution success, target satisfaction and validation apart; † the notes extend the separation to desired, designed, predicted, planned, realized, executed, observed, verified and validated.

The continuity loop. † The write-ups read a fourth loop at civilisational scale: identity, memory, provenance, transformation, a new substrate or embodiment, continuity. Its components in the record are yadein, which records, TransEg, which continues a constituted identity, and the provenance systems that keep what happened.

In every loop the gates do work in both directions. Verification is not only a last inspection: a failed gate stops the next step. Provenance is not only storage: it fixes what can legitimately be claimed about an artifact. Observation is not only telemetry: it changes the state from which the next action is chosen. VGC states the rule plainly: the gate is a verified artifact, and an AI's report that something passed is not evidence that it passed.

One component, several cycles

A component is not exhausted by the place it appears in one diagram. ILM belongs to the linguistic line, yet language shapes how knowledge and intent are represented. FAKIR is a knowledge kernel, yet the knowledge it supplies shapes what is done. PEDLER is a cognitive architecture, yet its Act, Inclination and Natural Justice are also the primitives of the constitutional corpus. VGC concerns verification, yet verification decides what may become accepted knowledge. GramSheel is a foundation for villages, yet it touches education, livelihood, health and technology at once. The graph records relations of dependence, interaction and constraint, not modules stacked in a tower.

Names are interfaces, and analogies are doorways

The density of names cannot be removed without removing distinctions; each name denotes a different object. What helps is a functional handle at first meeting, which the term-by-term chapter and Appendix A give, and an analogy used as a doorway rather than a definition. † In the write-ups' phrase, the analogy opens the door; the formal definition decides what is inside the room.

C. For researchers in the humanities and social sciences

In plain words. A guide for researchers in the humanities and social sciences: the questions each field can ask of this work, and where to find and test the answers.

Each discipline below is given a question the architecture makes concrete, where to look, what can be examined or run, and what must be kept distinct. These are research questions, not claims that any discipline has validated the architecture. The technical components specify relations; the disciplines bring the questions through which those relations can be investigated. Neither vocabulary replaces the other: PEDLER is a computational object that psychologists and philosophers can interrogate as a model of agency, not psychology; GramSheel is an infrastructure that sociologists can interrogate as a model of capability distribution, not a development programme; ILM is a computational object that linguists can interrogate as a model of multilingual representation, not linguistics.

Political science, law and governance

The question. Who holds agency and computational authority, and what counts as admissible evidence when machines take part in decisions? **Where to look.** The Proclamation of Individual Equity and recursive sovereignty (Chapter 4); Humanesque and its projections (Chapter 10); PATAL, GramSheel and Project VIKRAM; Natural Justice as a gate in PEDLER; VGC and the provenance order seal, clear, attest, mint (Chapter 25). **What can be examined.** The text of PoIE (zistgah/poie, doi:10.5281/zenodo.21397274); the executable contracts that gate every repository (zistgah/governance), each a set of clauses a machine checks; the Candor receipts, which record signed intent as in-toto statements. **Keep distinct.** PATAL is not democracy, and Humanesque is not political theory; Natural Justice in PEDLER is a gate on the passage from intent to Act, not a general ethics.

Sociology and development studies

The question. Can technological capability be distributed through communities rather than concentrated in distant institutions, and can villages produce knowledge and technology rather than only consume it? **Where to look.** GramSheel and its principles, health, education, justice, livelihood and peace; Project VIKRAM and its thirteen sections; FAKIR's lattice of every lawful economic activity under ISIC, ISCO and ISCED; the reference lab. **What can be examined.** VIKRAM's journal and its section pages (pvjournal.github.io), launched on 3 May 2020; FAKIR's explorer, where each point seeds a task for a domain, a language and a layer; the open contribution model of the repositories, through issues and pull requests. **Keep distinct.** GramSheel is the foundation and its principles; VIKRAM is the technical platform; neither is a market segment or a programme of delivery to passive recipients.

Psychology and cognitive science

The question. How does an agent move from a state and an inclination to an action, and how does a machine's projected authority affect what people accept? **Where to look.** PEDLER's loop and its primitives (Appendix B); COPA and its Authority Projection Index (Chapter 8); the Quantum Neuromorphic BCI programme, with its active research and its directions needing clinical partners. **What can be examined.** The PEDLER engine in PANINIq, with its tests (Chapter 23); COPA's index, which can be computed on any transcript as explanatory tokens before the first empirical action over all tokens. **Keep distinct.** Inclination is not intention, and intention is not an Act; COPA measures projected authority, not accuracy or intent, and a single reading proves nothing about intent.

Linguistics

The question. How are language, script and meaning preserved when linguistic structures become computational representations, and can any language travel in any script? **Where to look.** Writing systems and reversible transliteration (Chapter 19); Pāṇini's grammar as a formal system (Chapter 20); Hindawi and the symbol bridge (Chapter 21); ILM and tajziya. **What can be examined.** Every Romenagri result in this guide can be rerun on your own corpus with the commands of Chapter 15: the Brahmi hub, the round trip, the Urdu projection with and without tashkil. ILM's registry lists 7,867 languages and 226 scripts, and its keyword standards number 13,668, every one a valid

ASCII identifier. **Open.** The share of Urdu words whose vowels must come from the language layer, which the author estimates at 15 to 20 per cent, awaits measurement, and the reverse filter into Perso-Arabic script is a first cut. **Keep distinct.** ILM is not linguistics and Romenagri is not romanisation; script, language and standard are three axes.

Anthropology and cultural studies

The question. How do identity, situated knowledge and cultural plurality survive technological transformation?

Where to look. The three projections, Kaivalyik, Zistgah and Cosmopolis, and the names table (Chapter 10); Jyotish as the Kaivalyik projection of CHAKRA; the second pillar of linguistic equity, any language in any script. **What can be examined.** What the Devanagari hub keeps and loses when Urdu passes through it, printed in Chapter 21; how a component keeps one stem and takes a culturally sensitive name in each projection. **Keep distinct.** A projection is a complete traversal of the whole architecture, not a translation of a master copy; Zistgah's Persian and Islamicate dimension stays explicit.

Philosophy

The question. What constitutes knowledge, action, identity and continuity when the substrate changes? **Where to look.** PEDLER's primitives; VGC; the provenance systems; TransEg; Synthematic Pragmatic Realism; the rule that a record not yet retrieved is not thereby false. **What can be examined.** An applied epistemology already in use: GENIE's nine epistemic tags (definition, axiom, assumption, derivation, conjecture, hypothesis, empirical, open, failed), the four provenance states of a statement (retrieved, inferred, proposed, unresolved), and the status words of this guide (tested, executed, written, mocked, planned). † The write-ups put the underlying problem as an unconstrained system making an epistemic state transition without an admissible reason. **Keep distinct.** SPR is a stated primitive of the architecture; generic metaphysics must not be substituted for it.

Education

The question. Does technology hand out answers, or does it increase human capability? **Where to look.** The reference lab and the GATE chapter (Chapters 16 and 26); the reader views of this guide; the cyclers as recorded, reproducible method; VIDYA, the bridge to AyeAM. **What can be examined.** The component cards of Chapter 15 as laboratory sheets, each with what to type and what you should see. † The write-ups also name AYE Learn and AYE Sum among the educational work.

History and the history of computing

The question. What was built, when, and how can priority be checked without argument? **Where to look.** The dated record (Chapter 5) and its primary sources: the SourceForge and Savannah projects, the announcement of 19 December 2005, and the Zenodo records. For context rather than lineage, Gandhi's *Constructive Programme* (1941) placed village industries, education and the provincial and national languages inside the work of self-rule; Chapter 10 sets out where the estate's questions meet it.

Economics and political economy

The question. Can advanced capability become locally reproducible and sustainable rather than externally concentrated? **Where to look.** FAKIR's ISIC lattice as a map of economic activity; the open licences of the repositories (Chapter 25); the reference lab's cost target. † The notes treat monetisation, through products, services, courses, labs and licensing, as the mechanism that sustains the open work.

D. For researchers in science, technology, engineering and mathematics

In plain words. A guide for researchers in science, technology, engineering and mathematics: what can be read, run and tested now, and what is still open.

Each field below lists what in the estate can be read, run and tested now, and what is open. Status words keep their meaning from Chapter 1: tested, executed, written, mocked, planned.

Computer science

Programming languages and compilers. Hindawi's composed transducers over unmodified host toolchains; the Romenagri symbol bridge through ELF, DWARF and GDB; the PANINI construct model, with 39 constructs in 27 languages and 216 decorators across 8 host languages; the bootstrap and its self-hosting fixed point (Chapters 21 and 6). Executed here. **Open:** `std2hin` doubles every underscore, so vowel-initial names do not round-trip back from C; the guru shaili emits no `#line` directives; the YACC shaili awaits a fresh conformance run.

Systems. The hinlin kernel tree and the mass conversion of real C sources, including kernel trace code, into Hindi and back (hindawiai/chintamani, `Romenagri/mass_hindawi`).

AI and machine learning. The cyclers and GENIE as AI-agnostic, human-inspected method (Chapter 8); Romenagri as a tokenisation substrate, where bytes fall 1.76 times and the alphabet from 51 to 24 symbols without loss, but token counts did not fall at matched merges, so the predicted gain is still to be shown; COPA's Authority Projection Index as a measurable property of a model's behaviour.

Software engineering. Executable contracts and gates in every repository; the provenance chain, Tok DOI, spiguard, Candor and Misty DOI; VGC; the component cards of Chapter 15, each the result of a real run.

Data. ILM's registry of 7,867 languages and 226 scripts; 13,668 keyword standards; the corpora in the Romenagri tree.

Electrical and electronic engineering

The open chip-design flow and the 28 nm tape-out catalogue, jugaad28, with its evidence-gated descriptors (Chapter 22); Zamin's ternary cells in SPICE, whose testbench needs attention on ngspice 42 (Chapter 15); the PRATIK kernel, whose CPU build passes and whose CUDA build needs an NVIDIA toolchain; device control on ESP32, zasab, built but not yet pushed.

Physics

QEDLER, the event-hypergraph research programme for physics, self-declared as a framework with open problems; oscillator Ising machines and Kuramoto dynamics in PANINIq, with their Lyapunov function (Chapter 23); superconducting qubits and the physics of the cold (Chapter 22).

Mathematics

Coding theory for reversible transliteration: Kraft and McMillan, and Sardinas and Patterson (Chapter 19); formal languages and rewriting in Pāṇini's grammar, with pratyāhāras as interval classes over an ordered list (Chapter 20); graph theory in the typed hypergraph and in Max-Cut; dynamical systems in the order parameter and in the oscillators' energy; the formal tuples of the CEM kernel, $\mathcal{E} = (\mathcal{E}, \mathcal{C}, \mathcal{I}, \mathcal{W})$, and of the AyeAI Triad.

Biology and medicine

PANINIb, the biological realization arm, which lowers intent to sequence IR and claims no wet-laboratory result; CEMb, operated through Interglial Healthcare; the Quantum Neuromorphic BCI programme, with brain-to-text, cognitive prosthetics and intelligent rehabilitation as directions needing clinical collaboration; VIKRAM's Health section and its

tools for an accessible virtual hospital; the earlier HMSEI wearable and the Dr Rho telepresence platform.

Cognitive science and neuroscience

PEDLER's two branches, through LVF and MLCNE to the cognitive enablement modules, and through NI2A2 to ANGEL; its primitives; the brain-computer interface work of Act III.

Astronomy and the earth and space sciences

CHAKRA, the offline observatory and calendars, whose JavaScript and C99 versions produce byte-identical output; Jyotish, derived from it; the habitats, with a lander embodiment on the Moon and a quadroter on Mars inside the shared dome.

Robotics and manufacturing

PANINIphy, the physical realization arm, which compiles physical intent and chooses parts by force, stroke and voltage; pench, the embodied cyler; † the modular self-reconfiguring robots of the notes.

E. For practitioners: teachers, policy makers, communities, makers and engineers

In plain words. A guide for teachers, policy makers, community workers, artists and engineers: where to start and how to take part.

Teachers and lab builders. Start from the reference lab, the GATE chapter and the component cards, which double as laboratory sheets: each says what to install, what to type and what a student should see (Chapters 16, 26 and 15). The reader views let each student follow a path suited to their age and background, and remember what they have read.

Policy makers and administrators. The questions that matter are in Chapter 10 and Appendix C: sovereignty in its testable technical sense, the ability to rebuild, verify, modify and run (Chapter 21); provenance that makes claims checkable; and infrastructure that leaves capability with communities. The executable contracts of the repositories show what verifiable governance of software looks like in practice.

Community workers and health practitioners. Project VIKRAM's sections cover education, health, livelihood, justice, environment and habitat, food and agriculture, and peace; its health section accepts submissions as issues. The Urdu edition and the Brahmi hub show work in people's own scripts; CHAKRA and Jyotish run offline.

Artists, writers and makers. The cyclers produce print, film, sound, immersive work, embodied work and the record; matba turns posters into a sealed book; GENIE writes and paints; PANINIPHY turns a physical idea into parts. Every cycle is recorded as method, so a piece of work can be repeated, shared and credited.

Engineers in industry. Adopt one component at a time: read its card, clone it, run it, read its `CONTRACT.md` and its gate, check its DOI and provenance, then integrate it through its interface. † The notes give the full adoption ladder: the canonical artifact, a one-page explanation, the nearest industry analogue and the exact difference, a working demonstration, the repository, the DOI and timestamp, the interface, a case study, and the commercial route.

Contributors. Fork, open an issue, send a pull request. Every repository's gate runs on your change; every claim must point to its evidence; the AI you use is welcome, and its report that something passed is not the evidence that it did.

F. Formal definitions

In plain words. The same ideas written as mathematics, for readers who want the formal version.

This appendix states in symbols what the chapters state in words. Where the record gives a formal object, the PEDLER six-tuple, the CEM kernel or the Triad's tuples, it is reproduced as recorded. Where the record gives only words, the formalisation is this guide's reading, and is said to be.

The estate as a typed graph

The estate is a typed graph

$$H = (V, E, \tau_V, \tau_E), \quad E \subseteq V \times V \times R,$$

where V is the set of components, R the set of relations (projects, script layer, realized by, front end, mounts and the rest), $\tau_V : V \rightarrow \mathcal{F}$ assigns each component its family, and τ_E each edge its relation. The core is described as a typed hypergraph, in which one edge may join several components at once: $E \subseteq 2^V \times R$. A traversal is a walk $v_0 \xrightarrow{\tau_1} v_1 \xrightarrow{\tau_2} \dots \xrightarrow{\tau_k} v_k$; every sequence printed in this guide, and every reading path of the reader views, is one.

In this guide's reading, a projection $p \in \{\text{Kaivalyik, Zistgah, Cosmopolis}\}$ is a map $\pi_p : H \rightarrow H_p$ that keeps every component's stem and every relation and changes names: $\text{stem}(\pi_p(v)) = \text{stem}(v)$ and $(u, v, r) \in E \Rightarrow (\pi_p(u), \pi_p(v), r) \in E_p$. The Acts are a separate, partial labelling $\alpha : V \rightarrow \{\text{I, I|II, II, II|III, III}\}$, with PANINI in every Act.

Romenagri as a code

Let Σ be the Devanagari alphabet and $\Gamma = \{\mathbf{a}, \dots, \mathbf{z}, _ \}$, a subset of the alphabet of C identifiers. Romenagri's compiler form is a map

$$T : W \rightarrow \Gamma^*, \quad W \subseteq \Sigma^*,$$

injective on well-formed text W , with an inverse T^{-1} on $T(W)$. Text outside W is first brought to its canonical form by $c : \Sigma^* \rightarrow W$, so that for every string

$$T^{-1}(T(s)) = c(s), \quad c(s) = s \text{ whenever } s \in W.$$

On the 984 words of the Hindi corpus, $T^{-1}T(w) = w$ for 983; the one exception is a misspelling in the corpus.

For a Brahmi script σ whose Unicode block follows the ISCII layout from base b_σ , the hub map is $h_\sigma(x) = x - b_\sigma + \mathbf{0x900}$ on the aligned letters, a bijection onto the corresponding Devanagari letters; `flatten_uni_dev` applies it by table. A word in script σ therefore has the identity $T(h_\sigma(w))$, and returns by $h_\sigma^{-1}(T^{-1}(\cdot))$.

Urdu meets the hub through $\varphi : \Sigma_{\text{ur}}^* \rightarrow \Sigma^*$, which is many-to-one on unmarked text: the fibre $\varphi^{-1}(w)$ over a hub word can hold several spoken readings, because the short vowels are not written. With tashkil, the restriction of φ to marked text distinguishes them: unmarked `ہندوی` maps to `हन्दवी`, and marked `ہندوی` to `हिन्दवी`. In the corpus, 1 word in 794 carries a vowel mark.

The symbol bridge

For every name n in a program, the toolchain sees $\text{sym}(n) = T(n) \in \Gamma^*$. Each stage t_k between source and debugger (compiler; assembler; linker; ELF writer; DWARF writer; debugger) carries identifiers over Γ unchanged: $t_k(\text{sym}(n)) = \text{sym}(n)$. Hence

$$T^{-1}(t_m \circ \dots \circ t_1(\text{sym}(n))) = c(n),$$

which is the name as written. Chapter 21 runs this for every name of a program.

The construct model

Let $\mathcal{C}$ be the constructs (39), $\mathcal{L}$ the languages (27) and $\mathcal{H}$ the host languages (8). Each language ℓ has a keyword map $\kappa_\ell : \mathcal{C} \rightarrow \mathcal{W}_\ell$, and each host h a realisation $\rho_h : \mathcal{C} \rightarrow \mathcal{W}_h$, partial because not every host has every construct. Translation from ℓ to h is

$$\tau_{\ell \rightarrow h} = \rho_h \circ \kappa_\ell^{-1},$$

defined when κ_ℓ is injective, which is the reversibility check ILM runs on every keyword table. Direct pairings of a language with a host would number $|\mathcal{L}| \cdot |\mathcal{H}| = 27 \times 8 = 216$ tables; the construct model needs $|\mathcal{L}| + |\mathcal{H}| = 35$.

PEDLER, as PANINIq implements it

PANINIq implements PEDLER as the six-tuple

$$P = (I, G, U, S, F, *),$$

with events $U \subset \mathbb{R}^3$, states S , a feature map $F : S \rightarrow \mathbb{R}^3$, and the resolution kernel

$$G(v, r) = e^{-\gamma \|F(v) - r\|}.$$

The inclination of a state is its historical prior plus the weighted match, $I(v) = \eta(v) + w G(v, r)$. The operator $*$ grows and prunes the state set: a new state is spawned when the relaxed network's order parameter falls below a novelty threshold, $R < \theta$. The Act is the transition $A(I, S_t) \rightarrow S_{t+1}$. In this guide's reading, Natural Justice is an admissibility predicate on the passage from intent to Act: a transition fires only if $\mathbf{NJ}(\text{intent}, S_t)$ holds.

Kernels and closures, as recorded

The record gives the CEM kernel and the AyeAI Triad's closure as tuples, and they are reproduced here as given; the expansions of their letters are not part of the retrieved record:

$$\mathcal{E} = (E, \mathcal{C}, \Pi, W), \quad \text{AyeAM} = \langle S, R, C \rangle, \quad \text{AyeAI} = \langle M, I, G \rangle, \quad \text{AyeCNSe} = \langle T, Ch, \Sigma \rangle.$$

Gates and evidence

A gate is a function $g : \mathcal{A} \rightarrow \{\text{pass}, \text{fail}, \text{unjudged}\}$ on artifacts, computed by running it. An issue is done exactly when

$$g(a) = \text{pass} \wedge \text{merged}(a) \wedge (\text{published}(a) \Rightarrow \text{receipt}(a)).$$

VGC's rule is that a worker's report $r(a)$ is not an argument of g : the gate is computed from the artifact, never read from the claim.

Provenance

A repository's manifest is the list $m = [(p_i, H(f_i))]_i$ of every tracked path with the SHA-256 of its file. Sealing commits the digest $d = H(m)$ to OpenTimestamps, whose calendars aggregate many digests in a Merkle tree anchored in a Bitcoin block; the proof for one digest among N is a path of length $\lceil \log_2 N \rceil$, and verifying it recomputes the root. The four systems are ordered by implication:

$$\text{mint}(a) \Rightarrow \text{attest}(a) \Rightarrow \text{clear}(a) \Rightarrow \text{seal}(a).$$

The seed space and the reading paths

FAKIR seeds a task from a domain node, an AGI layer and a language: $S = D \times L \times \Lambda$, with $|D| = 1,505$ nodes across ISIC, ISCO and ISCED, $|L| = 10$ layers and $|\Lambda| = 7,867$ languages, so

$$|S| = 1,505 \times 10 \times 7,867 = 118,398,350.$$

A cross-domain task is seeded from a set of nodes, a subset of D , rather than one.

A reader view is a sequence of chapters $(c_1, \dots, c_k)$, a walk on the chapters. A chapter of w words is estimated at $t(c) = \max(1, \text{round}(w/200))$ minutes, and a path at $\sum_j t(c_j)$.

G. Review register

In plain words. A list of every part of this guide, showing which parts the author has checked by hand.

Every chapter, figure, component card, syllabus module, edition and proposed name in this guide is an object with an identifier, and each carries a mark: reviewed by the author, with the date, or not yet reviewed. The site shows the mark beside each object. The author marks an object with `bash ops/review.sh <identifier> [note]`, and the gate fails if any object is missing from this register.

0 of 424 objects reviewed by the author.

Object	Kind	Mark
ch-01 How to use this guide	chapter	not yet reviewed
ch-02 The work, not the worker	chapter	not yet reviewed
ch-start Start here: find your way in	chapter	not yet reviewed
ch-hist Where the work comes from	chapter	not yet reviewed
ch-05 The record, with dates	chapter	not yet reviewed
ch-03 Architecture: core, projections, layers	chapter	not yet reviewed
ch-graph The dependency graph	chapter	not yet reviewed
ch-onto The ecosystem, term by term	chapter	not yet reviewed
ch-acts The three Acts	chapter	not yet reviewed
ch-families Three families, set apart: the AyeAI Triad, GramSheel and Project VIKRAM, Humanesque and its projections	chapter	not yet reviewed
ch-projections Projections: names, and raising a component into Humanesque	chapter	not yet reviewed
ch-04 The repositories	chapter	not yet reviewed
ch-06 Choosing a platform	chapter	not yet reviewed
ch-07 The baseline environment and a first run	chapter	not yet reviewed
ch-run Running every component	chapter	not yet reviewed
ch-08 Labs: containers now, launch scripts next	chapter	not yet reviewed
ch-09 Numbers inside the machine	chapter	not yet reviewed
ch-10 Characters inside the machine	chapter	not yet reviewed
ch-11 Writing systems, sound, and reversible transliteration	chapter	not yet reviewed
ch-12 Pāṇini's grammar as a formal system	chapter	not yet reviewed
ch-13 Compilers, from lexer to microcode	chapter	not yet reviewed
ch-14 Below the software: logic, chips, ternary and qubits	chapter	not yet reviewed
ch-15 Quantum computing and PANINIQ, with the mathematics	chapter	not yet reviewed

Object	Kind	Mark
ch-16 AI, from MYCIN to transformers, and how to use any AI here	chapter	not yet reviewed
ch-17 Open source, provenance and contributing	chapter	not yet reviewed
ch-18 GATE: sit it, whatever your age	chapter	not yet reviewed
ch-syllabus Syllabus links, and our syllabus	chapter	not yet reviewed
ch-19 Questions that come up, answered technically	chapter	not yet reviewed
ch-roadmap The roadmap, for every collaborator and every domain	chapter	not yet reviewed
ch-estimate What the work would cost: a model-based estimate	chapter	not yet reviewed
ch-20 This release and the next	chapter	not yet reviewed
app-a The estate's terms, and the nearest common terms	chapter	not yet reviewed
app-read Reading the architecture without flattening it	chapter	not yet reviewed
app-hum For researchers in the humanities and social sciences	chapter	not yet reviewed
app-stem For researchers in science, technology, engineering and mathematics	chapter	not yet reviewed
app-prac For practitioners: teachers, policy makers, communities, makers and engineers	chapter	not yet reviewed
app-math Formal definitions	chapter	not yet reviewed
app-review Review register	chapter	not yet reviewed
app-refs References	chapter	not yet reviewed
fig-stack The PANINI stack, with Romenagri as the script layer beneath every front end	figure	not yet reviewed
fig-bridge The symbol bridge: a Hindi program through hincc, GCC and the debugger, its names carried in Romenagri and rendered back	figure	not yet reviewed
fig-hub One identity across scripts: Brahmi scripts meet Devanagari in one hub; Perso-Arabic meets it at its written, canonical form	figure	not yet reviewed
fig-acts The three Acts, with their dividers; PANINI runs through all three	figure	not yet reviewed
fig-projections Humanesque and its projections: parallel cultural traversals, not stages	figure	not yet reviewed
fig-triad The AyeAI Triad: a triad with AyeAI at the apex, not a loop and not a pipeline	figure	not yet reviewed
fig-pedler PEDLER's branches: cognition, embodiment, physics and silicon	figure	not yet reviewed

Object	Kind	Mark
fig-loops Loops at four scales: PEDLER's, GENIE's, the realization spine's, and continuity as the write-ups read it (†)	figure	not yet reviewed
fig-provenance Provenance in its fixed order: seal, clear, attest, mint	figure	not yet reviewed
fig-spine The realization spine the realization arms share	figure	not yet reviewed
fig-collab How any collaborator, in any domain, goes from the guide to a sealed contribution	figure	not yet reviewed
fig-cyclers The cyclers: one engine, configured in PANINI, and the protocol every cycle follows	figure	not yet reviewed
fig-vikram GramSheel and Project VIKRAM, with VIKRAM's sections as its own pages name them	figure	not yet reviewed
fig-plan The plan's phases; the lanes run alongside from P3	figure	not yet reviewed
dependency-graph The dependency graph	figure	not yet reviewed
c-hindawi Hindawi and Romenagri	component card	not yet reviewed
c-urdu The Urdu edition	component card	not yet reviewed
c-tajziya tajziya	component card	not yet reviewed
c-panini PANINI, the prompt-cycle language	component card	not yet reviewed
c-panini_by_claude PANINI, parser and interpreter	component card	not yet reviewed
c-panini_by_grok PANINI, the stage-0 bootstrap	component card	not yet reviewed
c-humansque The Humansque merged release	component card	not yet reviewed
c-paniniq PANINIq	component card	not yet reviewed
c-paninib PANINIB, biological realization	component card	not yet reviewed
c-paniniphy PANINIphy, physical realization	component card	not yet reviewed
c-matba matba, the press	component card	not yet reviewed
c-khwab khwab, the cutting room	component card	not yet reviewed
c-awaz awaz, the listening room	component card	not yet reviewed
c-studios tilasm, pench and yadein	component card	not yet reviewed
c-genie GENIE	component card	not yet reviewed
c-alam alam	component card	not yet reviewed
c-mez Mez, the Cognitive Workbench	component card	not yet reviewed
c-fakir FAKIR	component card	not yet reviewed
c-dhancha Dhancha	component card	not yet reviewed
c-ertabat ertabat	component card	not yet reviewed
c-jugaad28 jugaad28	component card	not yet reviewed

Object	Kind	Mark
c-transeg TransEg	component card	not yet reviewed
c-idgov TransEg identity governance	component card	not yet reviewed
c-chakra CHAKRA	component card	not yet reviewed
c-jyotish Jyotish	component card	not yet reviewed
c-pratik PRATIK kernel	component card	not yet reviewed
c-zamin Zamin	component card	not yet reviewed
c-qedler QEDLER	component card	not yet reviewed
c-misty Misty DOI	component card	not yet reviewed
c-kundali Research Kundali	component card	not yet reviewed
c-labs The linguistics lab	component card	not yet reviewed
c-rahnuma This guide	component card	not yet reviewed
mod-F1 Numbers inside the machine	syllabus module	not yet reviewed
mod-F2 Letters, scripts and your language	syllabus module	not yet reviewed
mod-F3 The sky and the calendar	syllabus module	not yet reviewed
mod-F4 Making with a cycler, safely	syllabus module	not yet reviewed
mod-S1 Binary, logic and circuits	syllabus module	not yet reviewed
mod-S2 Encoding and reversible transliteration	syllabus module	not yet reviewed
mod-S3 Programming in your own language	syllabus module	not yet reviewed
mod-S4 Using AI and checking it	syllabus module	not yet reviewed
mod-H1 Grammar as a formal system	syllabus module	not yet reviewed
mod-H2 From source code to machine	syllabus module	not yet reviewed
mod-H3 Open source, provenance and contributing	syllabus module	not yet reviewed
mod-H4 A project seeded from FAKIR	syllabus module	not yet reviewed
mod-U1 Theory of computation and compilers	syllabus module	not yet reviewed
mod-U2 Digital logic and architecture	syllabus module	not yet reviewed
mod-U3 Quantum and physical computing	syllabus module	not yet reviewed
mod-P1 Research on the open questions	syllabus module	not yet reviewed
mod-Q1 Adopting a component	syllabus module	not yet reviewed
board-gate-cs GATE 2027, Computer Science and Information Technology	syllabus link	not yet reviewed
board-gate-ra GATE 2027, Robotics and Automation	syllabus link	not yet reviewed
board-gate-in GATE 2027, Instrumentation Engineering	syllabus link	not yet reviewed
board-gate-ga GATE 2027, General Aptitude	syllabus link	not yet reviewed
board-gate GATE 2027, every paper	syllabus link	not yet reviewed

Object	Kind	Mark
<code>board-cbse</code> CBSE, classes 9 to 12	syllabus link	not yet reviewed
<code>board-ncert</code> NCERT textbooks	syllabus link	not yet reviewed
<code>board-icse</code> ICSE, class 10	syllabus link	not yet reviewed
<code>board-isc</code> ISC, class 12	syllabus link	not yet reviewed
<code>board-nios</code> NIOS	syllabus link	not yet reviewed
<code>board-igcse</code> Cambridge IGCSE	syllabus link	not yet reviewed
<code>board-alevel</code> Cambridge International AS and A Level	syllabus link	not yet reviewed
<code>board-ib-dp</code> IB Diploma Programme	syllabus link	not yet reviewed
<code>board-ib-myp</code> IB Middle Years Programme	syllabus link	not yet reviewed
<code>board-ap</code> Advanced Placement courses	syllabus link	not yet reviewed
<code>board-ccss</code> Common Core State Standards	syllabus link	not yet reviewed
<code>board-aqa</code> AQA GCSE and A level	syllabus link	not yet reviewed
<code>board-edexcel</code> Pearson Edexcel qualifications	syllabus link	not yet reviewed
<code>board-ocr</code> OCR qualifications	syllabus link	not yet reviewed
<code>board-acara</code> Australian Curriculum, version 9	syllabus link	not yet reviewed
<code>board-seab</code> Singapore national examinations	syllabus link	not yet reviewed
<code>board-ts</code> Telangana Board of Secondary Education	syllabus link	not yet reviewed
<code>board-ap-bse</code> Andhra Pradesh Board of Secondary Education	syllabus link	not yet reviewed
<code>board-mh</code> Maharashtra State Board	syllabus link	not yet reviewed
<code>board-ka</code> Karnataka School Examination and Assessment Board	syllabus link	not yet reviewed
<code>board-tn</code> Tamil Nadu Directorate of Government Examinations	syllabus link	not yet reviewed
<code>board-kl</code> Kerala SCERT	syllabus link	not yet reviewed
<code>board-wb</code> West Bengal Board of Secondary Education	syllabus link	not yet reviewed
<code>board-up</code> Uttar Pradesh Madhyamik Shiksha Parishad	syllabus link	not yet reviewed
<code>board-rj</code> Board of Secondary Education, Rajasthan	syllabus link	not yet reviewed
<code>board-gj</code> Gujarat Secondary and Higher Secondary Education Board	syllabus link	not yet reviewed
<code>board-pb</code> Punjab School Education Board	syllabus link	not yet reviewed
<code>ed-role-school</code> Rahnuma, the edition for: School student	edition	not yet reviewed
<code>ed-role-college</code> Rahnuma, the edition for: College student	edition	not yet reviewed

Object	Kind	Mark
<code>ed-role-pg</code> Rahnuma, the edition for: Postgraduate or doctoral researcher	edition	not yet reviewed
<code>ed-role-teacher</code> Rahnuma, the edition for: Teacher or lecturer	edition	not yet reviewed
<code>ed-role-researcher</code> Rahnuma, the edition for: Academic researcher	edition	not yet reviewed
<code>ed-role-professional</code> Rahnuma, the edition for: Working professional or engineer	edition	not yet reviewed
<code>ed-role-public</code> Rahnuma, the edition for: Public service, policy or law	edition	not yet reviewed
<code>ed-role-community</code> Rahnuma, the edition for: Community, NGO or village work	edition	not yet reviewed
<code>ed-role-creator</code> Rahnuma, the edition for: Artist, writer or maker	edition	not yet reviewed
<code>ed-role-industry</code> Rahnuma, the edition for: Industry partner, investor or funder	edition	not yet reviewed
<code>ed-field-00</code> Rahnuma, the edition for the field: Generic programmes and qualifications	edition	not yet reviewed
<code>ed-field-01</code> Rahnuma, the edition for the field: Education	edition	not yet reviewed
<code>ed-field-02</code> Rahnuma, the edition for the field: Arts and humanities	edition	not yet reviewed
<code>ed-field-03</code> Rahnuma, the edition for the field: Social sciences, journalism and information	edition	not yet reviewed
<code>ed-field-04</code> Rahnuma, the edition for the field: Business, administration and law	edition	not yet reviewed
<code>ed-field-05</code> Rahnuma, the edition for the field: Natural sciences, mathematics and statistics	edition	not yet reviewed
<code>ed-field-06</code> Rahnuma, the edition for the field: Information and Communication Technologies (ICTs)	edition	not yet reviewed
<code>ed-field-07</code> Rahnuma, the edition for the field: Engineering, manufacturing and construction	edition	not yet reviewed
<code>ed-field-08</code> Rahnuma, the edition for the field: Agriculture, forestry, fisheries and veterinary	edition	not yet reviewed
<code>ed-field-09</code> Rahnuma, the edition for the field: Health and welfare	edition	not yet reviewed
<code>ed-field-10</code> Rahnuma, the edition for the field: Services	edition	not yet reviewed
<code>name-kaivalyik-the-ground-balanced-ternary-silicon</code> The ground: balanced-ternary silicon: Bhūmi भूमि	proposed name	not yet reviewed
<code>name-cosmopolis-the-ground-balanced-ternary-silicon</code> The ground: balanced-ternary silicon: Gaia Γαῖα	proposed name	not yet reviewed

Object	Kind	Mark
name-datong-the-ground-balanced-ternary-silicon The ground: balanced-ternary silicon: 地 dì, chí, jì 지, ðja	proposed name	not yet reviewed
name-vaka-the-ground-balanced-ternary-silicon The ground: balanced-ternary silicon: Whenua, Honua	proposed name	not yet reviewed
name-kaivalyik-the-water-the-factory-process The water: the factory process: Jala जल	proposed name	not yet reviewed
name-cosmopolis-the-water-the-factory-process The water: the factory process: Hydōr ὕδωρ	proposed name	not yet reviewed
name-datong-the-water-the-factory-process The water: the factory process: 水 shuǐ, suǐ, su 水, thùy	proposed name	not yet reviewed
name-vaka-the-water-the-factory-process The water: the factory process: Wai	proposed name	not yet reviewed
name-kaivalyik-the-air-the-environmental-replica The air: the environmental replica: Vāyu वायु	proposed name	not yet reviewed
name-cosmopolis-the-air-the-environmental-replica The air: the environmental replica: Aēr αἴρ	proposed name	not yet reviewed
name-datong-the-air-the-environmental-replica The air: the environmental replica: 風 fēng, fū, pung 풍, phong	proposed name	not yet reviewed
name-vaka-the-air-the-environmental-replica The air: the environmental replica: Hau, Makani	proposed name	not yet reviewed
name-kaivalyik-the-sky-the-observatory The sky: the observatory: Jyotish ज्योतिष (established)	proposed name	not yet reviewed
name-cosmopolis-the-sky-the-observatory The sky: the observatory: Ouranos Οὐρανός	proposed name	not yet reviewed
name-datong-the-sky-the-observatory The sky: the observatory: 天 tiān, ten, cheon 천, thiên	proposed name	not yet reviewed
name-vaka-the-sky-the-observatory The sky: the observatory: Rangī, Lani	proposed name	not yet reviewed
name-kaivalyik-the-calendars-of-the-sky The calendars of the sky: Pañcāṅga पञ्चाङ्ग	proposed name	not yet reviewed
name-cosmopolis-the-calendars-of-the-sky The calendars of the sky: Hēmerologion ἡμερολόγιον	proposed name	not yet reviewed
name-datong-the-calendars-of-the-sky The calendars of the sky: 曆 lì, reki, ryeok 력, lịch	proposed name	not yet reviewed
name-vaka-the-calendars-of-the-sky The calendars of the sky: Maramataka	proposed name	not yet reviewed
name-kaivalyik-the-shared-virtual-dome The shared virtual dome: Maṇḍapa मण्डप	proposed name	not yet reviewed

Object	Kind	Mark
name-cosmopolis-the-shared-virtual-dome The shared virtual dome: Tholos θόλος	proposed name	not yet reviewed
name-datong-the-shared-virtual-dome The shared virtual dome: 穹 qióng, kyū, gung 궁, khung	proposed name	not yet reviewed
name-vaka-the-shared-virtual-dome The shared virtual dome: Whare, Hale	proposed name	not yet reviewed
name-kaivalyik-the-desk-the-cognitive- workbench The desk, the Cognitive Workbench: Pīṭha पीठ	proposed name	not yet reviewed
name-cosmopolis-the-desk-the-cognitive- workbench The desk, the Cognitive Workbench: Trapeza τράπεζα	proposed name	not yet reviewed
name-datong-the-desk-the-cognitive- workbench The desk, the Cognitive Workbench: 卓 zhuó, taku, tak 탁, trác	proposed name	not yet reviewed
name-vaka-the-desk-the-cognitive- workbench The desk, the Cognitive Workbench: Papa	proposed name	not yet reviewed
name-kaivalyik-the-press The press: Mudraṇa मुद्रण	proposed name	not yet reviewed
name-cosmopolis-the-press The press: Typographeion τυπογραφεῖον	proposed name	not yet reviewed
name-datong-the-press The press: 印 yìn, in, in 인, ăn	proposed name	not yet reviewed
name-vaka-the-press The press: Tā	proposed name	not yet reviewed
name-kaivalyik-the-cutting-room-the- visual-cycler The cutting room, the visual cyclor: Svapna स्वप्न	proposed name	not yet reviewed
name-cosmopolis-the-cutting-room-the- visual-cycler The cutting room, the visual cyclor: Oneiros ὄνειρος	proposed name	not yet reviewed
name-datong-the-cutting-room-the- visual-cycler The cutting room, the visual cyclor: 夢 mèng, mu, mong 몽, mông	proposed name	not yet reviewed
name-vaka-the-cutting-room-the-visual- cyclor The cutting room, the visual cyclor: Moemoeā	proposed name	not yet reviewed
name-kaivalyik-the-listening-room-the- audio-cycler The listening room, the audio cyclor: Svara स्वर	proposed name	not yet reviewed
name-cosmopolis-the-listening-room-the- audio-cycler The listening room, the audio cyclor: Phōnē φωνή	proposed name	not yet reviewed
name-datong-the-listening-room-the- audio-cycler The listening room, the audio cyclor: 聲 shēng, sei, seong 성, thanh	proposed name	not yet reviewed
name-vaka-the-listening-room-the-audio- cyclor The listening room, the audio cyclor: Reo	proposed name	not yet reviewed

Object	Kind	Mark
name-kaivalyik-the-immersive-cycler The immersive cycler: Indrajāla इन्द्रजाल	proposed name	not yet reviewed
name-cosmopolis-the-immersive-cycler The immersive cycler: Phantasia φαντασία	proposed name	not yet reviewed
name-datong-the-immersive-cycler The immersive cycler: 幻 huàn, gen, hwan 환, huyên	proposed name	not yet reviewed
name-vaka-the-immersive-cycler The immersive cycler: Whakaata	proposed name	not yet reviewed
name-kaivalyik-the-embodied-cycler The embodied cycler: Yantra यन्त्र	proposed name	not yet reviewed
name-cosmopolis-the-embodied-cycler The embodied cycler: Mēchanē μηχανή	proposed name	not yet reviewed
name-datong-the-embodied-cycler The embodied cycler: 機 jī, kī, gī 기, cō	proposed name	not yet reviewed
name-vaka-the-embodied-cycler The embodied cycler: Mīhini	proposed name	not yet reviewed
name-kaivalyik-the-record The record: Smṛti स्मृति	proposed name	not yet reviewed
name-cosmopolis-the-record The record: Mnēmē μνήμη	proposed name	not yet reviewed
name-datong-the-record The record: 記 jì, kī, gī 기, ký	proposed name	not yet reviewed
name-vaka-the-record The record: Mahara	proposed name	not yet reviewed
name-kaivalyik-small-work-cycled-with-any-ai Small work, cycled with any AI: Loka लोक	proposed name	not yet reviewed
name-cosmopolis-small-work-cycled-with-any-ai Small work, cycled with any AI: Kosmos κόσμος	proposed name	not yet reviewed
name-datong-small-work-cycled-with-any-ai Small work, cycled with any AI: 界 jiè, kai, gye 계, gióri	proposed name	not yet reviewed
name-vaka-small-work-cycled-with-any-ai Small work, cycled with any AI: Ao	proposed name	not yet reviewed
name-kaivalyik-the-cycles The cycles: Āvarta आवर्त	proposed name	not yet reviewed
name-cosmopolis-the-cycles The cycles: Kyklos κύκλος	proposed name	not yet reviewed
name-datong-the-cycles The cycles: 環 huán, kan, hwan 환, hoàn	proposed name	not yet reviewed
name-vaka-the-cycles The cycles: Porohita	proposed name	not yet reviewed
name-kaivalyik-the-book-template The book template: Grantha ग्रन्थ	proposed name	not yet reviewed
name-cosmopolis-the-book-template The book template: Biblos βιβλος	proposed name	not yet reviewed
name-datong-the-book-template The book template: 書 shū, sho, seo 서, thur	proposed name	not yet reviewed

Object	Kind	Mark
name-vaka-the-book-template The book template: Pukapuka	proposed name	not yet reviewed
name-kaivalyik-the-guide The guide: Mārgadarśaka मार्गदर्शक	proposed name	not yet reviewed
name-cosmopolis-the-guide The guide: Periēgētēs περιηγητής	proposed name	not yet reviewed
name-datong-the-guide The guide: 導 dǎo, dō, do 도, đạo	proposed name	not yet reviewed
name-vaka-the-guide The guide: Kaiārahi	proposed name	not yet reviewed
name-kaivalyik-the-domain-spine The domain spine: Merudaṇḍa मेरुदण्ड	proposed name	not yet reviewed
name-cosmopolis-the-domain-spine The domain spine: Rhachis ράχις	proposed name	not yet reviewed
name-datong-the-domain-spine The domain spine: 骨 gǔ, kotsu, gol 골, cốt	proposed name	not yet reviewed
name-vaka-the-domain-spine The domain spine: Tuarā	proposed name	not yet reviewed
name-kaivalyik-communication-engineering Communication engineering: Saṃvāda संवाद	proposed name	not yet reviewed
name-cosmopolis-communication-engineering Communication engineering: Angelia ἄγγελια	proposed name	not yet reviewed
name-datong-communication-engineering Communication engineering: 信 xìn, shin, sin 신, tín	proposed name	not yet reviewed
name-vaka-communication-engineering Communication engineering: Kōrero	proposed name	not yet reviewed
name-kaivalyik-device-control Device control: Nāḍī नाडी	proposed name	not yet reviewed
name-cosmopolis-device-control Device control: Neuron νεῦρον	proposed name	not yet reviewed
name-datong-device-control Device control: 脈 mài, myaku, maek 맥, mạch	proposed name	not yet reviewed
name-vaka-device-control Device control: Uaua	proposed name	not yet reviewed
name-kaivalyik-the-door-for-any-ai The door for any AI: Dvāra द्वार	proposed name	not yet reviewed
name-cosmopolis-the-door-for-any-ai The door for any AI: Thyra θύρα	proposed name	not yet reviewed
name-datong-the-door-for-any-ai The door for any AI: 門 mén, mon, mun 문, môn	proposed name	not yet reviewed
name-vaka-the-door-for-any-ai The door for any AI: Tatau	proposed name	not yet reviewed
name-kaivalyik-scan-and-store Scan and store: Citra चित्र	proposed name	not yet reviewed
name-cosmopolis-scan-and-store Scan and store: Eikōn εἰκών	proposed name	not yet reviewed

Object	Kind	Mark
name-datong-scan-and-store Scan and store: 像 xiàng, zǒ, sang 상, tượg	proposed name	not yet reviewed
name-vaka-scan-and-store Scan and store: Whakaahua	proposed name	not yet reviewed
name-kaivalyik-the-plates The plates: Phalaka फलक	proposed name	not yet reviewed
name-cosmopolis-the-plates The plates: Pinax πίναξ	proposed name	not yet reviewed
name-datong-the-plates The plates: 板 bǎn, ban, pan 판, bản	proposed name	not yet reviewed
name-vaka-the-plates The plates: Pou	proposed name	not yet reviewed
name-kaivalyik-the-parser-federation The parser federation: Vyākaraṇa व्याकरण	proposed name	not yet reviewed
name-cosmopolis-the-parser-federation The parser federation: Analysis ἀνάλυσις	proposed name	not yet reviewed
name-datong-the-parser-federation The parser federation: 析 xī, seki, seok 석, tích	proposed name	not yet reviewed
name-vaka-the-parser-federation The parser federation: Wetewete	proposed name	not yet reviewed
name-kaivalyik-the-geometry-of-becoming The Geometry of Becoming: Parivartana परिवर्तन	proposed name	not yet reviewed
name-cosmopolis-the-geometry-of-becoming The Geometry of Becoming: Genesis γένεσις	proposed name	not yet reviewed
name-datong-the-geometry-of-becoming The Geometry of Becoming: 變 biàn, hen, byeon 변, biến	proposed name	not yet reviewed
name-vaka-the-geometry-of-becoming The Geometry of Becoming: Panoni	proposed name	not yet reviewed
name-kaivalyik-the-geometry-of-living The Geometry of Living: Jivana जीवन	proposed name	not yet reviewed
name-cosmopolis-the-geometry-of-living The Geometry of Living: Bios βίος	proposed name	not yet reviewed
name-datong-the-geometry-of-living The Geometry of Living: 生 shēng, sei, saeng 생, sinh	proposed name	not yet reviewed
name-vaka-the-geometry-of-living The Geometry of Living: Ora	proposed name	not yet reviewed
name-kaivalyik-the-geometry-of-persisting The Geometry of Persisting: Sthiti स्थिति	proposed name	not yet reviewed
name-cosmopolis-the-geometry-of-persisting The Geometry of Persisting: Monē μονή	proposed name	not yet reviewed
name-datong-the-geometry-of-persisting The Geometry of Persisting: 恒 héng, kó, hang 항, hằng	proposed name	not yet reviewed

Object	Kind	Mark
name-vaka-the-geometry-of-persisting The Geometry of Persisting: Mau	proposed name	not yet reviewed
name-kaivalyik-topological-bildung Topological Bildung: Śikṣā शिक्षा	proposed name	not yet reviewed
name-cosmopolis-topological-bildung Topological Bildung: Paideia παιδεία	proposed name	not yet reviewed
name-datong-topological-bildung Topological Bildung: 育 yù, iku, yuk 육, dục	proposed name	not yet reviewed
name-vaka-topological-bildung Topological Bildung: Ako	proposed name	not yet reviewed
name-kaivalyik-kaivalyik-agi-the- posters Kaivalyik AGI, the posters: Karma कर्म	proposed name	not yet reviewed
name-cosmopolis-kaivalyik-agi-the- posters Kaivalyik AGI, the posters: Ergon ἔργον	proposed name	not yet reviewed
name-datong-kaivalyik-agi-the-posters Kaivalyik AGI, the posters: 業 yè, gyō, eop 업, nghiep	proposed name	not yet reviewed
name-vaka-kaivalyik-agi-the-posters Kaivalyik AGI, the posters: Mahi	proposed name	not yet reviewed
name-kaivalyik-the-nm-tape-out- catalogue The 28 nm tape-out catalogue: Yukti युक्ति	proposed name	not yet reviewed
name-cosmopolis-the-nm-tape-out- catalogue The 28 nm tape-out catalogue: Mētis μῆτις	proposed name	not yet reviewed
name-datong-the-nm-tape-out-catalogue The 28 nm tape-out catalogue: 巧 qiǎo, kǎo, gyo 교, xáo	proposed name	not yet reviewed
name-vaka-the-nm-tape-out-catalogue The 28 nm tape-out catalogue: Mōhio	proposed name	not yet reviewed
name-kaivalyik-the-estate-s-register The estate's register: Sūcī सूची	proposed name	not yet reviewed
name-cosmopolis-the-estate-s-register The estate's register: Katalogos κατάλογος	proposed name	not yet reviewed
name-datong-the-estate-s-register The estate's register: 錄 lù, roku, rok 록, lục	proposed name	not yet reviewed
name-vaka-the-estate-s-register The estate's register: Rārangi	proposed name	not yet reviewed
name-kaivalyik-the-contracts The contracts: Niyama नियम	proposed name	not yet reviewed
name-cosmopolis-the-contracts The contracts: Nomos νόμος	proposed name	not yet reviewed
name-datong-the-contracts The contracts: 律 lǜ, ritsu, yul 율, luật	proposed name	not yet reviewed
name-vaka-the-contracts The contracts: Ture	proposed name	not yet reviewed

Object	Kind	Mark
name-kaivalyik-a-research-map A research map: Kuṇḍalī कुण्डली	proposed name	not yet reviewed
name-cosmopolis-a-research-map A research map: Horoskopos ὠροσκόπος	proposed name	not yet reviewed
name-datong-a-research-map A research map: 譜 pǔ, fu, bo 博, phō	proposed name	not yet reviewed
name-vaka-a-research-map A research map: Mahere	proposed name	not yet reviewed
name-kaivalyik-the-paradox-of-innovation The Paradox of Innovation: Virodhābhāsa विरोधाभास	proposed name	not yet reviewed
name-cosmopolis-the-paradox-of-innovation The Paradox of Innovation: Paradoxon παράδοξον	proposed name	not yet reviewed
name-datong-the-paradox-of-innovation The Paradox of Innovation: 矛盾 máodùn, mujun, mosun 모순, mâu thuẫn	proposed name	not yet reviewed
name-vaka-the-paradox-of-innovation The Paradox of Innovation: to be named with speakers	proposed name	not yet reviewed
name-kaivalyik-inclinations Inclinations: Pravṛtti प्रवृत्ति	proposed name	not yet reviewed
name-cosmopolis-inclinations Inclinations: Rhopē ῥοπή	proposed name	not yet reviewed
name-datong-inclinations Inclinations: 傾 qīng, kei, gyeong 경, khuynh	proposed name	not yet reviewed
name-vaka-inclinations Inclinations: to be named with speakers	proposed name	not yet reviewed
name-kaivalyik-health Health: Auśadha औषध	proposed name	not yet reviewed
name-cosmopolis-health Health: Pharmakon φάρμακον	proposed name	not yet reviewed
name-datong-health Health: 藥 yào, yaku, yak 약, dược	proposed name	not yet reviewed
name-vaka-health Health: Rongoā	proposed name	not yet reviewed
name-kaivalyik-exercise Exercise: Vyāyāma व्यायाम	proposed name	not yet reviewed
name-cosmopolis-exercise Exercise: Gymnasia γυμνασία	proposed name	not yet reviewed
name-datong-exercise Exercise: 練 liàn, ren, yeon 련, luyện	proposed name	not yet reviewed
name-vaka-exercise Exercise: Korikori	proposed name	not yet reviewed
name-kaivalyik-waking Waking: Jāgrti जागृति	proposed name	not yet reviewed
name-cosmopolis-waking Waking: Egersis ἐγερσις	proposed name	not yet reviewed
name-datong-waking Waking: 醒 xǐng, sei, seong 성, tỉnh	proposed name	not yet reviewed

Object	Kind	Mark
name-vaka-waking Waking: Oho	proposed name	not yet reviewed
name-kaivalyik-ilm-s-lexicon ILM's lexicon: Śabdakośa शब्दकोश	proposed name	not yet reviewed
name-cosmopolis-ilm-s-lexicon ILM's lexicon: Lexikon λεξικόν	proposed name	not yet reviewed
name-datong-ilm-s-lexicon ILM's lexicon: 辭 cí, jì, sa 사, tì	proposed name	not yet reviewed
name-vaka-ilm-s-lexicon ILM's lexicon: Papakupu	proposed name	not yet reviewed
name-kaivalyik-ilm-s-phonology ILM's phonology: Dhvani ध्वनि	proposed name	not yet reviewed
name-cosmopolis-ilm-s-phonology ILM's phonology: Phthongos φθόγγος	proposed name	not yet reviewed
name-datong-ilm-s-phonology ILM's phonology: 音 yīn, on, eum 음, âm	proposed name	not yet reviewed
name-vaka-ilm-s-phonology ILM's phonology: Oro	proposed name	not yet reviewed
name-kaivalyik-ilm-s-orthography ILM's orthography: Lipi लिपि	proposed name	not yet reviewed
name-cosmopolis-ilm-s-orthography ILM's orthography: Graphē γραφή	proposed name	not yet reviewed
name-datong-ilm-s-orthography ILM's orthography: 字 zì, jì, ja 자, tʻ	proposed name	not yet reviewed
name-vaka-ilm-s-orthography ILM's orthography: Tuhi	proposed name	not yet reviewed
name-kaivalyik-the-programming-system The programming system: Hindawi (established)	proposed name	not yet reviewed
name-cosmopolis-the-programming-system The programming system: HPS, Latin and Greek	proposed name	not yet reviewed
name-datong-the-programming-system The programming system: HPS, Han characters	proposed name	not yet reviewed
name-vaka-the-programming-system The programming system: HPS, Te Reo Māori	proposed name	not yet reviewed
name-kaivalyik-humansque humansque: same	proposed name	not yet reviewed
name-cosmopolis-humansque humansque: same	proposed name	not yet reviewed
name-datong-humansque humansque: same	proposed name	not yet reviewed
name-vaka-humansque humansque: same	proposed name	not yet reviewed
name-kaivalyik-ilm ilm: same	proposed name	not yet reviewed
name-cosmopolis-ilm ilm: same	proposed name	not yet reviewed
name-datong-ilm ilm: same	proposed name	not yet reviewed
name-vaka-ilm ilm: same	proposed name	not yet reviewed

Object	Kind	Mark
name-kaivalyik-romenagri romenagri: same	proposed name	not yet reviewed
name-cosmopolis-romenagri romenagri: same	proposed name	not yet reviewed
name-datong-romenagri romenagri: same	proposed name	not yet reviewed
name-vaka-romenagri romenagri: same	proposed name	not yet reviewed
name-kaivalyik-panini panini: same	proposed name	not yet reviewed
name-cosmopolis-panini panini: same	proposed name	not yet reviewed
name-datong-panini panini: same	proposed name	not yet reviewed
name-vaka-panini panini: same	proposed name	not yet reviewed
name-kaivalyik-paniniq paniniq: same	proposed name	not yet reviewed
name-cosmopolis-paniniq paniniq: same	proposed name	not yet reviewed
name-datong-paniniq paniniq: same	proposed name	not yet reviewed
name-vaka-paniniq paniniq: same	proposed name	not yet reviewed
name-kaivalyik-paninib paninib: same	proposed name	not yet reviewed
name-cosmopolis-paninib paninib: same	proposed name	not yet reviewed
name-datong-paninib paninib: same	proposed name	not yet reviewed
name-vaka-paninib paninib: same	proposed name	not yet reviewed
name-kaivalyik-paniniphy paniniphy: same	proposed name	not yet reviewed
name-cosmopolis-paniniphy paniniphy: same	proposed name	not yet reviewed
name-datong-paniniphy paniniphy: same	proposed name	not yet reviewed
name-vaka-paniniphy paniniphy: same	proposed name	not yet reviewed
name-kaivalyik-pedler pedler: same	proposed name	not yet reviewed
name-cosmopolis-pedler pedler: same	proposed name	not yet reviewed
name-datong-pedler pedler: same	proposed name	not yet reviewed
name-vaka-pedler pedler: same	proposed name	not yet reviewed
name-kaivalyik-qedler qedler: same	proposed name	not yet reviewed
name-cosmopolis-qedler qedler: same	proposed name	not yet reviewed
name-datong-qedler qedler: same	proposed name	not yet reviewed
name-vaka-qedler qedler: same	proposed name	not yet reviewed
name-kaivalyik-pratik pratik: same	proposed name	not yet reviewed
name-cosmopolis-pratik pratik: same	proposed name	not yet reviewed
name-datong-pratik pratik: same	proposed name	not yet reviewed
name-vaka-pratik pratik: same	proposed name	not yet reviewed
name-kaivalyik-fakir fakir: same	proposed name	not yet reviewed
name-cosmopolis-fakir fakir: same	proposed name	not yet reviewed
name-datong-fakir fakir: same	proposed name	not yet reviewed

Object	Kind	Mark
name-vaka-fakir fakir: same	proposed name	not yet reviewed
name-kaivalyik-ukop ukop: same	proposed name	not yet reviewed
name-cosmopolis-ukop ukop: same	proposed name	not yet reviewed
name-datong-ukop ukop: same	proposed name	not yet reviewed
name-vaka-ukop ukop: same	proposed name	not yet reviewed
name-kaivalyik-genie genie: same	proposed name	not yet reviewed
name-cosmopolis-genie genie: same	proposed name	not yet reviewed
name-datong-genie genie: same	proposed name	not yet reviewed
name-vaka-genie genie: same	proposed name	not yet reviewed
name-kaivalyik-transeg transeg: same	proposed name	not yet reviewed
name-cosmopolis-transeg transeg: same	proposed name	not yet reviewed
name-datong-transeg transeg: same	proposed name	not yet reviewed
name-vaka-transeg transeg: same	proposed name	not yet reviewed
name-kaivalyik-transeg-idgov transeg-idgov: same	proposed name	not yet reviewed
name-cosmopolis-transeg-idgov transeg-idgov: same	proposed name	not yet reviewed
name-datong-transeg-idgov transeg-idgov: same	proposed name	not yet reviewed
name-vaka-transeg-idgov transeg-idgov: same	proposed name	not yet reviewed
name-kaivalyik-copa copa: same	proposed name	not yet reviewed
name-cosmopolis-copa copa: same	proposed name	not yet reviewed
name-datong-copa copa: same	proposed name	not yet reviewed
name-vaka-copa copa: same	proposed name	not yet reviewed
name-kaivalyik-vgc vgc: same	proposed name	not yet reviewed
name-cosmopolis-vgc vgc: same	proposed name	not yet reviewed
name-datong-vgc vgc: same	proposed name	not yet reviewed
name-vaka-vgc vgc: same	proposed name	not yet reviewed
name-kaivalyik-candor candor: same	proposed name	not yet reviewed
name-cosmopolis-candor candor: same	proposed name	not yet reviewed
name-datong-candor candor: same	proposed name	not yet reviewed
name-vaka-candor candor: same	proposed name	not yet reviewed
name-kaivalyik-misty-doi misty-doi: same	proposed name	not yet reviewed
name-cosmopolis-misty-doi misty-doi: same	proposed name	not yet reviewed
name-datong-misty-doi misty-doi: same	proposed name	not yet reviewed
name-vaka-misty-doi misty-doi: same	proposed name	not yet reviewed
name-kaivalyik-tok-doi tok-doi: same	proposed name	not yet reviewed

Object	Kind	Mark
name-cosmopolis-tok-doi tok-doi: same	proposed name	not yet reviewed
name-datong-tok-doi tok-doi: same	proposed name	not yet reviewed
name-vaka-tok-doi tok-doi: same	proposed name	not yet reviewed
name-kaivalyik-spiguard spiguard: same	proposed name	not yet reviewed
name-cosmopolis-spiguard spiguard: same	proposed name	not yet reviewed
name-datong-spiguard spiguard: same	proposed name	not yet reviewed
name-vaka-spiguard spiguard: same	proposed name	not yet reviewed
name-kaivalyik-dukedom dukedom: same	proposed name	not yet reviewed
name-cosmopolis-dukedom dukedom: same	proposed name	not yet reviewed
name-datong-dukedom dukedom: same	proposed name	not yet reviewed
name-vaka-dukedom dukedom: same	proposed name	not yet reviewed
name-kaivalyik-duke duke2: same	proposed name	not yet reviewed
name-cosmopolis-duke duke2: same	proposed name	not yet reviewed
name-datong-duke duke2: same	proposed name	not yet reviewed
name-vaka-duke duke2: same	proposed name	not yet reviewed
name-kaivalyik-janapad janapad: same	proposed name	not yet reviewed
name-cosmopolis-janapad janapad: same	proposed name	not yet reviewed
name-datong-janapad janapad: same	proposed name	not yet reviewed
name-vaka-janapad janapad: same	proposed name	not yet reviewed
name-kaivalyik-poie poie: same	proposed name	not yet reviewed
name-cosmopolis-poie poie: same	proposed name	not yet reviewed
name-datong-poie poie: same	proposed name	not yet reviewed
name-vaka-poie poie: same	proposed name	not yet reviewed

H. References

In plain words. The books and papers this guide draws on.

- Aho, A. V., Lam, M. S., Sethi, R. and Ullman, J. D. (2006). *Compilers: Principles, Techniques, and Tools*, 2nd edition. Addison-Wesley.
- Bureau of Indian Standards (1991). IS 13194:1991, Indian Script Code for Information Interchange (ISCII).
- Cardona, G. (1997). *Pāṇini: His Work and its Traditions*, volume 1, 2nd edition. Motilal Banarsidass.
- Chaitin, G. J. (1982). Register allocation and spilling via graph coloring. *Proceedings of the SIGPLAN Symposium on Compiler Construction*, 98 to 105.
- Cusumano, M. A. (1991). *Japan's Software Factories: A Challenge to U.S. Management*. Oxford University Press.
- Cytron, R., Ferrante, J., Rosen, B. K., Wegman, M. N. and Zadeck, F. K. (1991). Efficiently computing static single assignment form and the control dependence graph. *ACM Transactions on Programming Languages and Systems* 13(4), 451 to 490.
- Farhi, E. and Gutmann, S. (1998). Quantum computation and decision trees. *Physical Review A* 58, 915.
- Farhi, E., Goldstone, J. and Gutmann, S. (2014). A quantum approximate optimization algorithm. arXiv:1411.4028.
- Gandhi, M. K. (1909). *Hind Swaraj, or Indian Home Rule*. Navajivan Publishing House, Ahmedabad.
- Gandhi, M. K. (1941). *Constructive Programme: Its Meaning and Place*. Navajivan Publishing House, Ahmedabad.
- Goemans, M. X. and Williamson, D. P. (1995). Improved approximation algorithms for maximum cut and satisfiability problems using semidefinite programming. *Journal of the ACM* 42(6), 1115 to 1145.
- Ingerman, P. Z. (1967). “Pāṇini-Backus Form” suggested. *Communications of the ACM* 10(3), 137.
- Katre, S. M. (1987). *Aṣṭādhyāyī of Pāṇini*. University of Texas Press.
- Kiparsky, P. (1991). Economy and the construction of the Śivasūtras. In M. M. Deshpande and S. Bhate (eds.), *Pāṇinian Studies*. University of Michigan.
- Koch, J. and colleagues (2007). Charge-insensitive qubit design derived from the Cooper pair box. *Physical Review A* 76, 042319.
- Kuramoto, Y. (1984). *Chemical Oscillations, Waves, and Turbulence*. Springer.
- McIlroy, M. D. (1969). “Mass produced software components”. In P. Naur and B. Randell (eds.), *Software Engineering: Report of a conference sponsored by the NATO Science Committee, Garmisch, 1968*, 138 to 155.
- McMillan, B. (1956). Two inequalities implied by unique decipherability. *IRE Transactions on Information Theory* 2(4), 115 to 116.
- Nielsen, M. A. and Chuang, I. L. (2010). *Quantum Computation and Quantum Information*, 10th anniversary edition. Cambridge University Press.
- Petersen, W. (2004). A mathematical analysis of Pāṇini's Śivasūtras. *Journal of Logic, Language and Information* 13, 471 to 489.
- Sardinas, A. A. and Patterson, G. W. (1953). A necessary and sufficient condition for the unique decomposition of coded messages. *IRE International Convention Record* 8, 104 to 108.
- Shortliffe, E. H. (1976). *Computer-Based Medical Consultations: MYCIN*. Elsevier.
- Strogatz, S. H. (2000). From Kuramoto to Crawford: exploring the onset of synchronization in populations of coupled oscillators. *Physica D* 143, 1 to 20.
- Thompson, K. (1984). Reflections on trusting trust. *Communications of the ACM* 27(8), 761 to 763.
- Unicode Consortium. *The Unicode Standard*; Unicode Standard Annex #29, Unicode Text Segmentation; #31, Unicode Identifiers and Syntax; Unicode Technical Standard #39, Unicode Security Mechanisms.
- Vaswani, A. and colleagues (2017). Attention is all you need. *Advances in Neural Information Processing Systems* 30.
- Wang, T. and Roychowdhury, J. (2019). OIM: oscillator-based Ising machines for solving combinatorial optimisation problems. *Unconventional Computation and Natural Computation*, Lecture Notes in Computer Science 11493, 232 to 256.
- Wheeler, D. A. (2009). *Fully Countering Trusting Trust through Diverse Double-Compiling*. PhD dissertation, George Mason University.

